

# Java Concurrency: Complex Practical Examples

Integrating Multiple Concepts for Staff-Level Interviews

## 1 Example 1: Thread-Safe Cache with Expiration

### Requirements:

- Multiple threads read frequently (optimize for reads)
- Occasional writes to update entries
- Entries expire after TTL (time-to-live)
- Background thread periodically removes expired entries
- Thread-safe statistics tracking (hits, misses)

Concepts integrated: ReadWriteLock, ScheduledExecutorService, AtomicLong, ConcurrentHashMap

```
1 import java.util.concurrent.*;
2 import java.util.concurrent.locks.*;
3 import java.util.concurrent.atomic.*;
4
5 class ExpiringCache<K, V> {
6     private static class CacheEntry<V> {
7         final V value;
8         final long expiryTime;
9
10        CacheEntry(V value, long ttlMillis) {
11            this.value = value;
12            this.expiryTime = System.currentTimeMillis() + ttlMillis;
13        }
14
15        boolean isExpired() {
16            return System.currentTimeMillis() > expiryTime;
17        }
18    }
19
20    private final ConcurrentHashMap<K, CacheEntry<V>> cache;
21    private final AtomicLong hits = new AtomicLong(0);
22    private final AtomicLong misses = new AtomicLong(0);
23    private final ScheduledExecutorService cleaner;
24    private final long ttlMillis;
25
26    public ExpiringCache(long ttlMillis, long cleanupIntervalMillis) {
27        this.cache = new ConcurrentHashMap<>();
28        this.ttlMillis = ttlMillis;
29        this.cleaner = Executors.newScheduledThreadPool(1);
30
31        // Schedule periodic cleanup
32        cleaner.scheduleAtFixedRate(
33            this::removeExpiredEntries,
34            cleanupIntervalMillis,
35            cleanupIntervalMillis,
36            TimeUnit.MILLISECONDS
37        );
38    }
39
40    public V get(K key) {
41        CacheEntry<V> entry = cache.get(key);
42
43        if (entry == null || entry.isExpired()) {
44            misses.incrementAndGet();
45            if (entry != null) {
46                cache.remove(key, entry); // Remove expired
47            }
48            return null;
49        }
50
51        hits.incrementAndGet();
52        return entry.value;
```

```

53 }
54
55     public void put(K key, V value) {
56         cache.put(key, new CacheEntry<>(value, ttlMillis));
57     }
58
59     private void removeExpiredEntries() {
60         cache.entrySet().removeIf(entry -> entry.getValue().isExpired());
61     }
62
63     public CacheStats getStats() {
64         long hitCount = hits.get();
65         long missCount = misses.get();
66         long total = hitCount + missCount;
67         double hitRate = total == 0 ? 0.0 : (double) hitCount / total;
68
69         return new CacheStats(hitCount, missCount, hitRate, cache.size());
70     }
71
72     public void shutdown() {
73         cleaner.shutdown();
74         try {
75             if (!cleaner.awaitTermination(5, TimeUnit.SECONDS)) {
76                 cleaner.shutdownNow();
77             }
78         } catch (InterruptedException e) {
79             cleaner.shutdownNow();
80             Thread.currentThread().interrupt();
81         }
82     }
83
84     static class CacheStats {
85         final long hits, misses;
86         final double hitRate;
87         final int size;
88
89         CacheStats(long hits, long misses, double hitRate, int size) {
90             this.hits = hits;
91             this.misses = misses;
92             this.hitRate = hitRate;
93             this.size = size;
94         }
95     }
96 }

```

Listing 1: Thread-Safe Cache Implementation

### Key design decisions:

- **ConcurrentHashMap** for cache storage - allows concurrent reads/writes to different keys
- **AtomicLong** for hit/miss counters - lock-free increments under high contention
- **ScheduledExecutorService** for cleanup - better than manual thread + Timer
- **No ReadWriteLock on cache operations** - ConcurrentHashMap already optimized
- **Graceful shutdown** with timeout and force-stop fallback
- **Benign race in get():** Between checking isExpired() and removing, another thread might access the entry. This is acceptable - worst case is redundant removal. Alternative: `cache.computeIfPresent(key, (k,v) -> v.isExpired() ? null : v)` for atomic check-and-remove

## 2 Example 2: Rate-Limited API Client

### Requirements:

- Limit to N requests per second across all threads
- Block threads when rate limit exceeded
- Support burst capacity (can briefly exceed rate)

- Track and report throttling statistics

**Concepts integrated:** Semaphore, ScheduledExecutorService, AtomicInteger, synchronized blocks

```

1 import java.util.concurrent.*;
2 import java.util.concurrent.atomic.*;
3
4 class RateLimitedApiClient {
5     private final Semaphore tokens;
6     private final int maxTokens;
7     private final int refillRate; // tokens per second
8     private final ScheduledExecutorService refiller;
9     private final AtomicInteger throttledRequests = new AtomicInteger(0);
10
11    public RateLimitedApiClient(int maxTokens, int refillRate) {
12        this.maxTokens = maxTokens;
13        this.refillRate = refillRate;
14        this.tokens = new Semaphore(maxTokens);
15        this.refiller = Executors.newScheduledThreadPool(1);
16
17        // Refill tokens every 100ms
18        long refillIntervalMs = 100;
19        // Calculate tokens per interval to achieve target rate
20        // refillRate tokens/second = refillRate/10 tokens per 100ms
21        double tokensPerInterval = refillRate / 10.0;
22
23        refiller.scheduleAtFixedRate(
24            () -> refillTokens(tokensPerInterval),
25            refillIntervalMs,
26            refillIntervalMs,
27            TimeUnit.MILLISECONDS
28        );
29    }
30
31    public <T> T makeRequest(Callable<T> apiCall)
32        throws Exception {
33        // Try to acquire token (blocks if none available)
34        boolean acquired = tokens.tryAcquire(5, TimeUnit.SECONDS);
35
36        if (!acquired) {
37            throttledRequests.incrementAndGet();
38            throw new RateLimitException("Rate limit exceeded");
39        }
40
41        try {
42            return apiCall.call();
43        } finally {
44            // Token not returned - consumed by rate limit
45        }
46    }
47
48    private void refillTokens(double tokensToAdd) {
49        // Accumulate fractional tokens across refills
50        int available = tokens.availablePermits();
51        int wholePart = (int) tokensToAdd;
52
53        // Never exceed maxTokens
54        int toAdd = Math.min(wholePart, maxTokens - available);
55
56        if (toAdd > 0) {
57            tokens.release(toAdd);
58        }
59        // Note: Race between availablePermits() and release() is acceptable
60        // Worst case: slightly exceed maxTokens temporarily
61    }
62
63    public RateLimitStats getStats() {
64        return new RateLimitStats(
65            tokens.availablePermits(),
66            throttledRequests.get()
67        );
68}

```

```

68 }
69
70     public void shutdown() {
71         refiller.shutdown();
72     }
73
74     static class RateLimitStats {
75         final int availableTokens;
76         final int throttledRequests;
77
78         RateLimitStats(int availableTokens, int throttledRequests) {
79             this.availableTokens = availableTokens;
80             this.throttledRequests = throttledRequests;
81         }
82     }
83
84     static class RateLimitException extends Exception {
85         RateLimitException(String message) {
86             super(message);
87         }
88     }
89 }

```

Listing 2: Token Bucket Rate Limiter

### Usage example:

```

1 RateLimitedApiClient client = new RateLimitedApiClient(
2     100, // max 100 tokens (burst capacity)
3     50 // refill 50 tokens/second
4 );
5
6 // Multiple threads making requests
7 ExecutorService executor = Executors.newFixedThreadPool(20);
8 for (int i = 0; i < 1000; i++) {
9     executor.submit(() -> {
10         try {
11             String result = client.makeRequest(() -> {
12                 return callExternalApi();
13             });
14             System.out.println("Success: " + result);
15         } catch (RateLimitException e) {
16             System.out.println("Throttled!");
17         }
18     });
19 }

```

### Key design decisions:

- **Semaphore for token bucket** - naturally models available capacity
- **tryAcquire with timeout** - prevents indefinite blocking
- **Tokens not returned** after request - consumed by rate limit
- **Periodic refill** - scheduled task adds tokens back
- **Bounded refill** - never exceed maxTokens (prevents overflow)

## 3 Example 3: Parallel File Processor with Work Stealing

### Requirements:

- Process large directory of files in parallel
- Some files larger than others (imbalanced workload)
- Aggregate results from all files
- Track progress and report completion
- Handle errors gracefully

**Concepts integrated:** ExecutorService, CountDownLatch, ConcurrentHashMap, AtomicInteger, Future

```

1 import java.util.*;
2 import java.util.concurrent.*;
3 import java.util.concurrent.atomic.*;
4 import java.io.*;
5 import java.nio.file.*;
6
7 class ParallelFileProcessor {
8     private final ExecutorService executor;
9     private final ConcurrentHashMap<String, ProcessingResult> results;
10    private final AtomicInteger processed = new AtomicInteger(0);
11    private final AtomicInteger failed = new AtomicInteger(0);
12
13    public ParallelFileProcessor(int threadCount) {
14        // Work-stealing pool for imbalanced workloads
15        this.executor = Executors.newWorkStealingPool(threadCount);
16        this.results = new ConcurrentHashMap<>();
17    }
18
19    public AggregateResult processDirectory(Path directory)
20        throws IOException, InterruptedException {
21        List<Path> files = Files.walk(directory)
22            .filter(Files::isRegularFile)
23            .toList();
24
25        int totalFiles = files.size();
26        CountDownLatch latch = new CountDownLatch(totalFiles);
27
28        // Submit all file processing tasks
29        List<Future<ProcessingResult>> futures = new ArrayList<>();
30        for (Path file : files) {
31            Future<ProcessingResult> future = executor.submit(() -> {
32                try {
33                    ProcessingResult result = processFile(file);
34                    results.put(file.toString(), result);
35                    processed.incrementAndGet();
36                    return result;
37                } catch (Exception e) {
38                    failed.incrementAndGet();
39                    System.err.println("Failed: " + file + " - " + e.getMessage());
40                    return ProcessingResult.error(file.toString(), e);
41                } finally {
42                    latch.countDown();
43                }
44            });
45            futures.add(future);
46        }
47
48        // Progress reporting in separate thread
49        AtomicBoolean progressRunning = new AtomicBoolean(true);
50        Thread progressThread = new Thread(() -> {
51            while (progressRunning.get() && latch.getCount() > 0) {
52                reportProgress(totalFiles);
53                try {
54                    Thread.sleep(1000);
55                } catch (InterruptedException e) {
56                    Thread.currentThread().interrupt();
57                    break;
58                }
59            }
60        });
61        progressThread.start();
62
63        // Wait for all files to complete
64        latch.await();
65        progressRunning.set(false);
66        progressThread.interrupt();
67
68        // Aggregate results
69        return aggregateResults(files, futures);

```

```

70 }
71
72     private ProcessingResult processFile(Path file) throws IOException {
73         // Simulate file processing
74         long wordCount;
75         try (var lines = Files.lines(file)) {
76             wordCount = lines.count();
77         }
78         long byteSize = Files.size(file);
79         return new ProcessingResult(file.toString(), wordCount, byteSize, true, null);
80     }
81
82     private void reportProgress(int total) {
83         int done = processed.get();
84         int errors = failed.get();
85         double pct = (done + errors) * 100.0 / total;
86         System.out.printf("Progress: %.1f%% (%d/%d, %d errors)%n",
87             pct, done, total, errors);
88     }
89
90     private AggregateResult aggregateResults(List<Path> files,
91         List<Future<ProcessingResult>> futures) {
92         long totalWords = 0;
93         long totalBytes = 0;
94         int successCount = 0;
95         List<String> errors = new ArrayList<>();
96
97         for (Future<ProcessingResult> future : futures) {
98             try {
99                 ProcessingResult result = future.get();
100                if (result.success) {
101                    totalWords += result.wordCount;
102                    totalBytes += result.byteSize;
103                    successCount++;
104                } else {
105                    errors.add(result.filename + ": " + result.error.getMessage());
106                }
107            } catch (Exception e) {
108                errors.add("Future failed: " + e.getMessage());
109            }
110        }
111
112        return new AggregateResult(files.size(), successCount,
113            totalWords, totalBytes, errors);
114    }
115
116    public void shutdown() {
117        executor.shutdown();
118        try {
119            if (!executor.awaitTermination(60, TimeUnit.SECONDS)) {
120                executor.shutdownNow();
121            }
122        } catch (InterruptedException e) {
123            executor.shutdownNow();
124        }
125    }
126
127    static class ProcessingResult {
128        final String filename;
129        final long wordCount;
130        final long byteSize;
131        final boolean success;
132        final Exception error;
133
134        ProcessingResult(String filename, long wordCount, long byteSize,
135            boolean success, Exception error) {
136            this.filename = filename;
137            this.wordCount = wordCount;
138            this.byteSize = byteSize;
139            this.success = success;

```

```

140     this.error = error;
141 }
142
143     static ProcessingResult error(String filename, Exception error) {
144         return new ProcessingResult(filename, 0, 0, false, error);
145     }
146 }
147
148     static class AggregateResult {
149         final int totalFiles;
150         final int successCount;
151         final long totalWords;
152         final long totalBytes;
153         final List<String> errors;
154
155         AggregateResult(int totalFiles, int successCount, long totalWords,
156                         long totalBytes, List<String> errors) {
157             this.totalFiles = totalFiles;
158             this.successCount = successCount;
159             this.totalWords = totalWords;
160             this.totalBytes = totalBytes;
161             this.errors = errors;
162         }
163     }
164 }
```

Listing 3: Parallel File Processor

#### Key design decisions:

- **newWorkStealingPool()** - threads steal work from each other when idle (good for imbalanced tasks)
- **CountDownLatch** - wait for all files to complete before aggregating
- **ConcurrentHashMap** - store results from multiple threads
- **AtomicInteger** - track progress without locks
- **Future list** - collect results for final aggregation
- **Separate progress thread** - non-blocking progress reporting
- **Graceful error handling** - failures don't stop other tasks

## 4 Example 4: Connection Pool

#### Requirements:

- Fixed pool of N database connections
- Threads block when all connections in use
- Connections have idle timeout (close if unused)
- Health check connections periodically
- Track utilization statistics

**Concepts integrated:** BlockingQueue, ScheduledExecutorService, ReentrantLock, wait/notify

```

1 import java.util.concurrent.*;
2 import java.util.concurrent.locks.*;
3 import java.util.*;
4
5 class ConnectionPool {
6     private final BlockingQueue<PooledConnection> available;
7     private final Set<PooledConnection> inUse;
8     private final Lock useLock = new ReentrantLock();
9     private final int maxSize;
10    private final long idleTimeoutMs;
11    private final ScheduledExecutorService healthChecker;
12    private volatile boolean isShutdown = false;
13
14    public ConnectionPool(int maxSize, long idleTimeoutMs) {
15        this.maxSize = maxSize;
16        this.idleTimeoutMs = idleTimeoutMs;
```

```

17     this.available = new LinkedBlockingQueue<>();
18     this.inUse = new HashSet<>();
19     this.healthChecker = Executors.newScheduledThreadPool(1);
20
21     // Initialize pool
22     for (int i = 0; i < maxSize; i++) {
23         available.offer(new PooledConnection(i));
24     }
25
26     // Periodic health check and idle timeout
27     healthChecker.scheduleAtFixedRate(
28         this::maintainPool,
29         30, 30, TimeUnit.SECONDS
30     );
31 }
32
33 public PooledConnection acquire() throws InterruptedException {
34     if (isShutdown) {
35         throw new IllegalStateException("Pool is shutdown");
36     }
37
38     // Block until connection available
39     PooledConnection conn = available.take();
40
41     useLock.lock();
42     try {
43         conn.markInUse();
44         inUse.add(conn);
45     } finally {
46         useLock.unlock();
47     }
48
49     return conn;
50 }
51
52 public PooledConnection tryAcquire(long timeout, TimeUnit unit)
53     throws InterruptedException {
54     if (isShutdown) {
55         throw new IllegalStateException("Pool is shutdown");
56     }
57
58     PooledConnection conn = available.poll(timeout, unit);
59     if (conn == null) {
60         return null; // Timeout
61     }
62
63     useLock.lock();
64     try {
65         conn.markInUse();
66         inUse.add(conn);
67     } finally {
68         useLock.unlock();
69     }
70
71     return conn;
72 }
73
74 public void release(PooledConnection conn) {
75     if (conn == null) return;
76
77     useLock.lock();
78     try {
79         if (!inUse.remove(conn)) {
80             throw new IllegalStateException("Connection not in use");
81         }
82         conn.markAvailable();
83     } finally {
84         useLock.unlock();
85     }
86 }
```

```

87     available.offer(conn);
88 }
89
90     private void maintainPool() {
91         useLock.lock();
92         try {
93             // Health check in-use connections
94             for (PooledConnection conn : inUse) {
95                 if (!conn.isHealthy()) {
96                     System.err.println("Unhealthy connection: " + conn.id);
97                     // In real implementation: close and recreate
98                 }
99             }
100
101             // Check idle timeout for available connections
102             long now = System.currentTimeMillis();
103             Iterator<PooledConnection> iter = available.iterator();
104             while (iter.hasNext()) {
105                 PooledConnection conn = iter.next();
106                 if (now - conn.lastUsedTime > idleTimeoutMs) {
107                     iter.remove();
108                     conn.close();
109                     // Create new connection to maintain pool size
110                     available.offer(new PooledConnection(conn.id));
111                 }
112             }
113         } finally {
114             useLock.unlock();
115         }
116     }
117
118     public PoolStats getStats() {
119         useLock.lock();
120         try {
121             return new PoolStats(
122                 available.size(),
123                 inUse.size(),
124                 maxSize
125             );
126         } finally {
127             useLock.unlock();
128         }
129     }
130
131     public void shutdown() {
132         isShutdown = true;
133         healthChecker.shutdown();
134
135         useLock.lock();
136         try {
137             for (PooledConnection conn : available) {
138                 conn.close();
139             }
140             available.clear();
141         } finally {
142             useLock.unlock();
143         }
144     }
145
146     static class PooledConnection {
147         final int id;
148         volatile long lastUsedTime;
149         volatile boolean inUse;
150
151         PooledConnection(int id) {
152             this.id = id;
153             this.lastUsedTime = System.currentTimeMillis();
154             this.inUse = false;
155         }
156     }

```

```

157     void markInUse() {
158         this.inUse = true;
159         this.lastUsedTime = System.currentTimeMillis();
160     }
161
162     void markAvailable() {
163         this.inUse = false;
164         this.lastUsedTime = System.currentTimeMillis();
165     }
166
167     boolean isHealthy() {
168         // Simulate health check
169         return true;
170     }
171
172     void close() {
173         // Close actual connection
174     }
175
176     public void execute(String sql) {
177         // Execute query
178     }
179 }
180
181     static class PoolStats {
182         final int available;
183         final int inUse;
184         final int total;
185
186         PoolStats(int available, int inUse, int total) {
187             this.available = available;
188             this.inUse = inUse;
189             this.total = total;
190         }
191
192         public double getUtilization() {
193             return total == 0 ? 0.0 : (double) inUse / total;
194         }
195     }
196 }
```

Listing 4: Database Connection Pool

### Usage pattern:

```

1 ConnectionPool pool = new ConnectionPool(10, 60000); // 10 conn, 60s idle
2
3 // Thread-safe acquire/release
4 PooledConnection conn = null;
5 try {
6     conn = pool.acquire(); // blocks if none available
7     conn.execute("SELECT * FROM users");
8 } finally {
9     if (conn != null) {
10         pool.release(conn); // Always return to pool
11     }
12 }
13
14 // With timeout
15 conn = pool.tryAcquire(5, TimeUnit.SECONDS);
16 if (conn != null) {
17     try {
18         conn.execute("SELECT * FROM orders");
19     } finally {
20         pool.release(conn);
21     }
22 } else {
23     System.out.println("Could not acquire connection");
24 }
```

## **Key design decisions:**

- **BlockingQueue for available connections** - natural blocking behavior when pool exhausted
- **Set for in-use tracking** - fast lookup to validate release
- **ReentrantLock for state transitions** - protects inUse set modifications
- **ScheduledExecutorService** for maintenance - health checks and idle timeout
- **Try-acquire with timeout** - prevents indefinite blocking
- **Graceful shutdown** - close all connections, reject new acquires

## **5 Interview Discussion Points**

When presenting these examples in interviews, discuss:

### **Trade-offs made:**

- Why BlockingQueue vs custom wait/notify?
- When to use ConcurrentHashMap vs synchronized Map?
- Lock granularity: coarse-grained vs fine-grained
- Memory overhead vs performance

### **Edge cases handled:**

- What happens during shutdown?
- How are errors propagated?
- What about thread interruption?
- Resource cleanup in finally blocks

### **Scalability considerations:**

- How does it perform under high contention?
- Where are the bottlenecks?
- How would you monitor/tune in production?

### **Alternative designs:**

- Could use different primitives?
- What about lock-free approaches?
- Trade-offs between approaches