

Java Concurrency: Complex Practical Examples

Integrating Multiple Concepts for Staff-Level Interviews

1 Example 1: Thread-Safe Cache with Expiration

Requirements:

- Multiple threads read frequently (optimize for reads)
- Occasional writes to update entries
- Entries expire after TTL (time-to-live)
- Background thread periodically removes expired entries
- Thread-safe statistics tracking (hits, misses)

Concepts integrated: ReadWriteLock, ScheduledExecutorService, AtomicLong, ConcurrentHashMap

```
1 import java.util.concurrent.*;
2 import java.util.concurrent.locks.*;
3 import java.util.concurrent.atomic.*;
4
5 class ExpiringCache<K, V> {
6     private static class CacheEntry<V> {
7         final V value;
8         final long expiryTime;
9
10        CacheEntry(V value, long ttlMillis) {
11            this.value = value;
12            this.expiryTime = System.currentTimeMillis() + ttlMillis;
13        }
14
15        boolean isExpired() {
16            return System.currentTimeMillis() > expiryTime;
17        }
18    }
19
20    private final ConcurrentHashMap<K, CacheEntry<V>> cache;
21    private final AtomicLong hits = new AtomicLong(0);
22    private final AtomicLong misses = new AtomicLong(0);
23    private final ScheduledExecutorService cleaner;
24    private final long ttlMillis;
25
26    public ExpiringCache(long ttlMillis, long cleanupIntervalMillis) {
27        this.cache = new ConcurrentHashMap<>();
28        this.ttlMillis = ttlMillis;
29        this.cleaner = Executors.newScheduledThreadPool(1);
30
31        // Schedule periodic cleanup
32        cleaner.scheduleAtFixedRate(
33            this::removeExpiredEntries,
34            cleanupIntervalMillis,
35            cleanupIntervalMillis,
36            TimeUnit.MILLISECONDS
37        );
38    }
39
40    public V get(K key) {
41        CacheEntry<V> entry = cache.get(key);
42
43        if (entry == null || entry.isExpired()) {
44            misses.incrementAndGet();
45            if (entry != null) {
46                cache.remove(key, entry); // Remove expired
47            }
48            return null;
49        }
50
51        hits.incrementAndGet();
52        return entry.value;
```

```

53 }
54
55     public void put(K key, V value) {
56         cache.put(key, new CacheEntry<>(value, ttlMillis));
57     }
58
59     private void removeExpiredEntries() {
60         cache.entrySet().removeIf(entry -> entry.getValue().isExpired());
61     }
62
63     public CacheStats getStats() {
64         long hitCount = hits.get();
65         long missCount = misses.get();
66         long total = hitCount + missCount;
67         double hitRate = total == 0 ? 0.0 : (double) hitCount / total;
68
69         return new CacheStats(hitCount, missCount, hitRate, cache.size());
70     }
71
72     public void shutdown() {
73         cleaner.shutdown();
74         try {
75             if (!cleaner.awaitTermination(5, TimeUnit.SECONDS)) {
76                 cleaner.shutdownNow();
77             }
78         } catch (InterruptedException e) {
79             cleaner.shutdownNow();
80             Thread.currentThread().interrupt();
81         }
82     }
83
84     static class CacheStats {
85         final long hits, misses;
86         final double hitRate;
87         final int size;
88
89         CacheStats(long hits, long misses, double hitRate, int size) {
90             this.hits = hits;
91             this.misses = misses;
92             this.hitRate = hitRate;
93             this.size = size;
94         }
95     }
96 }

```

Listing 1: Thread-Safe Cache Implementation

Key design decisions:

- **ConcurrentHashMap** for cache storage - allows concurrent reads/writes to different keys
- **AtomicLong** for hit/miss counters - lock-free increments under high contention
- **ScheduledExecutorService** for cleanup - better than manual thread + Timer
- **No ReadWriteLock on cache operations** - ConcurrentHashMap already optimized
- **Graceful shutdown** with timeout and force-stop fallback
- **Benign race in get():** Between checking isExpired() and removing, another thread might access the entry. This is acceptable - worst case is redundant removal. Alternative: `cache.computeIfPresent(key, (k,v) -> v.isExpired() ? null : v)` for atomic check-and-remove

2 Example 2: Rate-Limited API Client

Requirements:

- Limit to N requests per second across all threads
- Block threads when rate limit exceeded
- Support burst capacity (can briefly exceed rate)

- Track and report throttling statistics

Concepts integrated: Semaphore, ScheduledExecutorService, AtomicInteger, synchronized blocks

```

1 import java.util.concurrent.*;
2 import java.util.concurrent.atomic.*;
3
4 class RateLimitedApiClient {
5     private final Semaphore tokens;
6     private final int maxTokens;
7     private final int refillRate; // tokens per second
8     private final ScheduledExecutorService refiller;
9     private final AtomicInteger throttledRequests = new AtomicInteger(0);
10
11    public RateLimitedApiClient(int maxTokens, int refillRate) {
12        this.maxTokens = maxTokens;
13        this.refillRate = refillRate;
14        this.tokens = new Semaphore(maxTokens);
15        this.refiller = Executors.newScheduledThreadPool(1);
16
17        // Refill tokens every 100ms
18        long refillIntervalMs = 100;
19        int tokensPerInterval = Math.max(1, refillRate / 10);
20
21        refiller.scheduleAtFixedRate(
22            () -> refillTokens(tokensPerInterval),
23            refillIntervalMs,
24            refillIntervalMs,
25            TimeUnit.MILLISECONDS
26        );
27    }
28
29    public <T> T makeRequest(Callable<T> apiCall)
30        throws Exception {
31        // Try to acquire token (blocks if none available)
32        boolean acquired = tokens.tryAcquire(5, TimeUnit.SECONDS);
33
34        if (!acquired) {
35            throttledRequests.incrementAndGet();
36            throw new RateLimitException("Rate limit exceeded");
37        }
38
39        try {
40            return apiCall.call();
41        } finally {
42            // Token not returned - consumed by rate limit
43        }
44    }
45
46    private void refillTokens(int count) {
47        int available = tokens.availablePermits();
48        int toAdd = Math.min(count, maxTokens - available);
49
50        if (toAdd > 0) {
51            tokens.release(toAdd);
52        }
53    }
54
55    public RateLimitStats getStats() {
56        return new RateLimitStats(
57            tokens.availablePermits(),
58            throttledRequests.get()
59        );
60    }
61
62    public void shutdown() {
63        refiller.shutdown();
64    }
65
66    static class RateLimitStats {
67        final int availableTokens;

```

```

68     final int throttledRequests;
69
70     RateLimitStats(int availableTokens, int throttledRequests) {
71         this.availableTokens = availableTokens;
72         this.throttledRequests = throttledRequests;
73     }
74 }
75
76     static class RateLimitException extends Exception {
77         RateLimitException(String message) {
78             super(message);
79         }
80     }
81 }
```

Listing 2: Token Bucket Rate Limiter

Usage example:

```

1 RateLimitedApiClient client = new RateLimitedApiClient(
2     100, // max 100 tokens (burst capacity)
3     50 // refill 50 tokens/second
4 );
5
6 // Multiple threads making requests
7 ExecutorService executor = Executors.newFixedThreadPool(20);
8 for (int i = 0; i < 1000; i++) {
9     executor.submit(() -> {
10         try {
11             String result = client.makeRequest(() -> {
12                 return callExternalApi();
13             });
14             System.out.println("Success: " + result);
15         } catch (RateLimitException e) {
16             System.out.println("Throttled!");
17         }
18     });
19 }
```

Key design decisions:

- **Semaphore for token bucket** - naturally models available capacity
- **tryAcquire with timeout** - prevents indefinite blocking
- **Tokens not returned** after request - consumed by rate limit
- **Periodic refill** - scheduled task adds tokens back
- **Bounded refill** - never exceed maxTokens (prevents overflow)

3 Example 3: Parallel File Processor with Work Stealing

Requirements:

- Process large directory of files in parallel
- Some files larger than others (imbalanced workload)
- Aggregate results from all files
- Track progress and report completion
- Handle errors gracefully

Concepts integrated: ExecutorService, CountDownLatch, ConcurrentHashMap, AtomicInteger, Future

```

1 import java.util.*;
2 import java.util.concurrent.*;
3 import java.util.concurrent.atomic.*;
4 import java.io.*;
5 import java.nio.file.*;
6
7 class ParallelFileProcessor {
```

```

8     private final ExecutorService executor;
9     private final ConcurrentHashMap<String, ProcessingResult> results;
10    private final AtomicInteger processed = new AtomicInteger(0);
11    private final AtomicInteger failed = new AtomicInteger(0);
12
13    public ParallelFileProcessor(int threadCount) {
14        // Work-stealing pool for imbalanced workloads
15        this.executor = Executors.newWorkStealingPool(threadCount);
16        this.results = new ConcurrentHashMap<>();
17    }
18
19    public AggregateResult processDirectory(Path directory)
20        throws IOException, InterruptedException {
21        List<Path> files = Files.walk(directory)
22            .filter(Files::isRegularFile)
23            .toList();
24
25        int totalFiles = files.size();
26        CountDownLatch latch = new CountDownLatch(totalFiles);
27
28        // Submit all file processing tasks
29        List<Future<ProcessingResult>> futures = new ArrayList<>();
30        for (Path file : files) {
31            Future<ProcessingResult> future = executor.submit(() -> {
32                try {
33                    ProcessingResult result = processFile(file);
34                    results.put(file.toString(), result);
35                    processed.incrementAndGet();
36                    return result;
37                } catch (Exception e) {
38                    failed.incrementAndGet();
39                    System.err.println("Failed: " + file + " - " + e.getMessage());
40                    return ProcessingResult.error(file.toString(), e);
41                } finally {
42                    latch.countDown();
43                }
44            });
45            futures.add(future);
46        }
47
48        // Progress reporting in separate thread
49        Thread progressThread = new Thread(() -> {
50            while (latch.getCount() > 0) {
51                reportProgress(totalFiles);
52                try {
53                    Thread.sleep(1000);
54                } catch (InterruptedException e) {
55                    break;
56                }
57            }
58        });
59        progressThread.start();
60
61        // Wait for all files to complete
62        latch.await();
63        progressThread.interrupt();
64
65        // Aggregate results
66        return aggregateResults(files, futures);
67    }
68
69    private ProcessingResult processFile(Path file) throws IOException {
70        // Simulate file processing
71        long wordCount = Files.lines(file).count();
72        long byteSize = Files.size(file);
73        return new ProcessingResult(file.toString(), wordCount, byteSize, true, null);
74    }
75
76    private void reportProgress(int total) {
77        int done = processed.get();

```

```

78     int errors = failed.get();
79     double pct = (done + errors) * 100.0 / total;
80     System.out.printf("Progress: %.1f%% (%d/%d, %d errors)%n",
81                       pct, done, total, errors);
82   }
83
84   private AggregateResult aggregateResults(List<Path> files,
85                                             List<Future<ProcessingResult>> futures) {
86     long totalWords = 0;
87     long totalBytes = 0;
88     int successCount = 0;
89     List<String> errors = new ArrayList<>();
90
91     for (Future<ProcessingResult> future : futures) {
92       try {
93         ProcessingResult result = future.get();
94         if (result.success) {
95           totalWords += result.wordCount;
96           totalBytes += result.byteSize;
97           successCount++;
98         } else {
99           errors.add(result.filename + ": " + result.error.getMessage());
100        }
101      } catch (Exception e) {
102        errors.add("Future failed: " + e.getMessage());
103      }
104    }
105
106    return new AggregateResult(files.size(), successCount,
107                               totalWords, totalBytes, errors);
108  }
109
110  public void shutdown() {
111    executor.shutdown();
112    try {
113      if (!executor.awaitTermination(60, TimeUnit.SECONDS)) {
114        executor.shutdownNow();
115      }
116    } catch (InterruptedException e) {
117      executor.shutdownNow();
118    }
119  }
120
121  static class ProcessingResult {
122    final String filename;
123    final long wordCount;
124    final long byteSize;
125    final boolean success;
126    final Exception error;
127
128    ProcessingResult(String filename, long wordCount, long byteSize,
129                      boolean success, Exception error) {
130      this.filename = filename;
131      this.wordCount = wordCount;
132      this.byteSize = byteSize;
133      this.success = success;
134      this.error = error;
135    }
136
137    static ProcessingResult error(String filename, Exception error) {
138      return new ProcessingResult(filename, 0, 0, false, error);
139    }
140  }
141
142  static class AggregateResult {
143    final int totalFiles;
144    final int successCount;
145    final long totalWords;
146    final long totalBytes;
147    final List<String> errors;

```

```

148     AggregateResult(int totalFiles, int successCount, long totalWords,
149                     long totalBytes, List<String> errors) {
150         this.totalFiles = totalFiles;
151         this.successCount = successCount;
152         this.totalWords = totalWords;
153         this.totalBytes = totalBytes;
154         this.errors = errors;
155     }
156 }
157 }
158 }
```

Listing 3: Parallel File Processor

Key design decisions:

- **newWorkStealingPool()** - threads steal work from each other when idle (good for imbalanced tasks)
- **CountDownLatch** - wait for all files to complete before aggregating
- **ConcurrentHashMap** - store results from multiple threads
- **AtomicInteger** - track progress without locks
- **Future list** - collect results for final aggregation
- **Separate progress thread** - non-blocking progress reporting
- **Graceful error handling** - failures don't stop other tasks

4 Example 4: Connection Pool

Requirements:

- Fixed pool of N database connections
- Threads block when all connections in use
- Connections have idle timeout (close if unused)
- Health check connections periodically
- Track utilization statistics

Concepts integrated: BlockingQueue, ScheduledExecutorService, ReentrantLock, wait/notify

```

1 import java.util.concurrent.*;
2 import java.util.concurrent.locks.*;
3 import java.util.*;
4
5 class ConnectionPool {
6     private final BlockingQueue<PooledConnection> available;
7     private final Set<PooledConnection> inUse;
8     private final Lock useLock = new ReentrantLock();
9     private final int maxSize;
10    private final long idleTimeoutMs;
11    private final ScheduledExecutorService healthChecker;
12    private volatile boolean isShutdown = false;
13
14    public ConnectionPool(int maxSize, long idleTimeoutMs) {
15        this.maxSize = maxSize;
16        this.idleTimeoutMs = idleTimeoutMs;
17        this.available = new LinkedBlockingQueue<>();
18        this.inUse = new HashSet<>();
19        this.healthChecker = Executors.newScheduledThreadPool(1);
20
21        // Initialize pool
22        for (int i = 0; i < maxSize; i++) {
23            available.offer(new PooledConnection(i));
24        }
25
26        // Periodic health check and idle timeout
27        healthChecker.scheduleAtFixedRate(
28            this::maintainPool,
29            30, 30, TimeUnit.SECONDS
30        );
31    }
32
33    private void maintainPool() {
34        while (!available.isEmpty() && inUse.size() < maxSize) {
35            PooledConnection connection = available.poll();
36            inUse.add(connection);
37            useLock.lock();
38            try {
39                connection.start();
40            } catch (Exception e) {
41                inUse.remove(connection);
42                available.offer(connection);
43                throw e;
44            }
45            useLock.unlock();
46        }
47    }
48
49    void shutdown() {
50        isShutdown = true;
51        healthChecker.shutdown();
52    }
53
54    void close() {
55        shutdown();
56        while (!available.isEmpty()) {
57            PooledConnection connection = available.poll();
58            connection.close();
59        }
60    }
61
62    void release(PooledConnection connection) {
63        useLock.lock();
64        try {
65            if (connection.isRunning())
66                connection.stop();
67            inUse.remove(connection);
68            available.offer(connection);
69        } finally {
70            useLock.unlock();
71        }
72    }
73
74    void acquire(PooledConnection connection) {
75        useLock.lock();
76        try {
77            if (connection.isRunning())
78                connection.stop();
79            inUse.add(connection);
80        } finally {
81            useLock.unlock();
82        }
83    }
84
85    void start(PooledConnection connection) {
86        useLock.lock();
87        try {
88            if (!connection.isRunning())
89                connection.start();
90        } finally {
91            useLock.unlock();
92        }
93    }
94
95    void stop(PooledConnection connection) {
96        useLock.lock();
97        try {
98            if (connection.isRunning())
99                connection.stop();
100        } finally {
101            useLock.unlock();
102        }
103    }
104}
```

```

31 }
32
33     public PooledConnection acquire() throws InterruptedException {
34         if (isShutdown) {
35             throw new IllegalStateException("Pool is shutdown");
36         }
37
38         // Block until connection available
39         PooledConnection conn = available.take();
40
41         useLock.lock();
42         try {
43             conn.markInUse();
44             inUse.add(conn);
45         } finally {
46             useLock.unlock();
47         }
48
49         return conn;
50     }
51
52     public PooledConnection tryAcquire(long timeout, TimeUnit unit)
53         throws InterruptedException {
54         if (isShutdown) {
55             throw new IllegalStateException("Pool is shutdown");
56         }
57
58         PooledConnection conn = available.poll(timeout, unit);
59         if (conn == null) {
60             return null; // Timeout
61         }
62
63         useLock.lock();
64         try {
65             conn.markInUse();
66             inUse.add(conn);
67         } finally {
68             useLock.unlock();
69         }
70
71         return conn;
72     }
73
74     public void release(PooledConnection conn) {
75         if (conn == null) return;
76
77         useLock.lock();
78         try {
79             if (!inUse.remove(conn)) {
80                 throw new IllegalStateException("Connection not in use");
81             }
82             conn.markAvailable();
83         } finally {
84             useLock.unlock();
85         }
86
87         available.offer(conn);
88     }
89
90     private void maintainPool() {
91         useLock.lock();
92         try {
93             // Health check in-use connections
94             for (PooledConnection conn : inUse) {
95                 if (!conn.isHealthy()) {
96                     System.err.println("Unhealthy connection: " + conn.id);
97                     // In real implementation: close and recreate
98                 }
99             }
100        }

```

```

101     // Check idle timeout for available connections
102     long now = System.currentTimeMillis();
103     Iterator<PooledConnection> iter = available.iterator();
104     while (iter.hasNext()) {
105         PooledConnection conn = iter.next();
106         if (now - conn.lastUsedTime > idleTimeoutMs) {
107             iter.remove();
108             conn.close();
109             // Create new connection to maintain pool size
110             available.offer(new PooledConnection(conn.id));
111         }
112     }
113 } finally {
114     useLock.unlock();
115 }
116 }
117
118 public PoolStats getStats() {
119     useLock.lock();
120     try {
121         return new PoolStats(
122             available.size(),
123             inUse.size(),
124             maxSize
125         );
126     } finally {
127         useLock.unlock();
128     }
129 }
130
131 public void shutdown() {
132     isShutdown = true;
133     healthChecker.shutdown();
134
135     useLock.lock();
136     try {
137         for (PooledConnection conn : available) {
138             conn.close();
139         }
140         available.clear();
141     } finally {
142         useLock.unlock();
143     }
144 }
145
146 static class PooledConnection {
147     final int id;
148     volatile long lastUsedTime;
149     volatile boolean inUse;
150
151     PooledConnection(int id) {
152         this.id = id;
153         this.lastUsedTime = System.currentTimeMillis();
154         this.inUse = false;
155     }
156
157     void markInUse() {
158         this.inUse = true;
159         this.lastUsedTime = System.currentTimeMillis();
160     }
161
162     void markAvailable() {
163         this.inUse = false;
164         this.lastUsedTime = System.currentTimeMillis();
165     }
166
167     boolean isHealthy() {
168         // Simulate health check
169         return true;
170     }

```

```

171     void close() {
172         // Close actual connection
173     }
174
175     public void execute(String sql) {
176         // Execute query
177     }
178 }
179
180 static class PoolStats {
181     final int available;
182     final int inUse;
183     final int total;
184
185     PoolStats(int available, int inUse, int total) {
186         this.available = available;
187         this.inUse = inUse;
188         this.total = total;
189     }
190
191     public double getUtilization() {
192         return total == 0 ? 0.0 : (double) inUse / total;
193     }
194 }
195 }
196

```

Listing 4: Database Connection Pool

Usage pattern:

```

1 ConnectionPool pool = new ConnectionPool(10, 60000); // 10 conn, 60s idle
2
3 // Thread-safe acquire/release
4 PooledConnection conn = null;
5 try {
6     conn = pool.acquire(); // blocks if none available
7     conn.execute("SELECT * FROM users");
8 } finally {
9     if (conn != null) {
10         pool.release(conn); // Always return to pool
11     }
12 }
13
14 // With timeout
15 conn = pool.tryAcquire(5, TimeUnit.SECONDS);
16 if (conn != null) {
17     try {
18         conn.execute("SELECT * FROM orders");
19     } finally {
20         pool.release(conn);
21     }
22 } else {
23     System.out.println("Could not acquire connection");
24 }

```

Key design decisions:

- **BlockingQueue for available connections** - natural blocking behavior when pool exhausted
- **Set for in-use tracking** - fast lookup to validate release
- **ReentrantLock for state transitions** - protects inUse set modifications
- **ScheduledExecutorService** for maintenance - health checks and idle timeout
- **Try-acquire with timeout** - prevents indefinite blocking
- **Graceful shutdown** - close all connections, reject new acquires

5 Interview Discussion Points

When presenting these examples in interviews, discuss:

Trade-offs made:

- Why BlockingQueue vs custom wait/notify?
- When to use ConcurrentHashMap vs synchronized Map?
- Lock granularity: coarse-grained vs fine-grained
- Memory overhead vs performance

Edge cases handled:

- What happens during shutdown?
- How are errors propagated?
- What about thread interruption?
- Resource cleanup in finally blocks

Scalability considerations:

- How does it perform under high contention?
- Where are the bottlenecks?
- How would you monitor/tune in production?

Alternative designs:

- Could use different primitives?
- What about lock-free approaches?
- Trade-offs between approaches