

Augment Code Interview Preparation

Network Ports & Process Management

Alex Yang

Prepared: November 26, 2024

Contents

1 Interview Overview	3
1.1 Problem Type	3
1.2 Interview Structure	3
2 Part 1: Network Ports Problem	5
2.1 Problem Statement	5
2.1.1 Data Structures	5
2.2 Requirements	5
2.3 Solution Approach	6
2.4 Implementation - Version 1 (Recursive DFS)	6
2.5 Complexity Analysis	7
2.6 Optimization Discussion	8
2.7 Implementation - Version 2 (Optimized)	8
2.8 Edge Cases	10
3 Part 2: Garbage Collection	11
3.1 Problem Statement	11
3.2 Array Analogy	11
3.3 Implementation	11
3.4 Visualization	12
3.5 Alternative: Stable Partition	13
4 Part 3: Thread Safety & Concurrency	15
4.1 Problem Context	15
4.2 Concurrency Concepts	15
4.3 Design Discussion Points	16
4.4 Common Mistakes	17
4.5 Advanced Considerations	17
5 Complete Solution Code	19
6 Interview Execution Strategy	24
6.1 Time Management (50 minutes)	24
6.2 Communication Tips	25
6.3 What to Write on Whiteboard	25

7 Key Concepts Review	27
7.1 Process Hierarchy	27
7.2 Data Structures	28
7.3 Recursion Patterns	29
7.4 Concurrency Primitives	30
7.5 Networking Basics	31
8 Common Mistakes & How to Avoid	32
9 Final Preparation Checklist	34
9.1 Day Before Interview	34
9.2 Interview Day Checklist	35
9.3 Key Reminders	36
10 Summary	37
10.1 Core Algorithms	37
10.2 Interview Mindset	37

1 Interview Overview

1.1 Problem Type

- **Category:** Systems Programming / Process Management
- **Difficulty:** Medium-Hard
- **Duration:** 50 minutes
- **Company:** Augment Code (AI Coding Startup)

1.2 Interview Structure

1. **Part 1 (20-25 min):** Implement network port finding function

- Uses recursion, hash set, array iteration
- Process hierarchy traversal
- Performance optimization discussion

2. **Part 2 (15-20 min):** Implement garbage collection logic

- Similar to moving 1s in binary array
- In-place swap approach
- Array manipulation efficiency

3. **Part 3 (5-10 min):** Thread safety discussion

- Read-Write (RW) locks
- Concurrency techniques
- Design discussion (no coding)

⚠ Critical Point**Key Interview Signals Based on Past Feedback:**
What they look for:

- Quick recognition of design issues (recursion, flat-list pitfalls)
- Understanding of space vs. time complexity tradeoffs
- Clean, correct implementations (not sloppy)
- Strong communication about design decisions
- Knowledge of concurrency primitives

Common pitfalls from past candidates:

- Missing that `get_children()` returns PIDs, not full records
- Spending too long trying to avoid copying the database
- Not knowing `del` from list is $O(n)$ operation
- Sloppy implementations with edge case bugs
- Weak concurrency knowledge (thinking single lock suffices)

2 Part 1: Network Ports Problem

2.1 Problem Statement

Context: You have a database of processes where each process can have child processes forming a tree hierarchy. Some processes are listening on network ports. You need to find all ports used by a given process and all its descendants.

2.1.1 Data Structures

Process Record:

```

1 class ProcessRecord:
2     def __init__(self, pid, parent_pid, ports, is_alive):
3         self.pid = pid                      # Process ID (unique)
4         self.parent_pid = parent_pid        # Parent PID (None for root)
5         self.ports = ports                  # List of port numbers [80, 443]
6         self.is_alive = is_alive           # Boolean: True/False

```

Database Interface:

```

1 class ProcessDatabase:
2     def __init__(self, processes: List[ProcessRecord]):
3         """Store processes in flat list (database table)"""
4         self.processes = processes
5
6     def get_process(self, pid: int) -> ProcessRecord:
7         """Get process by PID - O(n) lookup"""
8         for proc in self.processes:
9             if proc.pid == pid:
10                 return proc
11         return None
12
13     def get_children(self, pid: int) -> List[int]:
14         """Get child PIDs (not full records!) - O(n) scan"""
15         children_pids = []
16         for proc in self.processes:
17             if proc.parent_pid == pid:
18                 children_pids.append(proc.pid)
19
20         return children_pids

```

2.2 Requirements

Implement:

```

1 def find_all_ports(database: ProcessDatabase, root_pid: int) -> List[int]:
2     """
3         Find all network ports used by root_pid and all its descendants.
4
5     Args:
6         database: ProcessDatabase instance
7         root_pid: Starting process ID
8
9     Returns:
10        List of unique port numbers (no duplicates)
11
12    Notes:
13        - Only include ALIVE processes

```

```

14     - Traverse entire subtree recursively
15     - Handle dead processes gracefully
16 """
17 pass

```

Example:

Process Tree:

```

1 (ports=[80], alive=True)
+- 2 (ports=[443], alive=True)
|   +- 4 (ports=[8080], alive=False) # Dead - ignore
+- 3 (ports=[3000, 3001], alive=True)
    +- 5 (ports=[5000], alive=True)

find_all_ports(db, 1) -> [80, 443, 3000, 3001, 5000]
find_all_ports(db, 2) -> [443] # Ignores dead child 4
find_all_ports(db, 3) -> [3000, 3001, 5000]

```

2.3 Solution Approach

💡 Key Insight**Key Design Decisions:**

1. **Recursion:** Natural fit for tree traversal
2. **Hash set:** Track unique ports (automatic deduplication)
3. **Filter dead processes:** Check `is_alive` before processing
4. **Understand API:** `get_children()` returns PIDs, need `get_process()` to get records

2.4 Implementation - Version 1 (Recursive DFS)

```

1 def find_all_ports(database: ProcessDatabase, root_pid: int) -> List[int]:
2 """
3     Find all ports using recursive DFS with hash set for deduplication.
4
5     Time: O(n) where n = number of alive processes in subtree
6     Space: O(h + p) where h = height, p = unique ports
7 """
8
9     # Use set for automatic deduplication
10    ports_set = set()
11
12    def dfs(pid: int) -> None:
13        """Recursive helper to traverse process tree"""
14        # Get process record
15        process = database.get_process(pid)
16
17        # Handle missing or dead process
18        if process is None or not process.is_alive:
19            return
20
21        # Add this process's ports

```

```

21     for port in process.ports:
22         ports_set.add(port)
23
24     # Recursively process children
25     # CRITICAL: get_children returns PIDs, not ProcessRecords!
26     child_pids = database.get_children(pid)
27     for child_pid in child_pids:
28         dfs(child_pid)
29
30     # Start DFS from root
31     dfs(root_pid)
32
33     # Convert set to list
34     return list(ports_set)

```

⚠ Critical Point

Common Mistake - Past Candidate Error:

Many candidates miss that `get_children(pid)` returns **PIDs** (integers), not `ProcessRecord` objects!

Wrong:

```

1 children = database.get_children(pid)
2 for child in children:
3     for port in child.ports:  # ERROR: child is int, not ProcessRecord!
4         ports_set.add(port)

```

Correct:

```

1 child_pids = database.get_children(pid)
2 for child_pid in child_pids:
3     child_process = database.get_process(child_pid)
4     if child_process and child_process.is_alive:
5         # Now we have the actual ProcessRecord

```

2.5 Complexity Analysis

Time Complexity: $O(n \times m)$ where:

- n = number of alive processes in subtree
- m = average number of children per process
- Each `get_children()` scans entire database ($O(m)$)
- Total: n calls $\times O(m)$ per call = $O(n \times m)$

Space Complexity: $O(h + p)$ where:

- h = tree height (recursion stack)
- p = number of unique ports

2.6 Optimization Discussion

💡 Key Insight

Performance Bottleneck:

The `get_children()` method scans the entire database on every call. For large databases, this becomes expensive.

Interviewer might ask: "Can you optimize this?"

Key insight from feedback: One candidate proposed "filtering out dead processes" early but communicated it oddly. The interviewer wants to hear:

Two-pass optimization:

1. **Pass 1:** Build parent→children map once ($O(n)$)
2. **Pass 2:** DFS using map for instant child lookup ($O(1)$)

This reduces time from $O(n \times m)$ to $O(n)$.

2.7 Implementation - Version 2 (Optimized)

```

1 def find_all_ports_optimized(database: ProcessDatabase, root_pid: int) -> List[int]:
2     """
3         Optimized version: Build parent->children map first.
4
5         Time: O(n) total - single pass to build map + DFS
6         Space: O(n + p) - map storage + ports set
7     """
8
9     # Pass 1: Build parent->children mapping (O(n))
10    parent_to_children = {}
11    pid_to_process = {}
12
13    for process in database.processes:
14        pid_to_process[process.pid] = process
15        if process.parent_pid is not None:
16            if process.parent_pid not in parent_to_children:
17                parent_to_children[process.parent_pid] = []
18                parent_to_children[process.parent_pid].append(process.pid)
19
20    # Pass 2: DFS using precomputed map
21    ports_set = set()
22
23    def dfs(pid: int) -> None:
24        # Get process from our map
25        process = pid_to_process.get(pid)
26
27        if process is None or not process.is_alive:
28            return
29
30        # Add ports
31        ports_set.update(process.ports) # More pythonic than loop
32
33        # Get children from map (O(1) lookup)
34        child_pids = parent_to_children.get(pid, [])
        for child_pid in child_pids:

```

```
35         dfs(child_pid)
36
37     dfs(root_pid)
38     return list(ports_set)
39
40
41 def find_all_ports_space_optimized(database: ProcessDatabase, root_pid: int) ->
42     List[int]:
43     """
44     Space-Optimized Approach: "Traverse + Scan"
45
46     This addresses the interviewer's specific question about Space Complexity.
47     Instead of building a full O(N) map of the database, we only store PIDs
48     for the relevant subtree.
49
50     Strategy:
51     1. Collect target PIDs (recursively) -> Set of PIDs
52     2. Scan database once -> Filter against Set
53
54     Time: O(N + M) where N=database size, M=subtree size
55     Space: O(M) - Only store PIDs in the subtree (vs O(N) for full map)
56     """
57
58     # Pass 1: Collect all PIDs in subtree using get_children
59     subtree_pids = set()
60
61     def collect_pids(pid: int) -> None:
62         """Recursively collect all PIDs in subtree"""
63         subtree_pids.add(pid)
64         child_pids = database.get_children(pid)
65         for child_pid in child_pids:
66             collect_pids(child_pid)
67
68     collect_pids(root_pid)
69
70     # Pass 2: Single scan through database, lookup against PIDs
71     ports_set = set()
72     for process in database.processes:
73         # Check if this process is in our subtree
74         if process.pid in subtree_pids and process.is_alive:
75             ports_set.update(process.ports)
76
77     return list(ports_set)
```

✔ Action Item

What to communicate in interview:

"The naive approach calls `get_children()` repeatedly, scanning the entire database each time. I see two optimization approaches."

Approach 1 - Build parent→children map:

- Preprocess database to build lookup structure
- DFS with $O(1)$ child lookups
- Time: $O(m + n)$ where m =database size, n =subtree size

Approach 2 - Collect PIDs then scan (Space Optimization):

- First pass: DFS to collect only relevant PIDs in subtree
- Second pass: Single database scan, filter by PID set
- **Why this matters:** Reduces space from $O(\text{Database})$ to $O(\text{Subtree})$.
- *This matches the "aha" hint: "iterate database and lookup against pids"*

Space vs Time Tradeoff (Critical Discussion):

- **Naive DFS:** Time $O(N \times \text{Children})$, Space $O(\text{Height})$
- **Map Optimization:** Time $O(N)$, Space $O(N)$ (Index entire DB)
- **Traverse + Scan:** Time $O(N)$, Space $O(\text{Subtree})$ (Index only relevant items)
- *Candidate 1 missed this distinction when asked about space!*

2.8 Edge Cases

1. **Root process is dead:** Return empty list
2. **Root has no children:** Return only root's ports
3. **All descendants dead:** Return only root's ports
4. **Duplicate ports in tree:** Set handles deduplication
5. **Process with no ports:** Empty `ports` list is valid
6. **Invalid root PID:** Handle gracefully (return empty)

3 Part 2: Garbage Collection

3.1 Problem Statement

After finding ports, we need to compact the process database by removing dead processes. The interviewer describes this as "similar to moving all 1s in a binary array to the beginning."

Task: Implement garbage collection that:

- Moves all alive processes to the front of the array
- Moves all dead processes to the back
- Maintains relative order of alive processes
- Operates in-place (minimal extra space)

3.2 Array Analogy

Concept Review

Binary Array Problem (Classic Pattern):

Input: [0, 1, 0, 1, 1, 0, 1, 0]

Output: [1, 1, 1, 1, 0, 0, 0, 0]

Approach: Two-pointer swap

- Left pointer: Next position for a 1
- Right pointer: Current element being examined
- Swap when we find a 1

For processes:

- 1 = alive process
- 0 = dead process
- Goal: Move all alive to front

3.3 Implementation

```

1 def garbage_collect(processes: List[ProcessRecord]) -> int:
2     """
3         Compact process list: move alive processes to front.
4
5     Returns:
6         Number of alive processes (boundary index)
7
8     Time: O(n) - single pass
9     Space: O(1) - in-place swaps
10    """
11    # Two-pointer approach
12    write_idx = 0 # Next position for alive process
13
14    # Scan through array

```

```

15     for read_idx in range(len(processes)):
16         if processes[read_idx].is_alive:
17             # Found alive process - swap to front
18             if read_idx != write_idx:
19                 processes[write_idx], processes[read_idx] = \
20                     processes[read_idx], processes[write_idx]
21             write_idx += 1
22
23     # write_idx now points to first dead process
24     # Processes[0:write_idx] are alive
25     # Processes[write_idx:] are dead
26     return write_idx
27
28
29 def garbage_collect_with_truncation(processes: List[ProcessRecord]) -> None:
30     """
31     Alternative: Truncate dead processes entirely.
32     Modifies list in-place by removing dead entries.
33     """
34     alive_count = garbage_collect(processes)
35     # Remove dead processes from end
36     del processes[alive_count:]

```

⚠ Critical Point

Performance Trap - Past Candidate Error:

One candidate "didn't know that `del` from a list was linear time and didn't effectively internalize that after explanation."

Why this matters:

```

1 # WRONG: O(n^2) approach
2 i = 0
3 while i < len(processes):
4     if not processes[i].is_alive:
5         del processes[i]  # O(n) operation!
6         # Don't increment i
7     else:
8         i += 1

```

Problem:

- Each `del` operation shifts all subsequent elements
- If half the list is dead: $n/2$ deletions $\times O(n)$ each = $O(n^2)$

Correct approach:

- Use two-pointer swap (shown above) - $O(n)$
- Single `del` at end to truncate - $O(k)$ where k = dead count
- Total: $O(n)$

3.4 Visualization

Initial: [A1, D1, A2, D2, A3, D3] (A=alive, D=dead)

```
w=0  r=0

Step 1: r=0, A1 is alive
[A1, D1, A2, D2, A3, D3]
w$\rightarrow$r
Swap A1 with itself, w++, r++

Step 2: r=1, D1 is dead
[A1, D1, A2, D2, A3, D3]
w   r
Skip, r++

Step 3: r=2, A2 is alive
[A1, D1, A2, D2, A3, D3]
w           r
Swap D1 $\leftrightarrow$ A2
[A1, A2, D1, D2, A3, D3]
w
w++, r++

Step 4: r=3, D2 is dead
Skip, r++

Step 5: r=4, A3 is alive
[A1, A2, D1, D2, A3, D3]
w           r
Swap D1 $\leftrightarrow$ A3
[A1, A2, A3, D2, D1, D3]
w
w++, r++

Final: [A1, A2, A3, D2, D1, D3]
^boundary (w=3)
Alive: [0:3], Dead: [3:6]
```

3.5 Alternative: Stable Partition

```
1 def garbage_collect_stable(processes: List[ProcessRecord]) -> int:
2     """
3         Alternative using Python's stable partition.
4         Maintains relative order more explicitly.
5     """
6     alive = [p for p in processes if p.is_alive]
7     dead = [p for p in processes if not p.is_alive]
8
9     # Rebuild list
10    processes.clear()
11    processes.extend(alive)
12    processes.extend(dead)
13
```

14 `return len(alive)`

💡 Key Insight

Trade-offs:

Two-pointer swap:

- Pros: True in-place ($O(1)$ extra space), fast
- Cons: Slightly more complex logic

List comprehension rebuild:

- Pros: Cleaner, more Pythonic, easier to understand
- Cons: $O(n)$ extra space for temporary lists

In interview: Mention both, explain trade-offs. Candidate who proposed "swap approach" in feedback was on the right track but implementation was "sloppy" with edge cases.

4 Part 3: Thread Safety & Concurrency

4.1 Problem Context

The interviewer asks: "How would you make this system thread-safe if multiple threads are reading/writing the process database simultaneously?"

No coding required - design discussion only.

4.2 Concurrency Concepts

❑ Concept Review

Read-Write Lock (RWLock):

A synchronization primitive that allows:

- **Multiple readers** simultaneously (read-shared)
- **Single writer** exclusively (write-exclusive)
- Writers block all readers and other writers

When to use:

- Read-heavy workloads (common for process databases)
- Reads far outnumber writes
- Want to maximize read concurrency

Python implementation:

```
1 import threading
2
3 class ProcessDatabase:
4     def __init__(self):
5         self.processes = []
6         self.lock = threading.RLock() # or RWLock from external lib
7
8     def get_process(self, pid):
9         with self.lock: # Acquire for read
10             # Search logic
11             pass
12
13     def update_process(self, process):
14         with self.lock: # Acquire for write
15             # Update logic
16             pass
```

4.3 Design Discussion Points

✓ Action Item

Key points to articulate:

1. Identify Read vs Write Operations:

- **Reads:** `find_all_ports()`, `get_process()`, `get_children()`
- **Writes:** `garbage_collect()`, adding/updating processes

2. RW Lock Strategy:

```

1 class ThreadSafeProcessDatabase:
2     def __init__(self):
3         self.processes = []
4         self.rw_lock = RWLock()
5
6     def find_all_ports(self, root_pid):
7         """Read operation - shared lock"""
8         with self.rw_lock.read():
9             # Multiple threads can execute concurrently
10            return self._find_all_ports_impl(root_pid)
11
12    def garbage_collect(self):
13        """Write operation - exclusive lock"""
14        with self.rw_lock.write():
15            # Only one thread at a time
16            # Blocks all readers
17            return self._garbage_collect_impl()

```

3. Granularity Considerations:

- **Coarse-grained:** Lock entire database
 - Simple to reason about
 - May limit concurrency
- **Fine-grained:** Lock individual processes
 - Better concurrency
 - Risk of deadlocks
 - More complex

4. Transaction Safety:

```

1 def safe_update_and_gc(self):
2     """Atomic operation: update + GC"""
3     with self.rw_lock.write():
4         # Both operations happen atomically
5         self.update_processes()
6         self.garbage_collect()
7         # No other thread can read inconsistent state

```

4.4 Common Mistakes

⚠ Critical Point

Past candidate error: "Seemed to think he could get the benefit of locking by having a `with thread.Lock() on gc()` only."

Why this is wrong:

```

1 # INSUFFICIENT: Only locks GC, not reads
2 class BadDatabase:
3     def __init__(self):
4         self.processes = []
5         self.lock = threading.Lock()
6
7     def find_all_ports(self, root_pid):
8         # NO LOCK - race condition!
9         for proc in self.processes: # Another thread might be GC'ing!
10            ...
11
12     def garbage_collect(self):
13         with self.lock: # Lock only here - not enough!
14             # Swapping elements while readers access them = crash
15             ...

```

Problem:

- `find_all_ports()` can run while `garbage_collect()` modifies array
- Reader might access invalid indices during swap
- Data races, undefined behavior, crashes

Correct approach:

- ALL reads must acquire read lock
- ALL writes must acquire write lock
- No exceptions

4.5 Advanced Considerations

1. Copy-on-Write (COW):

- Readers use immutable snapshot
- Writers create new version
- Atomic pointer swap
- Good for read-heavy workloads

2. Lock-Free Data Structures:

- Use atomic operations (CAS)
- Avoid locks entirely
- Complex to implement correctly

- Better for high-contention scenarios

3. Thread-Local Caching:

- Each thread caches recent lookups
- Reduce lock contention
- Trade-off: Stale data risk

💡 Key Insight

For this interview:

Start with RW lock explanation. Only mention COW or lock-free if you have time and interviewer seems interested. Keep it practical and focused on correctness first, optimization second.

Good answer structure:

1. Identify read vs write operations
2. Explain why simple lock isn't optimal (blocks readers)
3. Propose RW lock with clear acquire points
4. Discuss granularity trade-offs
5. Mention testing strategy (race detectors, stress tests)

5 Complete Solution Code

```

1 """
2 Augment Code Interview - Network Ports Problem
3 Complete solution with all three parts
4 """
5
6 import threading
7 from typing import List, Set, Dict
8 from dataclasses import dataclass
9
10
11 @dataclass
12 class ProcessRecord:
13     """Process information"""
14     pid: int
15     parent_pid: int # None for root
16     ports: List[int]
17     is_alive: bool
18
19
20 class ProcessDatabase:
21     """Process database interface"""
22
23     def __init__(self, processes: List[ProcessRecord]):
24         self.processes = processes
25
26     def get_process(self, pid: int) -> ProcessRecord:
27         """Get process by PID - O(n) scan"""
28         for proc in self.processes:
29             if proc.pid == pid:
30                 return proc
31         return None
32
33     def get_children(self, pid: int) -> List[int]:
34         """Get child PIDs - O(n) scan"""
35         return [proc.pid for proc in self.processes
36                 if proc.parent_pid == pid]
37
38
39 class ThreadSafeProcessDatabase(ProcessDatabase):
40     """Thread-safe version with RW lock"""
41
42     def __init__(self, processes: List[ProcessRecord]):
43         super().__init__(processes)
44         # In production, use threading.RLock() or external RLock library
45         self.lock = threading.RLock()
46
47     def find_all_ports_safe(self, root_pid: int) -> List[int]:
48         """Thread-safe port finding"""
49         with self.lock: # Acquire for read
50             return find_all_ports_optimized(self, root_pid)
51
52     def garbage_collect_safe(self) -> int:
53         """Thread-safe garbage collection"""
54         with self.lock: # Acquire for write
55             return garbage_collect(self.processes)
56

```

```

57
58 # =====
59 # Part 1: Find All Ports
60 # =====
61
62 def find_all_ports(database: ProcessDatabase, root_pid: int) -> List[int]:
63     """
64     Basic recursive solution.
65
66     Time: O(n * m) where n=processes, m=avg children
67     Space: O(h + p) where h=height, p=unique ports
68     """
69     ports_set = set()
70
71     def dfs(pid: int) -> None:
72         process = database.get_process(pid)
73
74         if process is None or not process.is_alive:
75             return
76
77         # Add this process's ports
78         ports_set.update(process.ports)
79
80         # Traverse children (remember: get_children returns PIDs!)
81         child_pids = database.get_children(pid)
82         for child_pid in child_pids:
83             dfs(child_pid)
84
85     dfs(root_pid)
86     return list(ports_set)
87
88
89 def find_all_ports_optimized(database: ProcessDatabase,
90                             root_pid: int) -> List[int]:
91     """
92     Optimized with preprocessing.
93
94     Time: O(n) - single pass to build map + DFS
95     Space: O(n + p) - map storage + ports set
96     """
97
98     # Preprocessing: Build maps
99     parent_to_children: Dict[int, List[int]] = {}
100    pid_to_process: Dict[int, ProcessRecord] = {}
101
102    for process in database.processes:
103        pid_to_process[process.pid] = process
104        if process.parent_pid is not None:
105            parent_to_children.setdefault(
106                process.parent_pid, []
107            ).append(process.pid)
108
109    # DFS with O(1) lookups
110    ports_set = set()
111
112    def dfs(pid: int) -> None:
113        process = pid_to_process.get(pid)
114
115        if process is None or not process.is_alive:

```

```
116     ports_set.update(process.ports)
117
118     # O(1) lookup from map
119     for child_pid in parent_to_children.get(pid, []):
120         dfs(child_pid)
121
122     dfs(root_pid)
123     return list(ports_set)
124
125
126
127 # =====
128 # Part 2: Garbage Collection
129 # =====
130
131 def garbage_collect(processes: List[ProcessRecord]) -> int:
132     """
133     Move alive processes to front using two-pointer swap.
134
135     Time: O(n) - single pass
136     Space: O(1) - in-place swaps
137
138     Returns:
139         Index of first dead process (= count of alive)
140     """
141     write_idx = 0
142
143     for read_idx in range(len(processes)):
144         if processes[read_idx].is_alive:
145             # Swap alive process to front
146             if read_idx != write_idx:
147                 processes[write_idx], processes[read_idx] = \
148                     processes[read_idx], processes[write_idx]
149             write_idx += 1
150
151     return write_idx
152
153
154 def garbage_collect_with_removal(processes: List[ProcessRecord]) -> None:
155     """
156     Variant: Actually remove dead processes from list.
157     """
158     alive_count = garbage_collect(processes)
159     # Single O(k) deletion at end, not O(n^2) loop
160     del processes[alive_count:]
161
162
163 # =====
164 # Testing
165 # =====
166
167 def test_basic_tree():
168     """Test basic process tree traversal"""
169     processes = [
170         ProcessRecord(1, None, [80], True),
171         ProcessRecord(2, 1, [443], True),
172         ProcessRecord(3, 1, [3000, 3001], True),
173         ProcessRecord(4, 2, [8080], False), # Dead
174         ProcessRecord(5, 3, [5000], True),
```

```
175     ]
176
177     db = ProcessDatabase(processes)
178
179     # Test from root
180     ports = find_all_ports(db, 1)
181     assert set(ports) == {80, 443, 3000, 3001, 5000}
182
183     # Test from subtree
184     ports = find_all_ports(db, 2)
185     assert set(ports) == {443} # Ignores dead child
186
187     # Test optimized version
188     ports = find_all_ports_optimized(db, 1)
189     assert set(ports) == {80, 443, 3000, 3001, 5000}
190
191     print("[PASS] Basic tree test passed")
192
193
194 def test_garbage_collection():
195     """Test GC compaction"""
196     processes = [
197         ProcessRecord(1, None, [80], True),
198         ProcessRecord(2, 1, [], False),
199         ProcessRecord(3, 1, [443], True),
200         ProcessRecord(4, 1, [], False),
201         ProcessRecord(5, 3, [8080], True),
202     ]
203
204     alive_count = garbage_collect(processes)
205
206     assert alive_count == 3
207     # First 3 should be alive
208     for i in range(3):
209         assert processes[i].is_alive
210     # Rest should be dead
211     for i in range(3, 5):
212         assert not processes[i].is_alive
213
214     print("[PASS] Garbage collection test passed")
215
216
217 def test_edge_cases():
218     """Test edge cases"""
219     # Dead root
220     processes = [ProcessRecord(1, None, [80], False)]
221     db = ProcessDatabase(processes)
222     ports = find_all_ports(db, 1)
223     assert ports == []
224
225     # Empty ports
226     processes = [ProcessRecord(1, None, [], True)]
227     db = ProcessDatabase(processes)
228     ports = find_all_ports(db, 1)
229     assert ports == []
230
231     # Single process
232     processes = [ProcessRecord(1, None, [80, 443], True)]
233     db = ProcessDatabase(processes)
```

```
234     ports = find_all_ports(db, 1)
235     assert set(ports) == {80, 443}
236
237     print("[PASS] Edge cases test passed")
238
239
240 if __name__ == "__main__":
241     test_basic_tree()
242     test_garbage_collection()
243     test_edge_cases()
244     print("\n[SUCCESS] All tests passed!")
```

6 Interview Execution Strategy

6.1 Time Management (50 minutes)

✓ Action Item

Recommended breakdown:

1. **Minutes 0-5:** Understand problem, ask clarifications
 - What does `get_children()` return? (PIDs!)
 - Should we filter dead processes? (Yes)
 - Can ports be duplicated? (Yes - use set)
2. **Minutes 5-10:** Explain approach
 - "I'll use recursive DFS with a hash set"
 - "Key insight: `get_children()` returns PIDs, need `get_process()`"
 - Draw small tree example
3. **Minutes 10-25:** Implement Part 1
 - Write basic recursive solution first
 - Test with simple example
 - Discuss optimization if interviewer asks
4. **Minutes 25-30:** Complexity analysis
 - Time: $O(n \times m)$ basic, $O(n)$ optimized
 - Space: $O(h)$ stack, $O(p)$ ports set
 - Explain trade-offs
5. **Minutes 30-40:** Implement Part 2 (GC)
 - Explain two-pointer approach
 - "Like partitioning 1s and 0s in array"
 - Mention `del` is $O(n)$ pitfall
6. **Minutes 40-50:** Part 3 discussion
 - Identify reads vs writes
 - Explain RW lock benefit
 - Discuss where to acquire locks
 - Mention testing strategy

6.2 Communication Tips

💡 Key Insight

Based on feedback - What interviewers want to hear:

Design recognition (early in interview):

- "I notice this is a tree traversal - recursion makes sense"
- "The flat list means lookups are $O(n)$ - we could optimize with a map"
- "Need to be careful: `get_children()` returns IDs, not objects"

During implementation:

- "I'm using a set here for automatic deduplication"
- "Need to check `is_alive` before processing"
- "This recurses down the tree, collecting ports at each level"

Performance discussion:

- "Current approach calls `get_children()` n times, each scanning the database"
- "Trade-off: We could precompute a map for $O(1)$ lookups using $O(n)$ space"
- "For production, the optimization is worth it"

Don't say (from feedback):

- Vague statements about "filtering" without clear explanation
- "I think this works" without walking through logic
- Proposing approach you don't understand

6.3 What to Write on Whiteboard

1. Small tree example:

```

    1 [80]
    / \
  2   3 [443]
  /     \
4(D)   5 [8080]
  
```

Output: {80, 443, 8080}

2. Two-pointer GC visualization:

`[A, D, A, D, A] $rightarrow$ [A, A, A, D, D]`

w r w=3

3. Lock hierarchy:

Read operations	Write operations
-----	-----
find_all_ports	\$\rightarrow\$ garbage_collect
get_process	\$\rightarrow\$ update_process
(multiple concurrent)	(exclusive)

7 Key Concepts Review

7.1 Process Hierarchy

❑ Concept Review

Process Tree Concepts:

- **PID (Process ID):** Unique identifier for each process
- **Parent PID (PPID):** ID of parent process
- **Root process:** Has `parent.pid = None` (e.g., init/systemd)
- **Child processes:** Created via fork/spawn
- **Orphan processes:** Parent died, re-parented to init
- **Zombie processes:** Dead but not yet reaped

Real-world examples:

- Web server (nginx) spawns worker processes
- Each worker inherits parent's listening ports
- `ps auxf` shows tree hierarchy
- `pstree` visualizes process relationships

7.2 Data Structures

Concept Review

Hash Set Properties:

- **Insertion:** $O(1)$ average
- **Lookup:** $O(1)$ average
- **Automatic deduplication:** Adding same element is no-op
- **No ordering:** Iteration order undefined
- **Python:** `set()` built-in type

When to use:

- Need unique elements
- Fast membership testing
- Collecting distinct values

Array vs Hash Set:

- Array: Order preserved, duplicates allowed, $O(n)$ search
- Set: No order, no duplicates, $O(1)$ search

7.3 Recursion Patterns

Concept Review

Tree Recursion Template:

```
1 def traverse_tree(node):
2     # Base case
3     if node is None:
4         return
5
6     # Process current node
7     process(node)
8
9     # Recursive case: visit children
10    for child in node.children:
11        traverse_tree(child)
```

Key considerations:

- **Base case:** When to stop recursing (null node, leaf)
- **Stack depth:** Recursion uses call stack - $O(h)$ space
- **Tail recursion:** Can be optimized by compiler (not Python)
- **Alternative:** Iterative with explicit stack

7.4 Concurrency Primitives

Concept Review

Lock Types:

1. Mutex (Mutual Exclusion):

- One thread at a time (read or write)
- Simple but restrictive
- Python: `threading.Lock()`

2. RWLock (Read-Write Lock):

- Multiple readers OR single writer
- Better for read-heavy workloads
- Python: Need external library (e.g., `readerwriterlock`)

3. Semaphore:

- N threads at a time
- Rate limiting, connection pools
- Python: `threading.Semaphore(n)`

4. Condition Variable:

- Wait for specific condition
- Producer-consumer patterns
- Python: `threading.Condition()`

Common patterns:

```
1 # Context manager (automatic release)
2 with lock:
3     # Critical section
4     access_shared_data()
5
6 # Manual (error-prone)
7 lock.acquire()
8 try:
9     access_shared_data()
10 finally:
11     lock.release()
```

7.5 Networking Basics

❑ Concept Review

Network Ports:

- **Port number:** 16-bit integer (0-65535)
- **Well-known:** 0-1023 (HTTP=80, HTTPS=443, SSH=22)
- **Registered:** 1024-49151 (Application-specific)
- **Dynamic:** 49152-65535 (Ephemeral)

Process port binding:

- Process calls `bind(port)` to listen
- Only one process per port (usually)
- `SO_REUSEADDR` allows exceptions
- Child processes can inherit listening sockets

Checking ports:

- Linux: `lsof -i :port` or `netstat -tulpn`
- macOS: `lsof -i :port`
- Windows: `netstat -ano`

8 Common Mistakes & How to Avoid

⚠ Critical Point

Mistake 1: Misunderstanding get_children() return type

Error:

```
1 children = database.get_children(pid)
2 for child in children:
3     ports.extend(child.ports) # child is int, not ProcessRecord!
```

Fix:

```
1 child_pids = database.get_children(pid)
2 for child_pid in child_pids:
3     child = database.get_process(child_pid)
4     if child and child.is_alive:
5         ports.extend(child.ports)
```

Prevention: Read method signatures carefully!

⚠ Critical Point

Mistake 2: Premature optimization

Error: Spending 10+ minutes trying to avoid copying database

Interviewer feedback: "Candidate spent 10min spinning wheels trying linear-time approach avoiding copying"

Better approach:

1. Write correct basic solution first (5 min)
2. Explain optimization opportunity (2 min)
3. Implement optimization if interviewer interested (5 min)

Remember: Correct ↳ Fast

⚠ Critical Point

Mistake 3: Using del in loop

Error:

```
1 i = 0
2 while i < len(processes):
3     if not processes[i].is_alive:
4         del processes[i] # O(n) operation in loop!
5     else:
6         i += 1
7 # Total: O(n^2)
```

Fix: Use two-pointer approach ($O(n)$) or single deletion at end

When interviewer explains: Acknowledge, internalize, and use knowledge

⚠ Critical Point

Mistake 4: Insufficient locking

Error:

```
1 def garbage_collect(self):
2     with self.lock: # Only lock writes
3         # Swap elements
4
5 def find_ports(self, pid):
6     # NO LOCK - race condition!
7     for proc in self.processes: # Being modified by GC!
8         ...
```

Fix: ALL access to shared data must be synchronized

Key principle: Locks protect data, not functions

⚠ Critical Point

Mistake 5: Sloppy implementation

Feedback: "Part 2 implementation showed right ideas but needlessly complex and generated edge-case problems"

How to avoid:

- Write pseudocode first
- Test with simple example before finalizing
- Check boundary conditions (empty list, single element)
- Walk through code line-by-line with example

9 Final Preparation Checklist

9.1 Day Before Interview

✓ Action Item

Practice drill (do this 2-3 times):

1. Part 1 (20 min):

- Implement `find_all_ports()` from scratch
- Write 2-3 test cases
- Explain optimization

2. Part 2 (15 min):

- Implement `garbage_collect()` with two-pointer
- Test with example array
- Explain why `del` in loop is bad

3. Part 3 (10 min):

- Explain RW lock concept
- Sketch where locks go in code
- Discuss granularity trade-offs

Total practice time: 45 minutes per iteration

9.2 Interview Day Checklist

✓ Action Item

Before interview:

- ✓ Review process tree concepts
- ✓ Review hash set properties
- ✓ Review recursion patterns
- ✓ Review RW lock behavior
- ✓ Practice explaining your thinking out loud

During interview:

- ✓ Ask clarifying questions upfront
- ✓ Explain approach before coding
- ✓ Communicate assumptions
- ✓ Test code with examples
- ✓ Discuss complexity analysis
- ✓ Be receptive to hints

Red flags to avoid:

- ✗ Silent coding without explanation
- ✗ Not testing your code
- ✗ Ignoring interviewer hints
- ✗ Defensive about mistakes
- ✗ Giving up when stuck

9.3 Key Reminders

💡 Key Insight

Success factors from past feedback:

What got candidates hired:

- Quick recognition of recursion and flat-list issues
- Clear communication of design decisions
- Clean, correct implementations
- Strong understanding of complexity trade-offs
- Immediately picking up on swapping approach for GC

What led to rejection:

- Consistently substantial gaps in implementation
- Not understanding performance implications
- Sloppy code with edge case bugs
- Weak concurrency knowledge
- Poor communication about design choices

Bottom line:

- Correctness first, optimization second
- Communicate clearly throughout
- Be receptive to feedback and hints
- Test your code with examples
- Show systems thinking (not just coding)

10 Summary

10.1 Core Algorithms

1. Part 1 - Find All Ports:

- DFS recursion through process tree
- Hash set for deduplication
- Filter dead processes
- Optional: Precompute parent→children map
- Time: $O(n)$ optimized, Space: $O(n)$

2. Part 2 - Garbage Collection:

- Two-pointer in-place swap
- Move alive to front, dead to back
- Single pass, $O(n)$ time, $O(1)$ space
- Avoid `del` in loop ($O(n^2)$)

3. Part 3 - Thread Safety:

- RW lock for read-heavy workloads
- Multiple readers, single writer
- All access must be synchronized
- Discuss granularity trade-offs

10.2 Interview Mindset

Technical competence:

- Solid understanding of data structures
- Clean implementation skills
- Performance awareness

Communication:

- Explain approach before coding
- Think out loud
- Acknowledge and learn from hints

Problem-solving:

- Ask clarifying questions
- Start with correct solution
- Optimize after basics work

You've got this!