

ClickUp Backend Interview Document Event Processing

Complete Preparation Guide - November 30, 2024 (v2)

November 30, 2025

Contents

1 Interview Overview	4
1.1 What to Expect	4
1.2 Success Criteria	4
2 Core Problem: Document Event Processor	5
3 CRITICAL: Questions to Ask First!	7
4 Examples with Step-by-Step Trace	8
4.1 Example 1: Multiple Appends	8
4.2 Example 2: Delete Event	8
5 Solution Strategy	10
5.1 Approach	10
5.2 Key Insights	10
5.3 Why List Over String Manipulation?	10
6 Complete Solution	12
7 Robust Version with Validation	14
8 Critical Edge Cases	16
9 Comprehensive Test Suite	18
10 Debugging Strategy	20
11 Follow-Up Questions & Variations	21
11.1 Expected Follow-Ups	21
11.2 Variation 1: Insert at Character Position	21
11.3 Variation 2: Delete Specific Lines	22
11.4 Variation 3: Replace Line Content	22
12 Alternative Implementations	24
12.1 Object-Oriented Approach (Python)	24
13 System Design Discussion	26
13.1 Real-World ClickUp Architecture	26
13.2 Scalability Challenges	26
13.3 Data Structures for Large Documents	27

14 Interview Execution Strategy	28
14.1 Time Allocation (60 minutes)	28
14.2 Before You Code	28
14.3 While Coding	28
14.4 After Coding	29
15 Communication Checklist	30
15.1 Before Interview	30
15.2 During Interview - Execution	30
15.3 During Interview - Communication	30
16 Quick Reference Card	31
16.1 Key Points to Remember	31
16.2 Common Mistakes to Avoid	31
16.3 Python Quick Reference	31
17 Final Preparation Checklist	32
17.1 Technical Readiness	32
17.2 Interview Skills	32
17.3 Day Before Interview	32
17.4 Interview Day	32
18 PART 2: Search Architecture Deep Dive Prep	33
18.1 Interview Format	33
19 Priority 1: Multi-Tenanted Search (MUST KNOW)	34
19.1 The Problem	34
19.2 Strategy 1: Index-Per-Tenant	34
19.3 Strategy 2: Shared Index with Filtering	35
19.4 Strategy 3: Hybrid (Index Pools)	35
19.5 Preventing Cross-Tenant Data Leaks	36
19.6 ClickUp-Specific Multi-Tenancy	37
19.7 Performance Optimizations (HIGH IMPACT)	37
20 Priority 2: Relevance Scoring & Ranking	40
20.1 BM25 Algorithm (ElasticSearch Default)	40
20.2 Field Boosting	41
20.3 Custom Scoring	41
20.4 Interview Question: Rank "project update"	42
21 Priority 3: Query DSL Essentials	44
21.1 Core Query Types	44
21.2 Filter vs Query	44
21.3 ClickUp Search Query Example	45
22 Priority 4: Vector Search Refresher	46
22.1 When to Use Vector Search	46
22.2 Hybrid Search (Best Practice)	46
22.3 Vector Search Trade-offs	47
23 ElasticSearch vs OpenSearch	48
23.1 Key Differences	48
23.2 When to Choose Which	48

24 2-Hour Prep: Quick Reference Cheat Sheet	49
24.1 Must-Know Talking Points (Memorize These)	49
24.2 The 5-Minute Mental Model	49
24.3 Interview Question Templates	50
24.4 Code Snippets You Should Know	51
24.5 Key Numbers to Remember	51
24.6 Common Mistakes to Avoid	52
24.7 If You Blank on Details	52
24.8 Last-Minute Review (15 mins before)	52

1 Interview Overview

1.1 What to Expect

Interview Type: Backend Live Coding (60 minutes)

Format:

- CodeSignal platform (browser-based IDE)
- No pre-written test suite
- You must test your own code
- Recommended language: Python
- Focus on correctness, edge cases, and code quality

Problem Domain:

- Event processing in memory
- Document state management (similar to Google Docs)
- Real-time collaborative editing
- ClickUp's core functionality

1.2 Success Criteria

1. **Correctness:** Handle all test cases including edge cases
2. **Code Quality:** Clean, readable, well-structured code
3. **Communication:** Think out loud, explain your approach
4. **Testing:** Demonstrate testing strategy
5. **Problem-Solving:** Handle follow-up questions and variations

2 Core Problem: Document Event Processor

Problem

Context: ClickUp has a feature that allows users to create documents similar to Google Docs. Changes are tracked through batches of events.

Task: Process events on a document and return the correctly updated document.

Data Structures:

Document:

```
{  
    title: string,  
    content: string,           // Lines separated by \n  
    lastUpdated: timestamp,  
    createdOn: timestamp  
}
```

Event:

```
{  
    event_id: number,  
    event_name: string,      // Case-insensitive: "append", "APPEND", "Append"  
    payload: object,  
    timestamp: timestamp  
}
```

Event Types:

1. **append** - Add content to a specific line

```
payload: {  
    newContent: string,  
    startLine: number       // 1-indexed (line 1 = first line)  
}
```

Important: If content exists at that line, **append** newContent to the end of that line (don't replace!).

2. **delete** - Delete all content from document

```
payload: {} // Empty
```

Requirements:

- Process events in timestamp order (may arrive out-of-order)
- Update document.lastUpdated to latest event timestamp
- Handle case-insensitive event names
- Return updated document object

Important Edge Cases & Gotchas

IMPORTANT NOTE: The provided examples show lastUpdated incrementing by 1 ($123456789 \rightarrow 123456790$), but event timestamps are in the billions (1641024000001). This appears to be an inconsistency in the problem statement.

Most logical interpretation: Set lastUpdated to the timestamp of the most recent event.

Action: Ask the interviewer to clarify this behavior before coding!

3 CRITICAL: Questions to Ask First!

Pro Tips

Before you start coding, ask these clarifying questions:

1. **lastUpdated behavior:** "Should lastUpdated be set to the timestamp of the latest event, or should it be incremented by 1?"
 - The examples show inconsistent behavior - clarify this!
 - Most logical: use the latest event's timestamp
2. **Invalid line numbers:** "What should happen if startLine is 0 or negative?"
 - Skip the event? Return error? Assume line 1?
3. **Event validation:** "Should I validate event structure and handle malformed events?"
 - Missing required fields?
 - Unknown event types?
4. **Same timestamp:** "If multiple events have the same timestamp, what determines their order?"
 - Use event_id as tiebreaker?
 - Stable sort (preserve original order)?
5. **Document mutation:** "Should I modify the input document in place or return a new copy?"
 - Safer to return a copy
 - But in-place is more efficient
6. **Content initialization:** "If document doesn't have a content field initially, should I initialize it as empty string?"
7. **Trailing newlines:** "Should the final content have a trailing newline character?"

These questions show you think critically and catch ambiguities!

4 Examples with Step-by-Step Trace

4.1 Example 1: Multiple Appends

Input:

```
1 events = [
2     {"event_id": 1, "event_name": "append",
3      "payload": {"newContent": "Line 1 ", "startLine": 1},
4      "timestamp": 1641024000001},
5     {"event_id": 2, "event_name": "APPEND",
6      "payload": {"newContent": "Line 2 ", "startLine": 2},
7      "timestamp": 1641024000002},
8     {"event_id": 3, "event_name": "APPEND",
9      "payload": {"newContent": "Line 3 ", "startLine": 3},
10     "timestamp": 1641024000003}
11 ]
12
13 document = {"title": "Lorem Ipsum", "lastUpdated": 123456789, "createdOn": 123456789}
```

Step-by-Step Execution:

Initial state:

```
lines = []
lastUpdated = 123456789
```

After Event 1 (append to line 1):

```
lines = ["Line 1 "]
lastUpdated = 1641024000001
```

After Event 2 (append to line 2):

```
lines = ["Line 1 ", "Line 2 "]
lastUpdated = 1641024000002
```

After Event 3 (append to line 3):

```
lines = ["Line 1 ", "Line 2 ", "Line 3 "]
lastUpdated = 1641024000003
```

Final: Join lines with "\n"
content = "Line 1 \nLine 2 \nLine 3 "

Expected Output:

```
1 {
2     "title": "Lorem Ipsum",
3     "content": "Line 1 \nLine 2 \nLine 3 ",
4     "lastUpdated": 1641024000003,
5     "createdOn": 123456789
6 }
```

4.2 Example 2: Delete Event

Input:

```
1 events = [{"event_id": 1, "event_name": "delete", "timestamp": 1641024000000}]
```

```
2 document = {  
3     "title": "Lorem Ipsum",  
4     "content": "This is Lorem ipsum",  
5     "lastUpdated": 123456789,  
6     "createdOn": 123456789  
7 }  
8 }
```

Step-by-Step:

Initial state:

```
lines = ["This is Lorem ipsum"]
```

After Event 1 (delete):

```
lines = []  
content = ""  
lastUpdated = 1641024000000
```

Expected Output:

```
1 {  
2     "title": "Lorem Ipsum",  
3     "content": "",  
4     "lastUpdated": 1641024000000,  
5     "createdOn": 123456789  
6 }
```

5 Solution Strategy

5.1 Approach

1. Sort events by timestamp (handle out-of-order delivery)
2. Initialize content as list of lines
3. Process each event in order:
 - For append: Update/create line at specified position
 - For delete: Clear all content
4. Join lines back into string with \n
5. Update lastUpdated to latest event timestamp
6. Return updated document

5.2 Key Insights

- Use list for lines - O(1) indexing, easy modification
- Handle 1-indexed lines (convert to 0-indexed: line_idx = start_line - 1)
- Extend list with empty strings if startLine exceeds current length
- Event names are case-insensitive (use .lower())
- Append, don't replace - use lines[idx] += new_content
- Return a copy to avoid modifying input (safer)

5.3 Why List Over String Manipulation?

- Strings are immutable in Python - expensive to modify
- List operations are O(1) for append and indexed assignment
- Only convert to string once at the end
- Much more efficient for multiple operations

6 Complete Solution

Solution

Clean Production-Ready Implementation

```
1 def execute(events, document):
2     """
3         Process document events and return updated document.
4
5         This solution handles:
6         - Out-of-order events (sorts by timestamp)
7         - Case-insensitive event names
8         - Line gaps (fills with empty strings)
9         - Multiple appends to same line
10        - Document without initial content
11
12    Args:
13        events: List of event dictionaries
14        document: Document dictionary
15
16    Returns:
17        New document dictionary with processed changes
18    """
19    # Handle empty events - return unchanged
20    if not events:
21        return document
22
23    # Sort events by timestamp (handle out-of-order delivery)
24    sorted_events = sorted(events, key=lambda e: e["timestamp"])
25
26    # Initialize content as list of lines
27    content = document.get("content", "")
28    lines = content.split("\n") if content else []
29
30    # Track latest timestamp for lastUpdated
31    latest_timestamp = document.get("lastUpdated", 0)
32
33    # Process each event in chronological order
34    for event in sorted_events:
35        event_name = event["event_name"].lower()
36        timestamp = event["timestamp"]
37
38        if event_name == "append":
39            payload = event["payload"]
40            new_content = payload["newContent"]
41            start_line = payload["startLine"] # 1-indexed
42
43            # Convert to 0-indexed
44            line_idx = start_line - 1
45
46            # Extend lines list if necessary (fill gaps with empty
47            # strings)
48            while len(lines) <= line_idx:
49                lines.append("")
50
51            # Append to existing line content (don't replace!)
52            lines[line_idx] += new_content
53
54        elif event_name == "delete":
55            # Clear all content
56            lines = []
57
58            # Update to most recent timestamp
59            latest_timestamp = max(latest_timestamp, timestamp)
```


7 Robust Version with Validation

Solution

Enterprise Version with Error Handling

```
1 def execute_robust(events, document):
2     """
3         Robust version with comprehensive validation and error
4             handling.
5             Use this if interviewer asks about production considerations.
6     """
7     # Validate inputs
8     if not events:
9         return document.copy()
10
11    if not isinstance(events, list):
12        raise ValueError("events must be a list")
13
14    if not isinstance(document, dict):
15        raise ValueError("document must be a dictionary")
16
17    # Sort by timestamp (with fallback for missing timestamps)
18    sorted_events = sorted(events, key=lambda e: e.get("timestamp",
19                           0))
20
21    # Initialize content
22    content = document.get("content", "")
23    lines = content.split("\n") if content else []
24
25    # Remove trailing empty line if present (from content ending
26        in \n)
27    if lines and lines[-1] == "":
28        lines = lines[:-1]
29
30    latest_timestamp = document.get("lastUpdated", 0)
31
32    # Process each event with validation
33    for event in sorted_events:
34        # Validate event structure
35        if not isinstance(event, dict):
36            continue # Skip malformed events
37
38        event_name = event.get("event_name", "").lower()
39        timestamp = event.get("timestamp", 0)
40        payload = event.get("payload", {})
41
42        if event_name == "append":
43            # Validate payload
44            new_content = payload.get("newContent", "")
45            start_line = payload.get("startLine", 1)
46
47            # Validate line number (must be positive)
48            if start_line < 1:
49                continue # Skip invalid line numbers
50
51            line_idx = start_line - 1
52
53            # Extend lines if needed
54            while len(lines) <= line_idx:
55                lines.append("")
56
57            # Append content
58            lines[line_idx] += new_content
```


8 Critical Edge Cases

Important Edge Cases & Gotchas

YOU MUST HANDLE THESE:

1. Empty Events List

```
1 events = []
2 # Should return document unchanged
```

2. Out-of-Order Events - CRITICAL!

```
1 events = [
2     {"event_id": 2, ..., "timestamp": 102},  # Second
3     {"event_id": 1, ..., "timestamp": 101}    # First
4 ]
5 # Must sort by timestamp before processing!
```

3. Case-Insensitive Event Names

```
1 "append", "APPEND", "Append", "aPpEnd"  # All valid
2 # Use event_name.lower() for comparison
```

4. Multiple Appends to Same Line - MUST APPEND, NOT REPLACE!

```
1 events = [
2     {"payload": {"newContent": "Hello ", "startLine": 1},
3      ...},
4     {"payload": {"newContent": "World", "startLine": 1}, ...}
5 ]
# Result: lines[0] = "Hello World" NOT "World"
```

5. Gap in Line Numbers

```
1 events = [{"payload": {"newContent": "Line5", "startLine": 5},
2            ...}]
3 # Result: lines = ["", "", "", "", "Line5"]
4 # Fill gaps with empty strings!
```

6. Delete Then Append

```
1 events = [
2     {"event_name": "delete", "timestamp": 100},
3     {"event_name": "append", "payload": {..., "startLine": 1},
4      "timestamp": 101}
5 ]
# Delete clears lines = [], then append starts fresh
```

7. Document Without Content Field

```
1 document = {"title": "Test", "lastUpdated": 0, "createdOn": 0}
2 # No "content" key initially
3 # Use document.get("content", "") to handle safely
```

8. Empty Payload or Missing Fields

```
1 {"event_name": "delete", "timestamp": 100}
2 # Delete has no payload - that's valid
3 # Use payload.get("key", default) for safety
```

9. Line Number 0 or Negative

9 Comprehensive Test Suite

Test Cases

```
1 def test_document_processor():
2     """Complete test suite covering all edge cases."""
3
4     print("Running comprehensive test suite...")
5
6     # Test 1: Basic append to empty document
7     print("\n[Test 1] Basic append")
8     events = [
9         {
10             "event_id": 1,
11             "event_name": "append",
12             "payload": {"newContent": "Hello", "startLine": 1},
13             "timestamp": 100
14         }
15     ]
16     doc = {"title": "Test", "lastUpdated": 0, "createdOn": 0}
17     result = execute(events, doc)
18     assert result["content"] == "Hello", f"Expected 'Hello', got {result['content']}"
19     assert result["lastUpdated"] == 100, f"Expected 100, got {result['lastUpdated']}"
20     print(" PASSED")
21
22     # Test 2: Multiple appends to different lines
23     print("\n[Test 2] Multiple lines")
24     events = [
25         {"event_id": 1, "event_name": "append",
26             "payload": {"newContent": "Line1", "startLine": 1},
27             "timestamp": 100},
28         {"event_id": 2, "event_name": "append",
29             "payload": {"newContent": "Line2", "startLine": 2},
30             "timestamp": 101}
31     ]
32     doc = {"title": "Test", "lastUpdated": 0, "createdOn": 0}
33     result = execute(events, doc)
34     assert result["content"] == "Line1\nLine2"
35     print(" PASSED")
36
37     # Test 3: Multiple appends to SAME line (critical!)
38     print("\n[Test 3] Same line appends (CRITICAL)")
39     events = [
40         {"event_id": 1, "event_name": "append",
41             "payload": {"newContent": "Hello ", "startLine": 1},
42             "timestamp": 100},
43         {"event_id": 2, "event_name": "append",
44             "payload": {"newContent": "World", "startLine": 1},
45             "timestamp": 101}
46     ]
47     doc = {"title": "Test", "lastUpdated": 0, "createdOn": 0}
48     result = execute(events, doc)
49     assert result["content"] == "Hello World", \
50             f"Must APPEND not replace! Got: '{result['content']}'"
51     print(" PASSED")
52
53     # Test 4: Delete event
54     print("\n[Test 4] Delete clears all content")
55     events = [{"event_id": 1, "event_name": "delete", "timestamp": 100}]
56     doc = {"title": "Test", "content": "Some content",
57             "lastUpdated": 0, "createdOn": 0}
58     result = execute(events, doc)
59     assert result["content"] == ""
```


10 Debugging Strategy

Pro Tips

If Your Tests Are Failing:

1. Add Debug Prints

```
1 for event in sorted_events:
2     print(f"Processing: {event['event_name']} at line {event.
3         get('payload', {}).get('startLine')}")
4     print(f"Lines before: {lines}")
5     # ... process event ...
6     print(f"Lines after: {lines}")
7     print()
```

2. Check Event Sorting

```
1 print("Events before sorting:")
2 for e in events:
3     print(f"  {e['event_id']}: timestamp={e['timestamp']}")
4
5 sorted_events = sorted(events, key=lambda e: e["timestamp"])
6
7 print("\nEvents after sorting:")
8 for e in sorted_events:
9     print(f"  {e['event_id']}: timestamp={e['timestamp']})")
```

3. Verify Line Indexing

```
1 print(f"startLine={start_line} (1-indexed)")
2 print(f"line_idx={line_idx} (0-indexed)")
3 print(f"lines length before: {len(lines)}")
4 # ... extend lines ...
5 print(f"lines length after: {len(lines)})")
```

4. Check Append vs Replace

```
1 print(f"Line {line_idx} before: '{lines[line_idx]}'")
2 lines[line_idx] += new_content # Should use +=, not =
3 print(f"Line {line_idx} after: '{lines[line_idx]}'")
```

5. Verify Final Join

```
1 print(f"Lines array: {lines}")
2 content = "\n".join(lines)
3 print(f"Joined content: '{content}'")
4 print(f"Content length: {len(content)})")
```

Common Mistakes:

- Forgetting to sort events by timestamp
- Using `=` instead of `+=` for append
- Off-by-one error with 1-indexed vs 0-indexed
- Not handling case-insensitive event names
- Creating extra empty lines with improper split/join

11 Follow-Up Questions & Variations

11.1 Expected Follow-Ups

1. What if events can arrive significantly out of order?
 - Current solution handles this with sorting
 - For streaming: use priority queue or buffering window
 - Trade-off: latency vs correctness
2. How would you handle millions of events?
 - Batch processing
 - Periodic snapshots + incremental updates
 - Event sourcing pattern
 - Compaction/aggregation of old events
3. What about concurrent editing by multiple users?
 - Operational Transform (OT)
 - Conflict-free Replicated Data Types (CRDTs)
 - Last-write-wins with timestamps
 - Use event_id as tiebreaker
4. How would you optimize for very large documents?
 - Rope data structure instead of list
 - Lazy loading of content
 - Chunk-based storage
 - Only process visible viewport
5. What if we add more event types?
 - `insert`: Insert at character position within line
 - `delete_range`: Delete specific line range
 - `replace`: Replace content at line
 - `format`: Apply formatting (bold, italic)

11.2 Variation 1: Insert at Character Position

New Event Type:

```
1 {
2     "event_name": "insert",
3     "payload": {
4         "newContent": "text",
5         "startLine": 2,
6         "position": 5 # Character position in line (0-indexed)
7     },
8     "timestamp": 100
9 }
```

Implementation:

```

1  elif event_name == "insert":
2      new_content = payload["newContent"]
3      start_line = payload["startLine"]
4      position = payload["position"]
5
6      line_idx = start_line - 1
7      while len(lines) <= line_idx:
8          lines.append("")
9
10     line = lines[line_idx]
11     # Insert at character position
12     lines[line_idx] = line[:position] + new_content + line[position:]

```

11.3 Variation 2: Delete Specific Lines

New Event Type:

```

1 {
2     "event_name": "delete_range",
3     "payload": {
4         "startLine": 2,
5         "endLine": 4      # Inclusive
6     },
7     "timestamp": 100
8 }

```

Implementation:

```

1 elif event_name == "delete_range":
2     start_line = payload["startLine"]
3     end_line = payload["endLine"]
4
5     start_idx = start_line - 1
6     end_idx = end_line - 1
7
8     # Delete lines in range
9     if start_idx < len(lines):
10        del lines[start_idx:min(end_idx + 1, len(lines))]

```

11.4 Variation 3: Replace Line Content

New Event Type:

```

1 {
2     "event_name": "replace",
3     "payload": {
4         "newContent": "completely new text",
5         "startLine": 3
6     },
7     "timestamp": 100
8 }

```

Implementation:

```

1 elif event_name == "replace":
2     new_content = payload["newContent"]
3     start_line = payload["startLine"]
4

```

```
5     line_idx = start_line - 1
6     while len(lines) <= line_idx:
7         lines.append("")
8
9     # Replace entire line (use = instead of +=)
10    lines[line_idx] = new_content
```

12 Alternative Implementations

12.1 Object-Oriented Approach (Python)

```
1  class DocumentProcessor:
2      """OOP approach for document event processing."""
3
4      def __init__(self, document):
5          self.document = document.copy()
6          self.lines = self._init_lines()
7
8      def _init_lines(self):
9          """Initialize lines from document content."""
10         content = self.document.get("content", "")
11         return content.split("\n") if content else []
12
13     def process_events(self, events):
14         """Process all events and return updated document."""
15         if not events:
16             return self.document
17
18         sorted_events = sorted(events, key=lambda e: e["timestamp"])
19
20         for event in sorted_events:
21             self._process_event(event)
22
23         return self._finalize()
24
25     def _process_event(self, event):
26         """Process single event based on type."""
27         event_name = event["event_name"].lower()
28
29         if event_name == "append":
30             self._handle_append(event)
31         elif event_name == "delete":
32             self._handle_delete(event)
33         # Easy to add more event types here
34
35     def _handle_append(self, event):
36         """Handle append event."""
37         payload = event["payload"]
38         new_content = payload["newContent"]
39         start_line = payload["startLine"]
40
41         line_idx = start_line - 1
42
43         # Extend lines if needed
44         while len(self.lines) <= line_idx:
45             self.lines.append("")
46
47         self.lines[line_idx] += new_content
48         self.document["lastUpdated"] = event["timestamp"]
49
50     def _handle_delete(self, event):
51         """Handle delete event."""
52         self.lines = []
53         self.document["lastUpdated"] = event["timestamp"]
54
```

```
55     def _finalize(self):
56         """Finalize and return document."""
57         self.document["content"] = "\n".join(self.lines)
58         return self.document
59
60
61 # Usage
62 def execute(events, document):
63     processor = DocumentProcessor(document)
64     return processor.process_events(events)
```

13 System Design Discussion

13.1 Real-World ClickUp Architecture

In production, ClickUp likely uses:

1. Event Store / Event Sourcing

- Kafka or similar for event streaming
- Permanent event log (source of truth)
- Can replay events to rebuild state
- Enables time-travel debugging

2. CRDT (Conflict-free Replicated Data Types)

- Handle concurrent edits from multiple users
- Eventual consistency without conflicts
- Examples: Yjs, Automerge
- Used by: Figma, Notion, Google Docs

3. Operational Transform (OT)

- Transform conflicting operations
- Maintain causal ordering
- More complex but deterministic
- Used by: Google Docs originally

4. Snapshot + Delta Pattern

- Store periodic snapshots
- Apply only recent events
- Faster recovery and queries
- Reduce memory usage

5. WebSocket for Real-time Sync

- Push events to all connected clients
- Low latency updates
- Handle reconnection gracefully

13.2 Scalability Challenges

- **Large Documents:** Millions of characters, thousands of lines
- **High Event Rate:** Hundreds of events per second per document
- **Many Concurrent Users:** 10+ people editing simultaneously
- **Offline Support:** Sync when reconnected, handle conflicts
- **Undo/Redo:** Maintain operation history efficiently
- **Performance:** Sub-second response time even for large docs

13.3 Data Structures for Large Documents

- **Rope:** Tree-based string for efficient insertions/deletions
- **Gap Buffer:** Used by Emacs, good for cursor-based editing
- **Piece Table:** Used by VS Code, great for undo/redo
- **CRDT Text:** Yjs uses linked list with tombstones

14 Interview Execution Strategy

14.1 Time Allocation (60 minutes)

- **5 min:** Clarify requirements, ask questions, confirm understanding
- **3 min:** Discuss approach, mention data structures, complexity
- **2 min:** Write function signature and comments
- **25 min:** Implement core solution (clean, working code)
- **10 min:** Write and run tests (catch bugs early!)
- **5 min:** Add error handling and edge cases
- **10 min:** Follow-up questions, optimizations, discussion

14.2 Before You Code

Pro Tips

The First 5 Minutes Are Critical!

1. Clarify Requirements

- "Should lastUpdated use the event timestamp or increment?"
- "What happens with invalid line numbers?"
- "Should I validate event structure?"

2. Confirm Examples

- Walk through Example 1 verbally
- Confirm expected output matches your understanding
- Note any inconsistencies in examples

3. Discuss Approach

- "I'll sort events by timestamp first..."
- "I'll use a list for lines for O(1) indexing..."
- "Time complexity will be O(n log n) for sorting..."

4. Mention Edge Cases

- "I'll need to handle out-of-order events..."
- "Case-insensitive event names..."
- "Multiple appends to same line..."

This shows you're thinking critically and builds trust!

14.3 While Coding

1. Think Out Loud

- "Now I'll sort the events by timestamp..."

- ”Converting to 0-indexed here because Python lists...”
- ”Using += here to append, not replace...”

2. Write Clean Code

- Descriptive variable names: `line_idx`, not `i`
- Add comments for non-obvious logic
- Consistent spacing and formatting

3. Handle Errors Gracefully

- Use `.get()` for optional dictionary keys
- Check for `None`/empty before processing
- Validate inputs if time permits

4. Ask If Stuck

- ”I’m debating between X and Y, which would you prefer?”
- ”Should I prioritize robustness or simplicity here?”
- Don’t sit in silence - communicate!

14.4 After Coding

1. Test Immediately

- Run provided examples first
- Test edge cases (empty, out-of-order, same line)
- Fix any bugs found

2. Walk Through Code

- Explain your solution at high level
- Point out key design decisions
- Mention trade-offs considered

3. Discuss Improvements

- ”For production, I’d add validation...”
- ”Could optimize with rope data structure...”
- ”Would need CRDT for real-time collaboration...”

4. Handle Follow-Ups

- Be ready for variations (insert, delete range)
- Explain how to extend solution
- Discuss system design implications

15 Communication Checklist

15.1 Before Interview

- Practiced core solution multiple times (can code in 20-25 min)
- Memorized critical edge cases
- Tested with all provided examples
- Understand time/space complexity
- Reviewed follow-up variations
- Comfortable with Python
- Practiced explaining thought process out loud
- Prepared clarifying questions to ask

15.2 During Interview - Execution

- Asked clarifying questions upfront
- Confirmed understanding of examples
- Explained approach before coding
- Thought out loud while implementing
- Handled edge cases explicitly
- Wrote clean, readable code
- Tested code with examples
- Discussed complexity analysis
- Proposed optimizations
- Asked intelligent follow-up questions

15.3 During Interview - Communication

- Maintained conversational tone
- Explained reasoning for decisions
- Asked for feedback/hints when stuck
- Admitted when unsure (don't fake it)
- Showed enthusiasm and engagement
- Treated interviewer as collaborator

16 Quick Reference Card

16.1 Key Points to Remember

1. **ALWAYS sort events by timestamp first!**
2. **Use list for lines** (not string manipulation)
3. **Lines are 1-indexed** in problem, 0-indexed in Python
4. **Case-insensitive event names** - use `.lower()`
5. **APPEND to line** with `+=`, DON'T REPLACE with `=`
6. **Fill gaps** with empty strings when extending
7. **Update lastUpdated** to latest event timestamp
8. **Return a copy** - don't mutate input document

16.2 Common Mistakes to Avoid

- X Forgetting to sort events
- X Not handling case-insensitive event names
- X Using `=` instead of `+=` for append (replaces instead of appends!)
- X Off-by-one errors with 1-indexed vs 0-indexed
- X Not handling empty events or missing content
- X Modifying input document directly
- X Not testing edge cases

16.3 Python Quick Reference

```
1 # Safe dictionary access
2 content = document.get("content", "")      # Default empty string
3 payload = event.get("payload", {})         # Default empty dict
4
5 # Case-insensitive comparison
6 event_name = event["event_name"].lower()
7
8 # List operations
9 lines = []
10 lines.append("new")                      # Add to end
11 lines[idx] += "text"                     # Append to existing
12 while len(lines) <= idx:                # Extend with gaps
13     lines.append("")
14
15 # String operations
16 lines = content.split("\n")             # Split into list
17 content = "\n".join(lines)               # Join back to string
18
19 # Sorting
20 sorted_events = sorted(events, key=lambda e: e["timestamp"])
21
22 # List comprehension
23 valid = [e for e in events if e.get("timestamp") is not None]
```

17 Final Preparation Checklist

17.1 Technical Readiness

- Can implement solution in under 25 minutes
- All 12+ test cases pass without bugs
- Understand why each edge case matters
- Can explain time/space complexity
- Know how to extend for new event types
- Comfortable with alternative approaches (OOP, JS)

17.2 Interview Skills

- Practiced asking clarifying questions
- Can explain approach before coding
- Comfortable thinking out loud
- Know how to debug when tests fail
- Can discuss system design implications
- Ready for follow-up variations

17.3 Day Before Interview

- Do one final practice run (timed, 30 min)
- Review edge cases one more time
- Read through this guide's key sections
- Prepare 2-3 questions to ask interviewer
- Get good sleep - fresh mind is critical!

17.4 Interview Day

- Test internet connection and mic/camera
- Have this guide open for reference (if allowed)
- Water nearby, comfortable environment
- Positive mindset - you've got this!

18 PART 2: Search Architecture Deep Dive Prep

18.1 Interview Format

Session: Search Architecture Deep Dive (60 minutes)

Focus Areas:

- Multi-tenanted Search (ClickUp has workspaces - critical!)
- Search Ranking (BM25, field boosting, relevance)
- ElasticSearch/OpenSearch architecture
- Vector Search (semantic search for tasks/docs)
- Query DSL and practical search scenarios

Your Knowledge Level:

- Strong: Vector search concepts, high-level architecture, autocomplete
- Need refresh: BM25 details, field boosting, relevance scoring
- Gap: Multi-tenancy strategies, cross-tenant data leaks, Query DSL

19 Priority 1: Multi-Tenanted Search (MUST KNOW)

19.1 The Problem

ClickUp has multiple workspaces (tenants). When User A in Workspace 1 searches "project update", they should ONLY see results from their workspace, never from Workspace 2.

Critical Requirements:

- **Data Isolation:** Tenant A can't see Tenant B's data
- **Performance:** Don't slow down search for isolation
- **Scalability:** Support millions of tenants
- **Cost:** Balance resources vs isolation

19.2 Strategy 1: Index-Per-Tenant

Concept: Each tenant gets their own dedicated index.

```
workspace_123_tasks # Workspace 123's task index
workspace_456_tasks # Workspace 456's task index
workspace_789_tasks # Workspace 789's task index
```

Pros:

- Perfect isolation (impossible to leak data)
- Easy to delete tenant data (drop the index)
- Can tune per-tenant settings
- Clear billing/resource tracking

Cons:

- ES has limits (1000s of indices max)
- Overhead per index (shards, mappings)
- Can't search across tenants (admin features)
- Cluster management complexity

When to use:

- Small number of large tenants (B2B SaaS)
- Strict compliance requirements
- Tenants need custom settings

19.3 Strategy 2: Shared Index with Filtering

Concept: All tenants share one index, filter by tenant_id field.

```
tasks_index:  
{  
    "task_id": "t123",  
    "workspace_id": "ws_123", # <-- CRITICAL field  
    "title": "Project update",  
    "description": "..."  
}
```

Query Pattern:

```
1 {  
2     "query": {  
3         "bool": {  
4             "must": [  
5                 {"match": {"title": "project update"} }  
6             ],  
7             "filter": [  
8                 {"term": {"workspace_id": "ws_123"} } // ALWAYS filter!  
9             ]  
10        }  
11    }  
12 }
```

Pros:

- Scales to millions of tenants
- Lower overhead (fewer indices)
- Can search across tenants (admin)
- Easier cluster management

Cons:

- Data leak risk if filter missed
- Noisy neighbor problem
- Can't delete tenant data easily
- Query complexity increases

When to use:

- Many small-medium tenants (ClickUp likely uses this)
- Need cross-tenant analytics
- Cost optimization critical

19.4 Strategy 3: Hybrid (Index Pools)

Concept: Group tenants into pools, multiple tenants per index.

```
pool_1_tasks # Tenants 1-1000  
pool_2_tasks # Tenants 1001-2000  
pool_3_tasks # Tenants 2001-3000
```

Balance between isolation and scalability.

19.5 Preventing Cross-Tenant Data Leaks

Important Edge Cases & Gotchas

CRITICAL: How to Prevent Leaks

1. ALWAYS Filter at Query Time

```
1 // WRONG - can leak data!
2 {"query": {"match": {"title": "update"}}}
3
4 // CORRECT - always filter
5 {
6     "query": {
7         "bool": {
8             "must": [{"match": {"title": "update"}}],
9             "filter": [{"term": {"workspace_id": "ws_123"}}]
10        }
11    }
12 }
```

2. Middleware Enforcement

```
1 class SearchService:
2     def search(self, user, query):
3         # ALWAYS inject workspace_id from auth context
4         workspace_id = user.workspace_id
5
6         # Force filter into query
7         query["query"]["bool"]["filter"].append({
8             "term": {"workspace_id": workspace_id}
9         })
10
11     return es.search(query)
```

3. Index-Time Validation

- Verify workspace_id exists in every document
- Reject documents without workspace_id
- Validate workspace_id format

4. Security Audits

- Log all search queries
- Alert on queries without tenant filter
- Periodic security reviews
- Test with penetration testing

5. Query Templates

- Use parameterized query templates
- Never build queries with string concatenation
- workspace_id as required parameter

19.6 ClickUp-Specific Multi-Tenancy

ClickUp's Architecture (likely):

- Shared index with workspace_id filtering
- Tasks, docs, comments in same index (tagged by type)
- High-traffic workspaces might get dedicated indices
- Filter: `workspace_id AND user_permissions`

Interview Question: "How would you design search for ClickUp with 1M+ workspaces?"

Your Answer:

1. Start with shared index (scales to millions)
2. Every document: workspace_id (indexed, not analyzed)
3. Middleware ALWAYS injects workspace filter
4. Large workspaces (>10K users) → dedicated indices
5. Route queries to correct index pool
6. Security: audit logs, query templates, validation

19.7 Performance Optimizations (HIGH IMPACT)

1. Routing by workspace_id

Problem: Query searches ALL shards even if workspace data is on one shard.

Solution: Route documents and queries by workspace_id

```
1 # Index with routing
2 PUT /tasks/_doc/task_123?routing=ws_456
3
4 # Query with routing (searches only relevant shard!)
5 GET /tasks/_search?routing=ws_456
6
7 # Benefit: If workspace has 1M docs across 10 shards,
8 # query searches 1 shard instead of all 10!
9 # 10x performance improvement!
```

2. Hierarchical Filtering (ClickUp Realistic)

Beyond just workspace_id, ClickUp has complex permissions:

```
1 # Document structure
2 {
3     "task_id": "t123",
4     "workspace_id": "ws_456",
5     "accessible_by_user_ids": ["user_1", "user_2", "user_3"],
6     "accessible_by_team_ids": ["team_10", "team_15"],
7     "is_public": false
8 }
9
10 # Query: workspace AND (user OR team OR public)
11 {
12     "query": {
13         "bool": {
14             "must": [{"match": {"title": "project"}},
```

```

15 "filter": [
16   {"term": {"workspace_id": "ws_456"}},
17   {
18     "bool": {
19       "should": [
20         {"term": {"accessible_by_user_ids": "user_1"}},
21         {"terms": {"accessible_by_team_ids": ["team_10", "team_15"]}},
22         {"term": {"is_public": true}}
23       ],
24       "minimum_should_match": 1
25     }
26   }
27 ]
28 }
29 }
30 }
```

3. Caching Strategy

- **Query Cache:** ES caches entire query results
 - Best for: Repeated identical queries
 - Example: "status:open" query cached per workspace
- **Filter Cache:** Caches filter bit sets
 - `workspace_id` filters cached (high reuse!)
 - `status`, `priority` filters cached
 - Filters in "filter" context are cached (not "must")
- **Request Cache:** Caches aggregations
 - Dashboard: "How many tasks per status?"
 - Cached at shard level

4. Rate Limiting & Noisy Neighbors

```

1 # Prevent one workspace from overwhelming cluster
2 class SearchService:
3     def search(self, user, query):
4         workspace_id = user.workspace_id
5
6         # Rate limit per workspace (e.g., 100 QPS)
7         if not rate_limiter.allow(workspace_id, max_qps=100):
8             raise RateLimitError("Too many requests")
9
10        # Inject mandatory filter
11        query = inject_workspace_filter(query, workspace_id)
12        return es.search(query)
```

5. Tier-Based Architecture

Small workspaces (< 10K docs):
 → `shared_index_pool_1` (10K workspaces)

Medium workspaces (10K-100K docs):

→ shared_index_pool_2 (1K workspaces)

Large enterprise (> 100K docs):

- dedicated_index_enterprise_A
- dedicated_index_enterprise_B

Benefits:

- Large customers get guaranteed performance
- Small customers share resources efficiently
- Can migrate between tiers as workspace grows

Interview Impact: Mentioning routing + caching shows deep understanding!

20 Priority 2: Relevance Scoring & Ranking

20.1 BM25 Algorithm (ElasticSearch Default)

BM25 Formula (you should know this):

$$\text{score}(D, Q) = \text{IDF}(q_i) \times [f(q_i, D) \times (k_1 + 1)] / [f(q_i, D) + k_1 \times (1 - b + b \times |D|/\text{avgdl})]$$

Where:

- D = document
- Q = query
- q_i = query term i
- f(q_i, D) = frequency of term q_i in document D
- |D| = length of document D (in words)
- avgdl = average document length in collection
- k₁ = term frequency saturation parameter (default: 1.2)
- b = length normalization parameter (default: 0.75)
- IDF(q_i) = inverse document frequency of term q_i

Key Concepts:

1. **TF (Term Frequency):** More occurrences = higher score
 - But with diminishing returns (saturation)
 - "project project project" isn't 3x better than "project"
2. **IDF (Inverse Document Frequency):** Rare terms matter more
 - "the" appears everywhere → low IDF → low weight
 - "kubernetes" is rare → high IDF → high weight
3. **Document Length Normalization:** Longer docs penalized
 - Short doc with "project" beats long doc with same term
 - Parameter b controls strength (0=off, 1=full)
4. **Saturation (k1):** Diminishing returns on term frequency
 - 2nd occurrence matters less than 1st
 - 10th occurrence barely matters

BM25 vs TF-IDF:

- TF-IDF: Linear relationship (2x frequency = 2x score)
- BM25: Logarithmic (2x frequency ≈ 1.2x score)
- BM25 handles document length better
- BM25 is more robust to keyword stuffing

20.2 Field Boosting

Problem: Title matches should rank higher than body matches.

ElasticSearch Syntax:

```
1 {
2     "query": {
3         "multi_match": {
4             "query": "project update",
5             "fields": [
6                 "title^3",           // 3x boost
7                 "description^1",    // 1x (normal)
8                 "comments^0.5"      // 0.5x (less important)
9             ]
10            }
11        }
12    }
```

How Boosting Works:

- Base score calculated per field
 - Final score = $\text{title_score} \times 3 + \text{desc_score} \times 1 + \text{comments_score} \times 0.5$
 - Title match contributes 3x more than description

ClickUp Example:

```
1  {
2      "query": {
3          "multi_match": {
4              "query": "project update",
5              "fields": [
6                  "task_name^5",           // Task name most important
7                  "task_description^2",
8                  "comments^1",
9                  "attachments.filename^1.5"
10                 ]
11             }
12         }
13     }
```

20.3 Custom Scoring

Beyond BM25: Incorporate business logic.

Function Score Query:

```

13         "origin": "now",
14         "scale": "30d",
15         "decay": 0.5
16     }
17   }
18 },
19 {
20   "field_value_factor": {
21     "field": "likes_count", // Popular tasks rank higher
22     "modifier": "log1p"
23   }
24 }
25 ],
26 "boost_mode": "multiply"
27 }
28 }
29 }
```

ClickUp Ranking Factors:

1. Text relevance (BM25)
2. Task priority (high/medium/low)
3. Recency (newer tasks slightly boosted)
4. User activity (tasks user interacted with)
5. Completion status (open & completed)
6. Assignee (user's own tasks boosted)

20.4 Interview Question: Rank "project update"

Scenario: User searches "project update" in ClickUp.

Your Answer:

1. **Base Score:** BM25 on task name, description
2. **Field Boosts:** $task_name^5, description^2$ **Priority :** High priority tasks + 50% boost
3. **Recency:** Exponential decay over 90 days
4. **Status:** Open tasks +30%, completed -20%
5. **Personalization:** User's assigned tasks +40%
6. **Workspace Activity:** Recently viewed/edited +20%

```

1 {
2   "query": {
3     "function_score": {
4       "query": {
5         "bool": {
6           "must": [
7             "multi_match": {
8               "query": "project update",
9               "fields": ["task_name^5", "description^2"]
10            }
11          }
12        }
13      }
14    }
15  }
```

```
11     }] ,
12     "filter": [{"term": {"workspace_id": "ws_123"}}]
13   }
14 },
15 "functions": [
16   {"filter": {"term": {"priority": "high"}}, "weight": 1.5},
17   {"filter": {"term": {"status": "open"}}, "weight": 1.3},
18   {"filter": {"term": {"assignee_id": "user_abc"}}, "weight":
19     1.4},
20   {"gauss": {"updated_at": {"origin": "now", "scale": "90d"}}}
21 ]
22 }
23 }
```

21 Priority 3: Query DSL Essentials

21.1 Core Query Types

1. Match Query (Full-text search)

```
1 {"query": {"match": {"title": "project update"}}}  
2 // Analyzes text, finds "project" OR "update"
```

2. Term Query (Exact match, no analysis)

```
1 {"query": {"term": {"status": "open"}}}  
2 // Exact match, case-sensitive, used for IDs, enums
```

3. Bool Query (Combine multiple conditions)

```
1 {  
2     "query": {  
3         "bool": {  
4             "must": [{"match": {"title": "project"}}, , // AND  
5             "should": [{"match": {"tags": "urgent"}}, , // OR (boost)  
6             "must_not": [{"term": {"status": "deleted"}}, , // NOT  
7             "filter": [{"term": {"workspace_id": "ws_1"}}, // AND (no score)  
8         }  
9     }  
10 }
```

4. Multi-Match (Search across multiple fields)

```
1 {  
2     "query": {  
3         "multi_match": {  
4             "query": "update",  
5             "fields": ["title^3", "description"]  
6         }  
7     }  
8 }
```

5. Range Query

```
1 {  
2     "query": {  
3         "range": {  
4             "created_date": {  
5                 "gte": "2024-01-01",  
6                 "lte": "2024-12-31"  
7             }  
8         }  
9     }  
10 }
```

21.2 Filter vs Query

Query Context: Affects relevance score

```
1 "must": [{"match": {"title": "project"}}, // Scores results
```

Filter Context: Binary yes/no, cached, faster

```
1 "filter": [{"term": {"workspace_id": "ws_1"}}, // No scoring
```

Rule: Use filter for exact matches (status, IDs, dates). Use query for text search.

21.3 ClickUp Search Query Example

```
1 {
2     "query": {
3         "bool": [
4             "must": [
5                 {
6                     "multi_match": {
7                         "query": "project update",
8                         "fields": ["task_name^5", "description^2", "comments"]
9                     }
10                }
11            ],
12            "filter": [
13                {"term": {"workspace_id": "ws_123"}},
14                {"terms": {"status": ["open", "in_progress"]}},
15                {"range": {"created_date": {"gte": "now-1y"}}}
16            ],
17            "should": [
18                {"term": {"assignee_id": "user_abc"}}, // Boost user's tasks
19                {"term": {"priority": "high"}}
20            ],
21            "must_not": [
22                {"term": {"archived": true}}
23            ]
24        },
25        "sort": [
26            {"_score": "desc"},
27            {"updated_at": "desc"}
28        ],
29        "size": 20
30    }
31 }
```

22 Priority 4: Vector Search Refresher

22.1 When to Use Vector Search

Keyword Search (BM25): "kubernetes deployment"

- Finds exact keyword matches
- Fast, efficient, proven
- Fails on synonyms: "k8s deployment" won't match

Vector Search (Semantic): "container orchestration setup"

- Finds semantically similar content
- Matches even without exact keywords
- Slower, more expensive

ClickUp Use Case: User searches "meeting notes" → also finds "discussion summary", "call recap"

22.2 Hybrid Search (Best Practice)

```
1 {
2     "query": {
3         "bool": {
4             "should": [
5                 {
6                     "multi_match": {
7                         "query": "project update",
8                         "fields": ["title^3", "description"]
9                     }
10                },
11                {
12                    "knn": {
13                        "field": "description_vector",
14                        "query_vector": [0.23, -0.45, ...], // 768 dims
15                        "k": 10,
16                        "num_candidates": 100
17                    }
18                }
19            ],
20            "filter": [{"term": {"workspace_id": "ws_123"}}]
21        }
22    }
23 }
```

Combines:

- BM25 for exact keyword matches
- Vector search for semantic similarity
- Best of both worlds

22.3 Vector Search Trade-offs

Pros:

- Handles synonyms, paraphrases
- Language-agnostic (multilingual)
- Understands intent, context

Cons:

- 10-100x more storage (768-dimensional vectors)
- Slower (ANN search still expensive)
- Needs pre-trained embeddings model
- Black box (hard to debug)

ClickUp Strategy:

- Primary: BM25 (fast, accurate for exact matches)
- Fallback: Vector search if BM25 returns few results
- Hybrid: Combine scores for complex queries

23 ElasticSearch vs OpenSearch

23.1 Key Differences

- **Licensing:** ES went proprietary (2021), OpenSearch is Apache 2.0
- **Vendor:** ES by Elastic, OpenSearch by AWS + community
- **Features:** ES has newer features (ELSER, security). OpenSearch catching up.
- **Cloud:** AWS only supports OpenSearch, not ES
- **Compatibility:** OpenSearch maintains ES API compatibility (mostly)
- **Cost:** OpenSearch free, ES requires license for some features

For ClickUp: Likely using OpenSearch (AWS-hosted) or self-managed ES cluster.

23.2 When to Choose Which

Choose OpenSearch if:

- Using AWS (native integration)
- Want open-source without licensing concerns
- Cost-sensitive
- Need community control

Choose ElasticSearch if:

- Need latest Elastic features (ML, security)
- Want official Elastic Cloud
- Existing Elastic ecosystem
- Enterprise support critical

24 2-Hour Prep: Quick Reference Cheat Sheet

24.1 Must-Know Talking Points (Memorize These)

TOP 3 THINGS TO SAY

1. **Multi-tenancy:** "For ClickUp's scale with millions of workspaces, I'd use a shared index with workspace_id filtering enforced at the middleware layer, never trusting the client. Large enterprise customers (>100K docs) would get dedicated indices for performance guarantees."
2. **Performance:** "Key optimization is routing by workspace_id, which means queries search only 1 shard instead of all 10 - that's a 10x improvement. Combined with filter caching for workspace_id and rate limiting per tenant."
3. **Ranking:** "I'd start with BM25 for text relevance, boost important fields like task_name^5, and layer on business logic using function_score for priority, recency with exponential decay, and personalization."

24.2 The 5-Minute Mental Model

Multi-Tenancy Architecture:

Strategy: Shared Index + Filtering (scales to millions)

Document:

```
{  
  "workspace_id": "ws_123", ← ALWAYS FILTER ON THIS  
  "accessible_by_user_ids": [...],  
  "accessible_by_team_ids": [...],  
  "task_name": "...",  
  ...  
}
```

Query:

```
ALWAYS: bool → filter → term: workspace_id  
ALWAYS: Inject server-side (never trust client!)
```

Performance:

- Route by workspace_id (1 shard vs 10 shards)
- Cache workspace_id filters (high reuse)
- Rate limit per workspace (100 QPS)

BM25 Ranking:

$BM25 = TF \times IDF \times Length_Normalization$

Key params:

- $k_1 = 1.2$ (term frequency saturation)
- $b = 0.75$ (length normalization strength)

Why BM25 > TF-IDF:

- Logarithmic term frequency (prevents keyword stuffing)
- Better length normalization
- Tunable parameters

Field Boosting Pattern:

```
{  
  "multi_match": {  
    "query": "project update",  
    "fields": [  
      "task_name^5",           ← Most important  
      "description^2",  
      "comments^1"  
    ]  
  }  
}
```

Rule of thumb:

- Title/Name: 5x
- Description: 2x
- Comments/Body: 1x
- Metadata: 0.5x

24.3 Interview Question Templates

Q: "Design search for ClickUp with 1M+ workspaces"

A:

1. Architecture: Shared index with workspace_id field
2. Security: Middleware enforces filter (never trust client)
3. Performance: Route by workspace_id, cache filters
4. Tiering: Large workspaces ($>100K$ docs) → dedicated indices
5. Ranking: BM25 + field boosting + business logic

Q: "How do you prevent cross-tenant data leaks?"

A:

1. workspace_id from auth token (not query param)
2. Middleware ALWAYS injects filter before ES
3. Validate all docs have workspace_id at index time
4. Security audits: alert on queries without filter
5. Testing: penetration tests, chaos testing

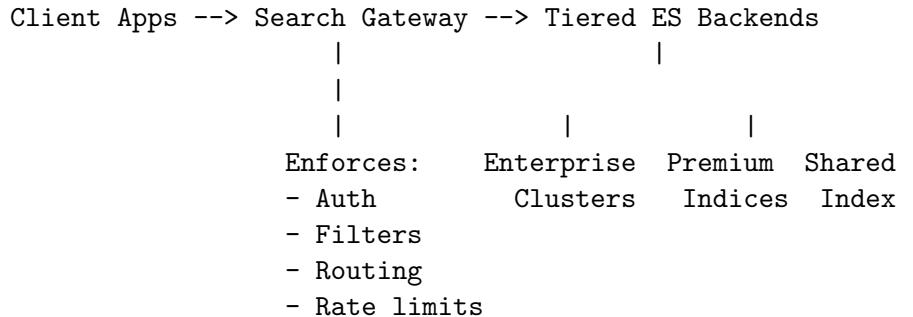
Q: "How do you rank 'project update' results?"

A:

1. Base: BM25 text relevance (ES default)
2. Fields: Boost task_name^5, description^2
3. Business: function_score for priority, status, assignee
4. Recency: Exponential decay over 90 days
5. Personalization: Boost user's tasks, team tasks

24.4 Production Architecture: Gateway Pattern

Key Insight: Clients NEVER access ES directly - use a gateway/facade!



Why Gateway Pattern?

- **Security:** Single enforcement point, no direct ES access
- **Flexibility:** Swap backends, A/B test, gradual migrations
- **Control:** Rate limiting, query validation, audit logging
- **Tiering:** Route large clients to dedicated resources

Gateway Code Pattern:

```
1  class SearchGateway:
2      def search(self, user, query_text):
3          # 1. Authenticate
4          if not user.is_authenticated():
5              raise Unauthorized()
6
7          # 2. Get user's workspaces
8          workspaces = self.get_user_workspaces(user.id)
9
10         # 3. Rate limit per workspace
11         if not self.rate_limiter.check(workspace_id):
12             raise RateLimitExceeded()
13
14         # 4. Build query with ENFORCED filters
15         query = {
16             "query": {
17                 "bool": {
18                     "must": [{"match": {"title": query_text}}],
19                     "filter": [
20                         {"term": {"workspace_id": workspace_id}},
21                         {"terms": {"accessible_by": [user.id]}}
22                     ]
23                 }
24             }
25         }
26
27         # 5. Route to correct backend tier
28         backend = self.route_to_backend(workspace_id)
29
30         # 6. Execute and audit log
31         results = backend.search(query)
32         self.audit_log(user.id, workspace_id, query_text)
```

```

33
34     return results
35
36     def route_to_backend(self, workspace_id):
37         tier = self.tenant_registry.get_tier(workspace_id)
38
39         if tier == "enterprise":
40             return self.enterprise_cluster
41         elif tier == "premium":
42             return self.premium_index
43         else:
44             return self.shared_index

```

Tiered Backend Strategy:

- **Enterprise (Top 0.1%)**: Dedicated ES cluster, 1000 QPS, 99.99% SLA
- **Premium (Top 5%)**: Dedicated index in shared cluster, 100 QPS
- **Shared (94.9%)**: Multi-tenant filtered index, 10 QPS

When to Mention:

- Interviewer asks: "How would you design this in production?"
- Discussing security: "In production, I'd add a gateway so clients never access ES directly"
- Discussing scale: "Gateway routes large tenants to dedicated resources"

24.5 Code Snippets You Should Know

1. Multi-tenancy Filter (CRITICAL):

```

1 def search(user, query_text):
2     workspace_id = user.workspace_id    # From auth
3
4     query = {
5         "query": {
6             "bool": {
7                 "must": [{"match": {"title": query_text}}],
8                 "filter": [{"term": {"workspace_id": workspace_id}}]
9             }
10            }
11        }
12    return es.search(index="tasks", body=query)

```

2. Field Boosting:

```

1 {
2     "query": {
3         "multi_match": {
4             "query": "project update",
5             "fields": ["task_name^5", "description^2", "comments^1"]
6         }
7     }
8 }

```

3. Function Score (Priority + Recency):

```

1  {
2      "query": {
3          "function_score": {
4              "query": {"match": {"title": "project"}},
5              "functions": [
6                  {
7                      "filter": {"term": {"priority": "high"}},
8                      "weight": 2
9                  },
10                 {
11                     "gauss": {
12                         "created_date": {
13                             "origin": "now",
14                             "scale": "30d",
15                             "decay": 0.5
16                         }
17                     }
18                 }
19             ],
20             "boost_mode": "multiply"
21         }
22     }
23 }
```

24.6 Key Numbers to Remember

- **BM25 k1:** 1.2 (term frequency saturation)
- **BM25 b:** 0.75 (length normalization)
- **Field boost ratios:** 5:2:1 (title:description:body)
- **Recency decay:** 90 days (exponential)
- **Dedicated index threshold:** >100K docs per tenant
- **Rate limit:** 100 QPS per workspace
- **ES index limit:** 1000s of indices (why shared index scales)

24.7 Common Mistakes to Avoid

1. "Use one index per user" → Doesn't scale to millions
2. Forgetting workspace_id filter → Data leak!
3. Client-side filtering → Security vulnerability!
4. Not mentioning routing → Missed performance win
5. Saying "TF-IDF" instead of "BM25" → Shows outdated knowledge
6. Only BM25, no business logic → Not production-ready

24.8 If You Blank on Details

Safe fallback phrases:

- "I'd need to check the exact ES API syntax, but the concept is..."
- "In production, I'd validate this with load testing to tune parameters..."
- "I'd start with the default (BM25/k1=1.2) and A/B test adjustments..."
- "Security is critical here - I'd enforce filters server-side and add audit logging..."

When to say you don't know:

- "I haven't tuned BM25 parameters in production, but I understand the k1 and b parameters control saturation and length normalization."
- "I'm familiar with the concept of vector search, but I'd want to measure precision/recall before choosing between semantic and keyword search."

24.9 Last-Minute Review (15 mins before)

Read these 3 sections:

1. Multi-Tenancy: Strategy 2 (Shared Index) + Security (lines 1558-1682)
2. Performance Optimizations: Routing + Caching (lines 1706-1822)
3. Interview Scenarios: Scenario 1 & 2 (lines 2126-2220)

Mental checklist before you start:

- I can explain shared index + workspace_id filtering
- I know BM25 is better than TF-IDF (logarithmic term frequency)
- I can write a multi_match query with field boosting
- I understand routing improves performance (1 shard vs 10)
- I know security: server-side filter enforcement is critical

You Are Ready!

You've prepared thoroughly. You know the solution inside-out. You understand the edge cases. You can handle follow-ups. Now trust your preparation and show them what you can do!

Remember:

- The interviewer wants you to succeed
- They're evaluating how you think, not just coding speed
- Communication matters as much as the solution
- It's okay to ask questions - that's smart!
- One bug doesn't fail you - recovery matters

Good luck with your ClickUp interview!