

DataHub Day 2 Interview Preparation

Project Deep Dive: Game Clickbait Detection System

Alex Yang
Principal Software Engineer, Roblox

Interview Date: November 20, 2024
Prepared: November 19, 2024

Contents

1 Interview Overview	4
1.1 Format	4
1.2 Selected Project	4
2 30-Second Executive Summary	4
3 Project Narrative (STAR Framework)	4
3.1 Situation (2-3 minutes)	4
3.1.1 The Problem	4
3.1.2 Previous Team's Approach (2 months of work)	5
3.1.3 Problem Reframing (First Principles)	5
3.2 Task (1-2 minutes)	7
3.2.1 My Responsibilities	7
3.2.2 Constraints	7
3.2.3 Success Criteria	7
3.3 Action (5-7 minutes)	8
3.3.1 Core Innovation: Protected Assets Library (PAL)	8
3.3.2 Signal A: Text Protection (Rule-Based Heuristics)	8
3.3.3 Signal B: Image Protection (CLIP + PAL)	9
3.3.4 Signal C: Logo Protection (YOLOv8 + CLIP)	9
3.3.5 Real-Time Detection Architecture	10
3.3.6 Offline Backfill Pipeline	11
3.3.7 Cross-Team Collaboration	12
3.4 Result (1-2 minutes)	13
3.4.1 Metrics	13
3.4.2 Timeline	13
3.4.3 Handoff	13
3.5 Technical Challenges & Lessons (2-3 minutes)	14
3.5.1 Challenge 1: Balancing Precision vs Coverage	14
3.5.2 Challenge 2: Real-Time Latency Requirements	14
3.5.3 Challenge 3: Fine-Tuning YOLOv8 for Logos	14
3.5.4 Challenge 4: Policy Rollout Without Creator Backlash	14

3.5.5 Challenge 5: Defining "Clickbait" vs "Poor Experience"	15
3.5.6 What I'd Do Differently	15
4 Connection to DataHub	16
5 Mock Q&A Scenarios	16
5.1 Technical Deep Dive Questions	16
5.1.1 Q: Why didn't you use a machine learning model for title detection?	16
5.1.2 Q: How did you handle the massive backfill of hundreds of millions of games?	17
5.1.3 Q: What metrics did you use to evaluate success?	18
5.1.4 Q: How did you ensure high precision to avoid false positives?	18
5.1.5 Q: Walk me through your collaboration with the Brand Safety team on logo detection.	19
5.2 Systems Design Questions	20
5.2.1 Q: Why event-driven architecture for real-time detection? What are the trade-offs?	20
5.2.2 Q: How would you handle a sudden 10x spike in game update traffic?	20
5.3 Project Management Questions	21
5.3.1 Q: How did you convince stakeholders to abandon 2 months of work?	21
5.3.2 Q: What was the biggest challenge in coordinating 5 teams?	22
5.3.3 Q: You mentioned leadership didn't fully understand the project. What happened there?	23
6 Practice Checklist	26
7 Key Talking Points (Memorize These)	27
8 Questions to Ask Abe	28
9 Final Reminders	29
10 Part 2: Search Architecture Deep Dive	30
10.1 Overview	30
10.2 Roblox Search Architecture: High-Level Overview	30
10.2.1 The Roblox Search Problem Space	30
10.2.2 End-to-End Search Flow	31
10.3 Your Contributions to Roblox Search	32
10.3.1 1. RSS (Roblox Similarity Search) - Vector Database Infrastructure	32
10.3.2 2. RAG Platform - Retrieval-Augmented Generation	33
10.3.3 3. Autocomplete Services (Avatar, Studio, Brand)	34
10.3.4 4. Spelling Correction Service	35
10.3.5 5. Game Clickbait Detection (Covered in Part 1)	36
10.4 Key Search Topics: Deep Dive	36
10.4.1 Indexing Pipeline	36
10.4.2 Query Understanding	37
10.4.3 Ranking & Relevance	38
10.4.4 Semantic Search (CLIP + RSS)	39
10.4.5 Personalization	40
10.4.6 Evaluation & Metrics	40

10.5 Search Architecture Trade-Offs	41
10.5.1 Latency vs Quality	41
10.5.2 Freshness vs Consistency	42
10.5.3 Recall vs Precision	42
10.5.4 Personalization vs Diversity	42
10.6 Potential Interview Questions & Answers	42
10.6.1 Q: Walk me through what happens when a user types a search query.	42
10.6.2 Q: How do you handle very popular queries that could overwhelm the system?	43
10.6.3 Q: What's your strategy for ranking results for a brand new user?	44
10.6.4 Q: How do you detect and fix search quality issues?	46
10.6.5 Q: How would you add a new ranking signal to the system?	47
10.6.6 Q: What's your biggest challenge with search at Roblox's scale?	49
10.7 Connection to DataHub (Search Context)	51
A Quick Reference: Key Numbers	52
B Quick Reference: Architecture Components	52
C Quick Reference: Contact Info	52

1 Interview Overview

1.1 Format

- **Interviewer:** Abe (DataHub)
- **Duration:** 45-60 minutes
- **Part 1:** Project Deep Dive (25-30 min)
- **Part 2:** Search Architecture Discussion (20-25 min)

1.2 Selected Project

Game Clickbait Detection System (Q1 2024, Roblox)

Why this project:

- Complete end-to-end ownership (problem definition -> architecture -> deployment)
- High-visibility impact (CEO escalation -> townhall recognition)
- Cross-team leadership (5 teams: DSA, Human Eval, Brand Safety, ML Platform, Game Discovery)
- Systems thinking (IP protection, search quality, policy framework)
- Measurable success (99%+ precision, real-time detection)
- Recent and fresh (2024)

2 30-Second Executive Summary

⚠ Critical Point

Practice delivering this in 30 seconds:

"In Q1 2024, our CEO escalated an urgent search quality issue: copycat games with identical names and images polluted results when searching for popular games like 'Brookhaven'. A team had spent 2 months on a CLIP-based image similarity approach, but when I evaluated it, precision was near-random.

I identified the core issue within days, threw away the previous work, and redesigned from first principles. The result: a real-time, multi-modal IP protection system with 99%+ precision that processes game updates in 10 seconds, protects hundreds of assets across billions of games, and was recognized at company townhall."

3 Project Narrative (STAR Framework)

3.1 Situation (2-3 minutes)

3.1.1 The Problem

CEO Observation (Holiday 2023):

- Searching "Brookhaven" (top Roblox game) returned pages of copycat games
- Copycats used identical/similar names and thumbnail images
- Search results looked "ugly" and hurt platform trust
- Users clicking copycats experienced low-quality, unrelated gameplay

3.1.2 Previous Team's Approach (2 months of work)

Method: CLIP-based image similarity detection

1. Pipeline to get top 100 popular games
2. Extract game thumbnail images
3. Generate CLIP embeddings for all game thumbnails (daily/weekly)
4. Find games with thumbnails semantically close to top 100
5. Demote these games in ranking

My Evaluation (within days):

- Ran human evaluation on their results with different thresholds (0.9, 0.95)
- **Result: Precision near-random (~50%)**
- Many false positives (legitimate games with similar art styles)
- Many false negatives (modified images evaded detection)
- **Critical flaw: Completely ignored title duplication signal**

💡 Key Insight

Why the previous approach failed:

- Treated as pure image similarity problem
- CLIP clusters by visual style, not deceptive intent
- False positives: Anime-style games, similar genres (tycoon, obby)
- False negatives: Color shifts, crops, overlays evade detection
- Ignored strongest signal: title text copying

3.1.3 Problem Reframing (First Principles)

Key question I asked: "Why is this even a problem?"

If authentic Brookhaven ranks #1, users can identify it. So what are we *really* solving?

Two distinct problems identified:

1. **IP Protection (Primary)**

- Platform responsibility to protect creator intellectual property
- Games stealing name + image = trademark/copyright violation
- Roblox's reputation as fair UGC platform at stake

2. Search Quality (Secondary)

- Copycats are typically low-quality (quick cash grabs)
- Crowd out higher-quality, diverse alternatives
- Users want: original game (#1) + quality similar games, NOT 50 low-effort clones

Design Principle: Multi-signal detection for IP violations, not just image matching

3.2 Task (1-2 minutes)

3.2.1 My Responsibilities

- **Technical architecture:** Design end-to-end detection system
- **Cross-team coordination:** Align 5 teams on approach and deliverables
- **Policy framework:** Define what constitutes IP violation
- **Implementation:** Build real-time detection service and offline backfill
- **Evaluation:** Establish metrics and validation pipeline
- **Deployment:** Launch to production with A/B testing

3.2.2 Constraints

- **Time:** CEO urgency - need results in Q1 2024
- **Scale:** Hundreds of millions of existing games to process
- **Latency:** Real-time updates (≤ 1 minute from game publish to ranking impact)
- **Precision:** Must avoid false positives (creator backlash risk)
- **Coverage:** Top 30 games covers 20-30% of search queries (Pareto principle)

3.2.3 Success Criteria

- 95%+ precision on clickbait detection
- Real-time detection (≤ 1 min latency)
- Backfill all existing games
- Zero creator backlash (clear, fair policies)
- CEO satisfaction (visible problem resolution)

3.3 Action (5-7 minutes)

3.3.1 Core Innovation: Protected Assets Library (PAL)

Concept: Centralized registry of IP assets to protect
Schema:

```
Field: type      (icon, title, logo, etc.)
Field: creatorId
Field: universeId
Field: rootPlaceId
Field: titleMatchPolicy (for title assets: exact, substring, fuzzy)
```

Violation Logic: Any signal triggered -> demotion in ranking

Three Asset Types:

- **Signal A:** Text (game names)
- **Signal B:** Images (full thumbnails)
- **Signal C:** Logos (brand marks within images)

3.3.2 Signal A: Text Protection (Rule-Based Heuristics)

Why rule-based, not semantic similarity?

- "Microhard" doesn't violate "Microsoft" IP (semantic but not legal violation)
- Need legal precision, not fuzzy matching
- Semantic similarity failed in evaluation (no practical utility)
- Heuristics achieved 99%+ precision on golden dataset

Three Match Policies (Flexible Framework):

1. **exact:** Strict 1:1 match
 - Use case: "Evade" (common English word, don't over-protect)
 - Normalization: lowercase, remove spaces, strip emojis/brackets
2. **substring:** Contains match
 - Use case: "Welcome to Bloxburg" prevents "Welcome to Bloxburg Christmas!"
 - Protects against title stuffing: "BrookhavenRP Adopt Me! Pet Simulator"
3. **fuzzy:** Edit distance tolerance
 - Use case: "Brookhaven" with fuzzy=2 catches "bro0khaven", "Br00khaven"
 - Prevents simple evasion tactics

Built-in Exceptions (Nuanced Policy):

- **Fan-made allowed:** "Brookhaven RP [Fan Version]" NOT clickbait

- **Same creator/group:** Can create own spinoffs (ownership check)
- **Typo tolerance:** "Fan Versoin" still triggers (prevents evasion)

Cross-functional Work:

- Collaborated with Community Program Manager on policy drafting
- Community announcement rollout (several weeks)
- Clear creator communication through official channels
- **Result:** Zero creator backlash

Impact: Solved ~30% of clickbait cases

3.3.3 Signal B: Image Protection (CLIP + PAL)

Key Difference from Previous Approach:

- **Previous:** Compare ALL games to top 100 (noisy, high false positive rate)
- **Mine:** Compare only against curated PAL library (precise, targeted)

Implementation:

- Human Eval team curates protected assets using Label Studio
- CLIP embeddings generated for full thumbnails
- Embeddings stored in RSS (Roblox Similarity Search - vector database)
- Fine-tuned similarity thresholds per PAL entry
- Monthly pipeline reviews top 30 games + ad-hoc updates

PAL Evolution:

- Started: ~30-100 entries (top games)
- Current: Hundreds of entries
- Planned: External brands (Disney, Coca-Cola, etc.)

Precision: 95%+ on image-based clickbait

3.3.4 Signal C: Logo Protection (YOLOv8 + CLIP)

Two-Stage Detection Pipeline:

1. Stage 1: YOLOv8 Logo Detection (Fine-tuned by me)

- Detect logo regions within game thumbnails
- Training: Human Eval drew bounding boxes in Label Studio
- Fine-tuned YOLOv8 on Roblox-specific logo corpus
- Accuracy: 95%+

2. Stage 2: CLIP Logo Comparison

- Extract detected logo regions
- Generate CLIP embeddings for logos
- Compare against PAL logo library (vector search)

Cross-Team Integration:

- Leveraged Brand Safety team's logo detection endpoint
- Established versioning protocol for future model updates:
 - Old endpoint maintained
 - New versioned endpoints for A/B testing
 - Enables safe model migrations

Combined Precision: 99.5%+ on logo-based clickbait

3.3.5 Real-Time Detection Architecture

Design Shift:

- **Previous:** Weekly batch processing (stale, reactive)
- **Mine:** Event-driven, real-time (10-second latency)

System Architecture:

SNS Topics (Game Lifecycle Events):

- PRODUCTION_PlaceEntityChangeEvent
- PRODUCTION_UniverseDisplayInformationChangeEvent
- PRODUCTION_UniverseEntityChangeEvent

|

SQS Queue (~10 QPS)

|

Queue Processor Service

|

Clickbait Detection Service (bedev2 microservice)

- |-- Text rule matching (heuristics)
- |-- CLIP embedding generation -> RSS query
- \-- YOLO logo detection -> Brand Safety endpoint

|

Frost Feature Store

Table: sdp_production_nonpii.clickbait_games_features_v0

Columns: universe_id, is_clickbait, metadata (JSON), updated_time_unix

|

Experience Document Builder (SNS topic)

|

Elasticsearch Index (for ranking)

|

Search Retrieval & Ranking Pipeline

Key Design Decisions:

- **Event-driven:** Triggered by game lifecycle, not polling (efficient)
- **SQS buffering:** Handles traffic spikes, decouples services (resilient)
- **Microservice:** Modular CLIP/YOLO inference (scalable)
- **Feature serving:** Signals propagate to ranking in ~10 seconds (fast)
- **Non-blocking:** Doesn't delay game creation/updates (user-friendly)

API Features:

- REST endpoints for on-demand detection
- Manual override API for edge cases (parodies, legitimate references)
- JSON metadata field for detailed signal breakdown

3.3.6 Offline Backfill Pipeline

Challenge: Process hundreds of millions of existing games

Solution: Modular MLP Pipelines

1. Separate Pipeline per Feature (title, thumbnail, logo)

- Load games corpus from Hive
- Load PAL
- Run similarity detection (text heuristics or vector search)
- Output to intermediate Hive tables

2. Powerhouse Aggregation Pipeline

- Load input tables: `sdp_icon_similarity`, `sdp_title_similarity`
- Aggregate scores into boolean `is_clickbait` column
- Publish batch to Frost feature store

Benefits of Modular Design:

- Scale compute independently per feature
- Iterate on one feature without reprocessing all
- Recover from failures without full reruns
- Clear separation of concerns (detection vs aggregation)

3.3.7 Cross-Team Collaboration

Five Teams Coordinated:

1. DSA (Jinlong Ji)

- Owned initial backfill pipeline architecture
- Top-K game collection logic
- Hive table management

2. Human Eval (Patricia Morales)

- Created PAL using Label Studio
- Drew logo bounding boxes for YOLOv8 training
- Ongoing evaluation pipeline (human-in-the-loop)

3. Brand Safety Team

- Provided logo detection endpoint
- Versioning protocol for future model updates

4. ML Platform Team

- Inference infrastructure (CLIP/YOLO serving)
- Model deployment support

5. Game Discovery Team

- Search ranking integration
- Final ownership post-launch
- Ongoing policy iteration

Key Coordination Challenge:

Handling abandonment of 2 months of work:

- Presented evaluation results showing near-random precision
- Data spoke for itself - approach wasn't salvageable
- Principal ML Engineer moved to other priorities (no friction)
- Key: Show problem clearly + propose concrete alternative (not just criticize)

3.4 Result (1-2 minutes)

3.4.1 Metrics

Precision/Recall:

- **Previous approach:** ~50% precision (near-random)
- **Text rules:** 99%+ precision, solved 30% of cases
- **CLIP (full image):** 95%+ precision
- **YOLO + CLIP (logos):** 99.5%+ precision

Scale:

- **PAL:** Hundreds of protected assets
- **Coverage:** 20-30% of top search queries
- **Backfill:** Hundreds of millions of existing games processed
- **Real-time:** ~10 QPS game updates, ~10-second latency end-to-end

Business Impact:

- **CEO satisfaction:** Problem visibly resolved for top queries
- **Townhall recognition:** Company-wide acknowledgment
- **Policy framework:** Established IP protection standards for platform
- **Expandable:** Foundation for other fraud detection (bait-and-switch teleport games)
- **Zero creator backlash:** Clear communication, fair policies

3.4.2 Timeline

- **Week 1:** Identified previous approach failure, proposed redesign
- **Milestone 1 (March 1):** Title text protection prod launch + A/B
- **Milestone 2 (March 15):** Thumbnail protection prod launch
- **Milestone 3 (March 22):** Logo protection prod launch
- **Q1 2024 End:** Evaluation pipeline, optimization, handoff to Game Discovery
- **Ongoing:** Maintained and evolved system for several quarters

3.4.3 Handoff

- Successfully transferred ownership to Game Discovery team
- Modular architecture enabled them to iterate on policies without rewriting core services
- Continued evolution: bait-and-switch detection, expanded PAL, etc.

3.5 Technical Challenges & Lessons (2-3 minutes)

3.5.1 Challenge 1: Balancing Precision vs Coverage

Problem: Can't manually curate PAL for millions of games

Solution:

- Focus on top 30 games (Pareto principle: 20-30% query coverage with minimal effort)
- Human-in-the-loop for high-profile additions
- Scalable pipeline for monthly reviews

Lesson: Perfect is the enemy of good - solve 80% of the problem with 20% of the effort

3.5.2 Challenge 2: Real-Time Latency Requirements

Problem: Game creation can't be blocked; search results must reflect changes quickly

Solution:

- Async event processing (SQS decoupling)
- Non-blocking: Detection happens post-creation
- Acceptable staleness: 10-second propagation vs instant blocking

Lesson: Understand business constraints - near real-time (seconds) \gg batch (days)

3.5.3 Challenge 3: Fine-Tuning YOLOv8 for Logos

Problem: Off-the-shelf YOLO detects generic objects, not game logos

Solution:

- Created custom dataset with Human Eval (bounding boxes in Label Studio)
- Fine-tuned YOLOv8 on Roblox-specific logo corpus
- Achieved 95%+ accuracy on domain-specific task

Lesson: Domain-specific ML requires domain-specific data - generic models aren't enough

3.5.4 Challenge 4: Policy Rollout Without Creator Backlash

Problem: New restrictions could anger creators (platform risk)

Solution:

- Clear, fair rules (spinoff patterns allowed)
- Transparent communication (official community announcements)
- Gradual enforcement (education before penalties)
- Explainable heuristics ("Your title matches X with fuzzy=1" vs "Model score 0.87")

Lesson: Technical solutions need social/policy layer - engineering alone isn't enough

3.5.5 Challenge 5: Defining "Clickbait" vs "Poor Experience"

Problem: Initial discussions blurred clickbait (pre-game deception) with poor in-game quality

My Framing:

- **Clickbait** = misleading *pre-game signals* (title, thumbnail) regardless of in-game content
- **Poor experience** = separate problem requiring in-game scene analysis (future scope with RFM models)

Why This Mattered:

- If evaluated based on in-game content, pre-game detection algorithms would falsely appear to fail
- Example: Game copies "Brookhaven" title/image but has transformative gameplay → still clickbait for IP protection

Solution:

- Created golden dataset focused *only* on pre-game signals
- Separated in-game content analysis to "poor experience detection" (out of scope for V1)
- Enabled clear algorithm evaluation and explainable enforcement

Lesson: Precise problem definition is foundational - ambiguity leads to misaligned evaluation

3.5.6 What I'd Do Differently

1. Earlier evaluation framework

- Built golden dataset evaluation late
- Should've been day-one priority to validate approach faster

2. More automated PAL maintenance

- Top 30 games required manual monthly review
- Could've built automated pipeline detecting trending games for Human Eval queue

3. Better internal documentation

- Leadership didn't fully understand project complexity (part of reason for seeking new opportunities)
- Should've created exec-level summary decks: problem → solution → impact

What Worked Well:

- Modular architecture paid off - Game Discovery could iterate on policies without rewriting core services
- Cross-team coordination was smooth due to clear work streams and ownership
- First principles thinking caught fundamental flaw early (saved months of wasted effort)

4 Connection to DataHub

💡 Key Insight

How This Project Maps to DataHub's Challenges:

"This project taught me that **metadata quality** - whether it's game titles, thumbnails, or IP assets - is foundational to platform trust. At DataHub, metadata quality is the *core product*.

I see direct parallels:

- **Protected Assets Library (PAL) ≈ DataHub's metadata registry**
 - Both are centralized sources of truth
 - Both require curation, governance, and maintenance
 - Both enable downstream systems (search ranking vs data discovery)
- **Cross-team adoption** was key to my success at Roblox
 - Got 5 teams (DSA, Human Eval, Brand Safety, ML Platform, Game Discovery) to trust and contribute to PAL
 - DataHub faces same challenge: getting engineering, data science, and business teams to trust and contribute metadata
- **Real-time propagation** (game updates → search ranking in 10 seconds)
 - DataHub needs similar: schema changes → lineage updates → downstream awareness
 - Event-driven architecture patterns directly applicable
- **Policy framework** (exact/substring/fuzzy matching policies)
 - Flexible policies ↳ one-size-fits-all
 - DataHub needs governance policies: who can edit what, approval workflows, data classification rules

I'm excited about DataHub because it's solving the *same fundamental problem* - metadata quality and governance - at the data infrastructure level rather than search/discovery."

5 Mock Q&A Scenarios

5.1 Technical Deep Dive Questions

5.1.1 Q: Why didn't you use a machine learning model for title detection?

Answer:

"I evaluated both approaches. For title matching, heuristics significantly outperformed semantic similarity models for three reasons:

1. **Explainability:** We needed to communicate to creators *exactly why* they were flagged for

policy compliance. 'Your title matches protected title X with fuzzy distance 1' is clear; 'model confidence 0.87' is not. This was critical for avoiding creator backlash.

2. **Semantic similarity misalignment:** ML models would flag 'Zombie Night' as similar to 'Halloween Night' based on semantic meaning, but that's not IP violation. Clickbait is about exact/near-exact copying, not conceptual similarity.
3. **Pattern recognition for evasion:** Common evasion tactics like 'title stuffing' - cramming multiple popular titles like 'BrookhavenRP Adopt Me! Pet Simulator' - are trivial to detect with substring rules but hard for models to learn without massive training data and constant retraining.
4. **Empirical results:** When we evaluated semantic similarity on our golden dataset, it had 'no practical utility' - near-random precision. Heuristics achieved 99%+ precision.

We *did* use ML where appropriate: CLIP for image similarity and YOLOv8 for logo detection, where embeddings and object detection excel. The key was choosing the right tool for each signal."

5.1.2 Q: How did you handle the massive backfill of hundreds of millions of games?

Answer:

"I designed separate MLP pipelines per feature - title, thumbnail, logo - so we could run them independently and iterate on one without reprocessing all. The architecture:

1. Per-Feature Pipelines:

- Each loads games corpus from Hive
- Loads PAL for that feature type
- Runs similarity detection (text heuristics or vector search)
- Outputs to intermediate Hive tables: `sdp_title_similarity`, `sdp_icon_similarity`

2. Powerhouse Aggregation Pipeline:

- Aggregates scores from intermediate tables
- Publishes final `is_clickbait` boolean to Frost in batch

This modular design had three benefits:

- **Independent scaling:** Could scale compute per feature based on workload (CLIP inference vs simple string matching)
- **Fault tolerance:** If one pipeline failed, didn't need full rerun - just reprocess that feature
- **Iterability:** Could improve title matching heuristics without re-running CLIP embeddings on billions of images

We coordinated with ML Platform team to ensure adequate CLIP inference capacity for the backfill burst."

5.1.3 Q: What metrics did you use to evaluate success?

Answer:

"We had multi-level metrics across model performance, business impact, and operations:

Model Performance (Precision/Recall):

- Golden dataset evaluation on 2000+ labeled examples
- Per-signal accuracy: Title (99%), CLIP (95%), YOLO+CLIP (99.5%)
- Confusion matrix analysis to identify systematic errors

Business Impact:

- CEO satisfaction: Problem visibly resolved for top queries like 'Brookhaven'
- Query coverage: Top 30 protected games covered 20-30% of search queries (Pareto principle validated)
- PAL growth: Started with ~30 entries, grew to hundreds (system scaling successfully)
- Creator feedback: Zero backlash (policy clarity worked)

Operational Metrics:

- Real-time latency: P50/P99 for end-to-end detection
- Throughput: Sustained ~10 QPS game updates without queue buildup
- Backfill completion: Time to process all existing games
- False positive rate: Manual override API usage as proxy

The combination gave us confidence in both technical correctness and business value."

5.1.4 Q: How did you ensure high precision to avoid false positives?

Answer:

"False positives were our biggest risk - flagging legitimate games would hurt creators and damage platform trust. I used a multi-layered approach:

1. Conservative Thresholds:

- Started with stringent similarity thresholds (high confidence bar)
- Fine-tuned per PAL entry based on that asset's visual distinctiveness
- Used confusion matrix to optimize precision vs recall trade-off

2. Human-in-the-Loop:

- Label Studio pipeline for ongoing evaluation
- High-profile games get human review before PAL addition
- Feedback loop: flagged games reviewed weekly, errors fed back to threshold tuning

3. Nuanced Policies:

- Built-in exceptions: fan-made content, same creator ownership
- Flexible match policies: 'Evade' uses exact (common word), 'Brookhaven' uses fuzzy=2
- Manual override API for edge cases (parodies, legitimate references)

4. Ranking Demotion, Not Filtering:

- Demoted in ranking rather than hard-filtered
- Users still have choice; we just deprioritize suspected clickbait
- Reduced impact of false positives

The result: 99%+ precision with zero creator backlash, which validated our approach."

5.1.5 Q: Walk me through your collaboration with the Brand Safety team on logo detection.

Answer:

"The Brand Safety team had an existing logo detection model that hadn't been updated since deployment. I needed to integrate it while planning for future model evolution. Here's how I approached it:

Discovery Phase:

- Met with Brand Safety team to understand their model (YOLOv8-based)
- Identified constraint: model was static, no update cadence
- Learned their model detected generic logos, not game-specific ones

Integration Strategy:

- Used their endpoint for initial logo region detection (bounding boxes)
- Fine-tuned my own YOLOv8 on Roblox-specific logo corpus (using Label Studio dataset)
- Two-stage pipeline: their model for generic detection → my model for game logo classification

Future-Proofing:

- Established versioning protocol: if they update model, maintain old endpoint
- Create new versioned endpoint (/v2/detect-logos)
- Enables A/B testing before migration
- Asked if they could incorporate PAL into their training (potential future improvement)

Result:

- Smooth integration, 99.5%+ precision on logo-based clickbait
- Reusable pattern for future cross-team ML integrations
- Brand Safety team appreciated the versioning protocol (adopted it for other consumers)

This taught me that cross-team ML integration requires planning for model evolution, not just initial integration."

5.2 Systems Design Questions

5.2.1 Q: Why event-driven architecture for real-time detection? What are the trade-offs?

Answer:

"I chose event-driven architecture for three reasons:

Benefits:

1. **Efficiency:** Only process games that changed (vs polling all games periodically)

- Game updates are sparse events (~10 QPS vs millions of games)
- Polling would waste 99.9% of compute checking unchanged games

2. **Low Latency:** React immediately to changes (10-second end-to-end)

- Users see updated ranking within seconds of game publish
- vs batch: hours/days of staleness

3. **Decoupling:** SQS buffer isolates detection service from upstream events

- Handles traffic spikes without dropping events
- Detection service can scale independently
- Failed processing goes to dead letter queue for recovery

Trade-Offs:

1. **Complexity:** More moving parts (SNS, SQS, multiple services)

- Mitigation: Used mature AWS services with high availability
- Clear ownership boundaries between teams

2. **Eventual Consistency:** 10-second delay before ranking reflects changes

- Acceptable: game creation isn't blocked, creators don't see immediate impact anyway
- Much better than batch (hours/days)

3. **Ordering:** SQS doesn't guarantee FIFO (standard queue)

- Not a problem: we only care about final state, not order of updates
- Each update is idempotent (upsert to Frost by universe_id)

For this use case, the benefits far outweighed the trade-offs. Event-driven was the right choice."

5.2.2 Q: How would you handle a sudden 10x spike in game update traffic?

Answer:

"The architecture has natural buffers, but I'd implement additional safeguards:

Existing Resilience:

- SQS queue absorbs spikes (messages wait until processed)
- Detection service can scale horizontally (more bedev2 instances)

- Frost writes are async (doesn't block upstream)

Additional Mitigations for 10x Spike:

1. Rate Limiting at SQS:

- Set max queue depth threshold (e.g., 100K messages)
- If exceeded, temporarily drop low-priority events (minor game updates)
- Prioritize new game creation events over minor metadata changes

2. Auto-Scaling Detection Service:

- CloudWatch alarm on SQS ApproximateNumberOfMessages
- Trigger auto-scaling policy: add bedev2 instances when queue depth \geq 10K
- Scale down when queue drains to save cost

3. Batching at Frost:

- Instead of per-message Frost write, batch 100 updates
- Reduces Frost write load by 100x
- Trade-off: slightly higher latency (batch flush interval)

4. Circuit Breaker for Downstream:

- If CLIP/YOLO inference services are overloaded, fail fast
- Requeue to SQS for retry with exponential backoff
- Prevents cascading failures

5. Monitoring & Alerting:

- P99 latency alerts (if \geq 60 seconds, something's wrong)
- Queue depth alerts (if \geq 50K, need manual intervention)
- Dead letter queue alerts (if messages failing, investigate)

The key principle: **graceful degradation**. Accept slightly higher latency under extreme load, but never drop data or crash services.”

5.3 Project Management Questions

5.3.1 Q: How did you convince stakeholders to abandon 2 months of work?

Answer:

”This was delicate - I had to show the problem clearly without making the previous team look bad. Here's how I approached it:

1. Data-First Approach (Week 1):

- Immediately ran human evaluation on their CLIP results
- Tested multiple thresholds (0.9, 0.95) on Brookhaven case
- Result: precision \sim 50% (near-random)

- Documented: false positives (legitimate anime-style games), false negatives (modified images)

2. Root Cause Analysis:

- **Not:** "Your approach is wrong" (accusatory)
- **Instead:** "We're solving the wrong problem" (reframe)
- Showed CLIP clusters by visual style, not deceptive intent
- Identified missing signal: title duplication (strongest clickbait indicator)

3. Concrete Alternative (Not Just Criticism):

- Presented multi-signal architecture: text + image + logo
- Showed how each signal addresses specific failure modes
- Estimated timeline: M1 (title) in 3 weeks, full system in Q1

4. Meeting Dynamics:

- Invited Principal ML Engineer to the discussion (respect)
- Framed as "we discovered this together through evaluation"
- He acknowledged CLIP-only approach had limitations
- Moved to other priorities (no friction)

5. Leadership Alignment:

- CEO urgency gave me leverage (problem needs solving *now*)
- PM supported pivot after seeing evaluation data
- Committed to clear milestones (M1, M2, M3) for accountability

Key Lesson: When proposing to scrap work, lead with **data**, offer **concrete alternatives**, and **avoid blame**. Make it about the problem, not the people."

5.3.2 Q: What was the biggest challenge in coordinating 5 teams?

Answer:

"The biggest challenge was **aligning on scope and deliverables** - each team had different priorities and timelines. Here's what I did:

Challenge 1: Diverging Priorities

- DSA team wanted comprehensive backfill (all games, all time)
- Human Eval had limited bandwidth (can't label millions of games)
- ML Platform focused on inference efficiency (latency optimization)
- Game Discovery wanted fast deployment (CEO pressure)

Solution: Clear Work Streams

- Created detailed milestones doc with team ownership (M1: title by DSA, M2: thumbnail by me + Human Eval, M3: logo by Brand Safety)

- Weekly sync meetings (30 min, focused agenda)
- Slack channel for async updates (`#clickbait-detection-project`)
- **Decision-making authority:** I had final call on architecture; PM had final call on launch timing

Challenge 2: Human Eval Bottleneck

- Label Studio tasks required expert judgment (what is clickbait?)
- Human Eval team had competing priorities (safety moderation)

Solution: Prioritization + Iteration

- Phase 1: Top 30 games only (minimal viable PAL)
- Created detailed labeling guidelines to reduce ambiguity
- Automated easy cases (exact title match) to reduce human load
- Phase 2: Expand PAL based on query volume analysis

Challenge 3: Brand Safety Model Dependency

- Their logo model was static (no update schedule)
- Risk: if model degrades, we're stuck

Solution: Versioning Protocol + Fine-Tuning

- Agreed on versioned endpoints for future updates
- Fine-tuned my own YOLOv8 as backup (not fully dependent)
- Offered to share PAL data for their model retraining (mutual benefit)

Key Lesson: Cross-team projects need **clear ownership, explicit dependencies, and fall-back plans.** Over-communicate, document decisions, and reduce critical path dependencies where possible.”

5.3.3 Q: You mentioned leadership didn't fully understand the project. What happened there?

Answer (Honest but Diplomatic):

”This was a learning experience for me about the importance of **executive communication**, not just technical execution.

What Happened:

- Project had high visibility (CEO escalation, townhall recognition)
- But my direct leadership didn't fully grasp the technical complexity or business impact
- This contributed to my decision to explore new opportunities

Where I Could've Done Better:

1. Executive-Level Artifacts:

- I wrote detailed technical docs (architecture, APIs, pipelines)
- Should've also created 1-pager for leadership: Problem → Solution → Impact
- Metrics dashboard showing business value (query coverage, precision, CEO satisfaction)

2. Proactive Updates:

- I communicated in team meetings and Slack
- Should've done monthly 1:1 highlights with skip-level manager
- "Here's what we shipped, here's the impact, here's what's next"

3. Connecting to Org Goals:

- I framed as "fixing clickbait problem"
- Could've framed as "establishing IP protection framework for platform" (strategic)
- Or "improving search quality metrics" (measurable org goal)

What I Learned:

- Technical excellence alone doesn't guarantee recognition
- Senior/Principal engineers must translate technical work into business value *for leadership*
- Documentation, metrics, and storytelling are as important as code

Looking Forward:

- At DataHub, I'd apply these lessons from day one
- Create exec-friendly artifacts alongside technical design docs
- Proactively communicate impact, not just features shipped

This experience taught me that **impact without visibility is wasted effort** - something I'm now very conscious of."

6 Practice Checklist

✓ Action Item

2-Day Prep Plan:

Day 1 (3-4 hours)

- Morning (2 hours):
 - Read this document cover-to-cover
 - Highlight 5-7 key talking points to memorize
 - Practice 30-second elevator pitch (time yourself)
- Afternoon (2 hours):
 - Draw architecture diagram from memory (whiteboard/paper)
 - Practice answering "Why heuristics?" question out loud
 - Write down 3 questions to ask Abe about DataHub

Day 2 (3-4 hours)

- Morning (2 hours):
 - Review Mock Q&A section
 - Practice STAR narrative out loud (record yourself, aim for 10-12 min)
 - Identify weak spots (what questions make you hesitate?)
- Afternoon (1-2 hours):
 - Mock interview with friend/colleague
 - Practice whiteboard diagram + verbal walkthrough
 - Final review: Key Insights and Critical Points boxes

Night Before Interview

- Review 30-second pitch (know it cold)
- Skim Technical Challenges section (3-5 lessons learned)
- Prepare 3 questions for Abe (show genuine interest)
- Get good sleep!

Interview Day Morning

- Skim Executive Summary (page 2)
- Review architecture diagram (can you draw it from memory?)
- Practice: "This project taught me..." (Connection to DataHub)
- Arrive 10 min early, bring water, have paper/pen ready

7 Key Talking Points (Memorize These)

1. First Principles Reframing

- Previous team: image similarity problem
- Me: IP protection + search quality problem
- Multi-signal detection (text + image + logo), not single model

2. Flexible Policy Framework

- Three match policies: exact, substring, fuzzy
- Not one-size-fits-all: "Evade" vs "Brookhaven" need different protection
- Built-in exceptions: fan-made, same creator, typo tolerance

3. Right Tool for Each Signal

- Heuristics for text (explainable, 99% precision)
- CLIP for images (semantic similarity works here)
- YOLO for logos (object detection + fine-tuning)

4. Event-Driven Real-Time Architecture

- SNS → SQS → Detection Service → Frost → Search Ranking
- 10-second latency, non-blocking, scalable
- vs previous: weekly batch (stale, reactive)

5. 99%+ Precision

- Systematic validation on golden dataset (2000+ examples)
- Conservative thresholds, human-in-the-loop, nuanced policies
- Zero creator backlash (proof of policy success)

6. Cross-Team Leadership

- Coordinated 5 teams (DSA, Human Eval, Brand Safety, ML Platform, Game Discovery)
- Clear work streams, weekly syncs, decision-making authority
- Successful handoff to Game Discovery post-launch

7. Connection to DataHub

- PAL ≈ DataHub metadata registry
- Cross-team adoption challenges (same at DataHub)
- Real-time propagation (schema changes → lineage updates)
- Policy framework (governance at DataHub)

8 Questions to Ask Abe

1. Product Strategy:

- "What's DataHub's biggest challenge in driving enterprise adoption? Is it technical (scale, latency) or organizational (getting teams to contribute metadata)?"

2. Technical Architecture:

- "How does DataHub handle real-time metadata propagation? If a schema changes in production, how quickly do downstream systems see that in lineage graphs?"

3. Team Dynamics:

- "For this role, what's the balance between building new platform features vs enabling customer integrations? How much is greenfield vs existing codebase evolution?"

4. Search/Discovery Expertise:

- "Given my background in search and RAG systems, where do you see opportunities for me to contribute to DataHub's metadata search and discovery experience?"

5. Personal Growth:

- "What does success look like for a Principal Engineer at DataHub in the first 6-12 months? What are the key milestones or projects that would demonstrate high impact?"

Strategy: Pick 2-3 questions that genuinely interest you. Let the conversation flow naturally - you may not need to ask all of them if the discussion covers the topics organically.

9 Final Reminders

⚠ Critical Point

Interview Execution Tips:

- Start High-Level, Then Drill Down
 - 30-second summary → problem reframing → architecture → technical details
 - Let Abe guide depth (if he asks about YOLO fine-tuning, dive deep; if not, stay high-level)
- Use Whiteboard/Paper
 - Draw architecture diagram as you explain
 - Visual aids help both you and interviewer follow along
 - Shows systems thinking
- Be Honest About Challenges
 - Don't pretend everything was perfect
 - "What I'd do differently" shows self-awareness and learning
 - Leadership recognition issue: frame as learning about exec communication
- Show Enthusiasm
 - This was a project you're proud of - let that show!
 - Technical problem-solving is fun, not drudgery
 - Connect excitement about this work to excitement about DataHub
- Ask Clarifying Questions
 - If Abe asks something unclear, ask for clarification
 - Shows careful thinking, not just rushing to answer
 - "Just to make sure I understand, you're asking about X vs Y?"
- Time Management
 - Aim for 10-12 min STAR narrative (not 20 min monologue)
 - Leave time for Abe's follow-up questions (the real signal)
 - If running long, ask: "Should I keep going or would you like to dive deeper into something specific?"

💡 Key Insight

You've Got This!

You have:

- 20+ years of engineering experience
- A compelling project with measurable impact (99%+ precision, CEO recognition)
- Deep technical expertise (search, ML, distributed systems)
- Clear narrative arc (problem → insight → solution → impact)
- Strong connection to DataHub's mission (metadata quality, cross-team adoption)

Trust your experience. Be yourself. Show your passion for solving hard problems.

Good luck!

10 Part 2: Search Architecture Deep Dive

10.1 Overview

💡 Key Insight

Your Search Expertise at Roblox:

You are uniquely positioned to discuss search architecture because you've worked across the entire search stack:

- **Retrieval Layer:** RSS (billion-scale vector database), semantic search
- **Ranking Layer:** Clickbait detection signal (quality scoring)
- **Query Understanding:** Spelling correction service, autocomplete
- **Search Quality:** Game clickbait detection, evaluation frameworks
- **Platform Infrastructure:** RAG Platform (retrieval + generation)

Key Message: You're not just a user of search systems - you're a *builder of search infrastructure* that others depend on.

10.2 Roblox Search Architecture: High-Level Overview

10.2.1 The Roblox Search Problem Space

What Makes Roblox Search Unique:

1. Scale:

- Millions of user-generated games (experiences)
- Billions of search queries annually
- Real-time content creation (games published every second)

2. User-Generated Content (UGC) Challenges:

- Noisy metadata (creators gaming SEO with keyword stuffing)
- Quality variance (AAA games vs low-effort clones)
- IP violations (clickbait, copycats)
- Sparse signals for new content (cold start problem)

3. Diverse User Base:

- Age range: 6-60+ years old
- Query types: navigational ("Brookhaven"), exploratory ("fun games"), developmental ("obby")
- Intent: play now, discover new, find friends' games

10.2.2 End-to-End Search Flow

Draw this architecture during the interview:

```
User Query: "Brookhaven"
|
[1] Query Understanding
|--- Spelling Correction (your service)
|--- Query Normalization (lowercase, trim)
|--- Intent Detection (navigational vs exploratory)
\--- Query Expansion (synonyms, related terms)
|
[2] Retrieval (Multi-Stage)
|--- [2a] Keyword Retrieval (Elasticsearch)
|   |--- Title/Description match (BM25 scoring)
|   \--- Returns: Top 1000 candidates
|--- [2b] Semantic Retrieval (RSS - your system)
|   |--- Query -> CLIP embedding
|   |--- Vector search against game thumbnails
|   \--- Returns: Top 500 candidates
\--- [2c] Hybrid Fusion
    \--- Merge keyword + semantic results -> Top 200 candidates
|
[3] Ranking
|--- [3a] Feature Generation (Frost)
|   |--- Engagement: CTR, play time, favorites
|   |--- Quality: clickbait score (your system), creator reputation
|   |--- Freshness: publish date, update recency
|   |--- Personalization: user history, age appropriateness
|--- [3b] Learning-to-Rank (LTR) Model
|   |--- Gradient-boosted trees (XGBoost/LightGBM)
|   |--- Trained on user feedback (clicks, play sessions)
|   \--- Outputs: relevance score per candidate
\--- [3c] Re-Ranking & Filtering
```

```

    |-- Diversity: avoid duplicate creators in top 10
    |-- Safety: filter age-inappropriate content
    \-- Policy: demote clickbait (your system)
    |
[4] Serving
    |-- Cache: Redis for popular queries
    |-- Latency: P99 < 500ms
    \-- A/B Testing: gradual rollout of ranking changes
    |
Search Results Displayed to User

```

10.3 Your Contributions to Roblox Search

10.3.1 1. RSS (Roblox Similarity Search) - Vector Database Infrastructure

Your Role: Initiator and lead engineer (2022-2023, V1)

Problem Solved:

- Roblox needed billion-scale vector search for semantic retrieval
- Off-the-shelf solutions (Pinecone, Weaviate) didn't meet requirements:
 - Cost at Roblox scale ($\sim 100M+$ vectors)
 - Latency for real-time serving (P99 \downarrow 100ms)
 - Data residency (must run in Roblox data centers)

Technical Approach:

1. Ported Milvus (Open Source) to Roblox Infrastructure

- Milvus: distributed vector database (built on FAISS)
- Adapted for Roblox's Nomad-based microservice environment
- Integration with Roblox storage layer (S3, persistent volumes)

2. Deployment Architecture

- Distributed across Roblox data centers
- Sharded by collection (games, avatars, assets)
- HNSW index for approximate nearest neighbor (ANN) search
- Replica sets for high availability

3. Scale Metrics

- Billions of vectors indexed
- P99 latency: $\sim 50\text{-}100\text{ms}$ for k-NN search ($k=100$)
- ~ 50 production use cases across Roblox (search, recommendations, safety)

First Major Use Case: Avatar Marketplace Semantic Search

- **Problem:** Keyword search failed for visual queries ("red hoodie", "anime hair")

- **Solution:**

- Generated CLIP embeddings for all avatar items (images)
- Indexed in RSS (multi-modal: text query → image results)
- Hybrid retrieval: keyword (Elasticsearch) + semantic (RSS)

- **Impact:** Improved discovery for long-tail items (less popular items surfaced)

Trade-Offs Discussion (Critical for Interview):

Decision	Benefits	Trade-Offs
Milvus vs Build from Scratch	Faster time-to-market, battle-tested at scale	Less control, vendor lock-in (mitigated by open source)
HNSW vs FLAT index	Sub-100ms latency (ANN)	Lower recall than exhaustive search (acceptable for top-k)
In-house vs SaaS	Cost savings, data residency, customization	Operational burden (maintenance, upgrades)
Sharding by collection	Isolation, independent scaling per use case	Cross-collection queries require multiple searches

10.3.2 2. RAG Platform - Retrieval-Augmented Generation

Your Role: Lead architect and engineer (2024, V1 MVP)

Problem Solved:

- Multiple teams building ad-hoc RAG systems (AI Safety, DevRel, ROS)
- Duplicated effort: data ingestion, embedding generation, vector DB integration
- No standardized evaluation or monitoring

Technical Approach:

1. Ported AnythingLLM (Open Source) to Roblox BEDEV2

- AnythingLLM: full-stack RAG platform (document processing, embeddings, LLM integration)
- Adapted for Nomad-based microservice deployment
- Integration with Roblox auth, data access controls

2. Platform Components

- **Data Ingestion:** PDF, DOCX, CSV parsers; chunking strategies
- **Embedding Generation:** CLIP, Sentence-BERT, domain-specific models
- **Vector Storage:** RSS integration (reused your vector DB)
- **LLM Orchestration:** Multi-LLM support (GPT-4, Claude, internal models)
- **Agent Framework:** Tool use, memory, prompt templates

3. Web Platform (No-Code RAG)

- UI for non-engineers to build RAG apps
- Upload documents → configure retrieval → deploy chatbot
- Reduced development cycle: 3 months → 1 week

Adoption & Impact:

- **10+ production use cases within 6 months**
- **Use Case 1: Moderation Mate (Safety Team)**
 - RAG over moderation policies and case histories
 - Helps moderators make consistent decisions
 - 5K QPS at peak
- **Use Case 2: Community Insights (DevRel)**
 - RAG over creator feedback, forum posts, support tickets
 - Helps DevRel team understand creator pain points
 - Research tool, not user-facing

RAG-Specific Challenges (Discuss if Asked):

1. Data Privacy & Access Control

- Problem: RAG system has access to sensitive documents
- Solution: Row-level security, document-level permissions
- Integration with Roblox's RBAC (role-based access control)

2. Retrieval Quality (Precision vs Recall)

- Problem: Retrieving irrelevant chunks hurts LLM output quality
- Solution: Hybrid search (keyword + semantic), reranking, chunk overlap
- Evaluation: human eval on relevance, LLM-as-judge for answer quality

3. Latency (P99 ~ 5 seconds for user-facing)

- Breakdown: Retrieval (100ms) + Reranking (200ms) + LLM (3-4s)
- Optimization: parallel retrieval, streaming LLM responses

10.3.3 3. Autocomplete Services (Avatar, Studio, Brand)

Your Role: Sole engineer (Avatar: 2020, Studio: 2021, Brand: 2025 as mentor)

Problem Solved:

- Users need instant query suggestions as they type
- Reduce typos, guide users to popular content
- Different domains: Avatar Marketplace, Studio (creator tools), Brand partnerships

Technical Approach:

1. Trie-Based Data Structure

- Prefix tree for fast prefix matching ($O(k)$ for query length k)
- Precomputed suggestions for top 10K prefixes
- In-memory (Redis) for P99 \leq 10ms latency

2. Suggestion Ranking

- Signals: query frequency, result CTR, recency
- Personalization: user history, age group
- A/B testing framework for ranking changes

3. Offline Pipeline (Daily)

- Aggregate query logs (Hive)
- Filter: remove PII, profanity, low-frequency queries
- Build trie, compute scores, publish to Redis

Scale Metrics:

- Avatar Autocomplete: All Roblox users ($\sim 70M$ DAU)
- Studio Autocomplete: Roblox creators ($\sim 2M+$ monthly)
- P99 latency: 10ms (critical for user experience)
- Update frequency: Daily (captures trending queries)

10.3.4 4. Spelling Correction Service

Your Role: Sole engineer (2022)

Problem Solved:

- User queries contain typos ("Brookhven" \rightarrow "Brookhaven")
- Children are primary users (higher typo rate)
- Poor spelling = zero results = bad user experience

Technical Approach:

1. Dictionary-Based Correction

- Dictionary: all valid game titles, popular queries
- Algorithm: Edit distance (Levenshtein), phonetic similarity (Soundex)
- Candidate generation: words within edit distance 2

2. Context-Aware Ranking

- Given "Brookhven", candidates: ["Brookhaven", "Brookhaven RP", "Brook Haven"]
- Rank by: query frequency, result availability, user context

3. Production Deployment

- Microservice (bedev2), REST API
- Integration: Search (query rewriting), Autocomplete (suggestion generation)
- 3 production adoptions: Avatar Marketplace, Studio, Community

10.3.5 5. Game Clickbait Detection (Covered in Part 1)

Search Quality Signal:

- Clickbait score (0-1) stored in Frost feature table
- Used by ranking model as quality signal (demote low-quality copycats)
- Real-time updates (10-second latency from game publish to ranking)

10.4 Key Search Topics: Deep Dive

10.4.1 Indexing Pipeline

Roblox Game Indexing Flow:

```

Game Created/Updated Event (SNS)
|
[1] Document Builder Service
|-- Fetch metadata: title, description, creator, genre, etc.
|-- Generate features: engagement stats, quality scores
|-- Enrich: clickbait score (your system), safety flags
\-- Output: JSON document
|
[2] Elasticsearch Indexing
|-- Index mapping: title (text), description (text), genre (keyword)
|-- Sharding: by universe_id (game ID)
|-- Replica sets: 3x for availability
\-- Refresh interval: 1 second (near real-time)
|
[3] Vector Indexing (RSS)
|-- Generate CLIP embedding for game thumbnail
|-- Index in RSS (separate collection: "game_thumbnails")
\-- Async (doesn't block Elasticsearch indexing)
|
[4] Feature Store (Frost)
|-- Compute ranking features (engagement, quality, freshness)
|-- Publish to Frost tables
\-- Used by ranking service at query time

```

Indexing Challenges:

1. Freshness vs Consistency

- Problem: Game metadata changes frequently (title updates, thumbnail swaps)
- Solution: Near real-time indexing (1-second refresh) vs eventual consistency
- Trade-off: Slight staleness acceptable for better performance

2. Schema Evolution

- Problem: Adding new fields (e.g., clickbait_score) requires reindexing

- Solution: Backward-compatible mappings, gradual rollout
- Zero-downtime reindexing: alias swap pattern

3. Scale (Millions of Games)

- Problem: Full reindex takes hours/days
- Solution: Incremental indexing (only changed documents)
- Backfill: parallel workers, checkpointing for recovery

10.4.2 Query Understanding

Roblox Query Understanding Pipeline:

1. Normalization

- Lowercase, trim whitespace
- Remove special characters (emojis, punctuation)
- Unicode normalization ($\acute{e} \rightarrow e$)

2. Spelling Correction (Your Service)

- Detect misspellings: edit distance, phonetic similarity
- Suggest corrections: "Brookhven" \rightarrow "Brookhaven"
- User experience: "Did you mean..." vs auto-correct

3. Intent Detection

- **Navigational:** User wants specific game ("Adopt Me")
- **Exploratory:** User wants category ("scary games", "obby")
- **Transactional:** User wants to create/edit ("new game template")

4. Query Expansion

- Synonyms: "scary" \rightarrow ["horror", "spooky", "creepy"]
- Related terms: "obby" \rightarrow ["obstacle course", "parkour"]
- Learned from user behavior (query reformulations)

Challenge: Roblox-Specific Terminology

- "Oddy" = obstacle course (Roblox slang, not in standard dictionaries)
- "Brookhaven RP" = roleplay (RP is domain-specific abbreviation)
- Solution: Custom dictionary, query log mining, creator tagging

10.4.3 Ranking & Relevance

Ranking Signal Categories:

1. Textual Relevance (Traditional IR)

- BM25 score (keyword match in title, description)
- TF-IDF (term frequency, inverse document frequency)
- Exact match boost (query = title → higher score)

2. Engagement Signals

- CTR (click-through rate): clicks / impressions
- Play time: average session duration
- Favorites: user saves to favorites list
- Retention: users return to game repeatedly

3. Quality Signals

- Clickbait score: 0-1 (your system, lower = higher quality)
- Creator reputation: established vs new creator
- Safety flags: age-appropriate, policy violations
- User ratings: thumbs up/down

4. Freshness Signals

- Publish date: new games get temporary boost
- Update recency: recently updated games prioritized
- Decay function: boost diminishes over time

5. Personalization Signals

- User history: games played, genres preferred
- Social: friends' games, group memberships
- Age appropriateness: filter by user's age group
- Language: localization, regional preferences

Learning-to-Rank (LTR) Model:

- **Model Type:** Gradient-boosted decision trees (XGBoost/LightGBM)
- **Training Data:** User interactions (clicks, play sessions, favorites)
- **Features:** ~50-100 features across categories above
- **Objective:** Maximize engagement (play time, retention) while maintaining quality
- **Evaluation:** Offline (NDCG, MRR) + Online (A/B tests on CTR, play time)

Trade-Off: Engagement vs Quality

- **Problem:** Clickbait games have high CTR (by design) but low quality
- **Naive approach:** Rank purely by CTR → clickbait dominates
- **Solution:** Multi-objective optimization
 - Objective 1: Maximize engagement (CTR, play time)
 - Objective 2: Maximize quality (low clickbait score, high ratings)
 - Weighted combination: $\alpha \cdot \text{engagement} + (1 - \alpha) \cdot \text{quality}$
- **Your contribution:** Clickbait score as quality signal (part of Objective 2)

10.4.4 Semantic Search (CLIP + RSS)

Why Semantic Search?

- Keyword search fails for visual/conceptual queries:
 - "red hoodie" (visual attribute, not in text metadata)
 - "anime style" (artistic style, hard to describe in keywords)
 - "games like Brookhaven" (conceptual similarity, not keyword match)

Semantic Search Pipeline:

```
User Query: "anime style games"
|
[1] Query Embedding
|-- CLIP text encoder: query -> 512-dim vector
|-- Normalize: L2 norm = 1
\-- Output: query embedding
|
[2] Vector Search (RSS)
|-- k-NN search: find top-k nearest game thumbnail embeddings
|-- Distance metric: cosine similarity (dot product for normalized vectors)
|-- Index: HNSW (approximate nearest neighbor)
\-- Returns: Top 500 candidates with similarity scores
|
[3] Hybrid Fusion
|-- Keyword results (Elasticsearch): Top 1000
|-- Semantic results (RSS): Top 500
|-- Fusion method: Reciprocal Rank Fusion (RRF)
\-- Output: Top 200 candidates for ranking
```

Hybrid Search Trade-Offs:

CLIP Model Details:

- **Base model:** OpenAI CLIP (ViT-B/32 or ViT-L/14)
- **Fine-tuning:** Considered but not done in V1 (off-the-shelf worked well)
- **Future:** Fine-tune on Roblox-specific data (game thumbnails + user clicks)

Approach	Strengths	Weaknesses
Keyword Only	Precise for exact matches, fast	Misses conceptual similarity
Semantic Only	Finds conceptual matches	Misses exact matches, slower
Hybrid (Keyword + Semantic)	Best of both worlds	Increased complexity, tuning needed

10.4.5 Personalization

Cold Start Problem (New Users):

- **Problem:** No user history → can't personalize
- **Solution:**
 - Default to popularity (most played games globally)
 - Age-based filtering (show age-appropriate content)
 - Gradual learning: collect clicks/plays, update profile

Cold Start Problem (New Games):

- **Problem:** No engagement data → can't rank well
- **Solution:**
 - Freshness boost: new games get temporary ranking boost
 - Creator reputation: leverage creator's past game quality
 - Content-based features: genre, description, thumbnail quality

Privacy Considerations:

- User history stored with consent (privacy policy)
- Anonymized for model training (no PII)
- User control: "Clear history" option

10.4.6 Evaluation & Metrics

Offline Metrics (Pre-Deployment):

1. **NDCG (Normalized Discounted Cumulative Gain)**
 - Measures ranking quality (higher-ranked relevant results = better)
 - NDCG@10: focus on top 10 results (most important for users)
2. **MRR (Mean Reciprocal Rank)**
 - Position of first relevant result (1/rank)
 - Good for navigational queries ("Adopt Me" → should be #1)

3. Precision@k / Recall@k

- Precision@10: % of top 10 results that are relevant
- Recall@10: % of relevant results found in top 10

Online Metrics (A/B Testing):

1. CTR (Click-Through Rate)

- Clicks / Impressions (how often users click results)
- Broken down by position (position 1 vs position 10)

2. Play Time

- Average session duration after search
- Measures result quality (did user find what they wanted?)

3. Zero Result Rate

- % of queries with no results
- Target: <5% (spelling correction helps)

4. Query Reformulation Rate

- % of users who search again after initial query
- High rate = initial results unsatisfactory

Human Evaluation:

- Periodic manual review of top queries (top 1000)
- Raters assess: relevance, quality, diversity
- Identify edge cases, systematic errors
- Feed back into model training, feature engineering

10.5 Search Architecture Trade-Offs

10.5.1 Latency vs Quality

Component	Latency Impact	Quality Impact
Retrieval Depth	More candidates = slower	More candidates = better recall
Ranking Model Complexity	Deeper trees = slower	More features = better accuracy
Semantic Search	Vector search adds 50-100ms	Improves conceptual matches
Personalization	User profile lookup adds 20ms	Improves relevance

Roblox's Choice:

- Target P99 latency: 500ms (acceptable for search)
- Budget breakdown: Retrieval (100ms) + Ranking (200ms) + Overhead (200ms)
- Trade-off: Accept slightly slower for better quality (user retention vs speed)

10.5.2 Freshness vs Consistency

- **Real-time indexing:** Game available in search within seconds
- **Eventual consistency:** Ranking features (engagement) update hourly/daily
- **Acceptable staleness:** Users don't expect instant engagement stats

10.5.3 Recall vs Precision

- **High recall (cast wide net):** Retrieve 1000 candidates, risk irrelevant results
- **High precision (narrow focus):** Retrieve 100 candidates, risk missing relevant results
- **Roblox's approach:** High recall in retrieval (1000), precision in ranking (filter to top 20)

10.5.4 Personalization vs Diversity

- **Problem:** Over-personalization creates filter bubbles
 - User only sees games similar to past plays
 - Misses discovery of new genres, creators
- **Solution:** Inject diversity
 - 70% personalized results + 30% diverse/popular results
 - Diversity constraints: no more than 2 games from same creator in top 10
 - Exploration: occasionally show random high-quality games

10.6 Potential Interview Questions & Answers

10.6.1 Q: Walk me through what happens when a user types a search query.

Answer (End-to-End Flow):

"Let me walk through a concrete example: user searches 'scary games'.

Stage 1: Query Understanding (50ms)

- Normalize: lowercase, trim spaces
- Check spelling: 'scarry' → 'scary' (my spelling service)
- Detect intent: exploratory (not navigational)
- Expand: add synonyms ['horror', 'spooky', 'creepy']

Stage 2: Retrieval (100ms)

- Keyword retrieval (Elasticsearch):
 - Match 'scary' in title, description, tags
 - BM25 scoring: games with 'scary' in title rank higher
 - Returns: Top 1000 candidates
- Semantic retrieval (RSS - my system):

- Generate CLIP embedding for 'scary games' text
- Vector search: find games with similar thumbnail embeddings
- Returns: Top 500 candidates
- Hybrid fusion: Merge results using Reciprocal Rank Fusion
 - Combines keyword + semantic scores
 - Output: Top 200 candidates for ranking

Stage 3: Ranking (200ms)

- Feature generation (Frost):
 - Engagement: CTR, play time, favorites
 - Quality: clickbait score (my system), creator reputation
 - Freshness: publish date, update recency
 - Personalization: user's age group, past scary game plays
- LTR model inference:
 - XGBoost model with ~100 features
 - Outputs relevance score per candidate
- Re-ranking & filtering:
 - Diversity: max 2 games from same creator in top 10
 - Safety: filter age-inappropriate content for this user
 - Policy: demote clickbait (score $<$ 0.8 threshold)

Stage 4: Serving (50ms)

- Check cache (Redis): popular queries cached for 5 minutes
- Format results: thumbnails, titles, metadata
- Log: query, user, results for offline analysis
- Return: Top 20 results to client

Total Latency: P99 \sim 400-500ms (within budget)"

10.6.2 Q: How do you handle very popular queries that could overwhelm the system?

Answer:

"Popular queries like 'Adopt Me' or 'Brookhaven' can spike to thousands of QPS. We use multi-layer caching:

Layer 1: CDN Edge Cache

- Results cached at edge nodes (close to users geographically)
- TTL: 5 minutes (balance freshness vs load)

- Invalidation: on ranking model updates or game policy changes
- Benefit: Offloads 80-90% of popular query traffic

Layer 2: Application Cache (Redis)

- Full result set cached in Redis (including ranking scores)
- TTL: 1 minute (more aggressive, needs fresher data)
- Personalization: cache key includes user age group (coarse-grained)
- Benefit: Handles cache misses from Layer 1

Layer 3: Partial Result Cache

- Cache retrieval candidates (top 200) separately from ranking
- Ranking still computed per-request (for personalization)
- Benefit: Saves expensive retrieval step, allows real-time personalization

Graceful Degradation:

- If ranking service is overloaded (P99 > 1s):
 - Fall back to simpler ranking (BM25 only, skip LTR model)
 - Still functional, just less optimal
- If Elasticsearch is overloaded:
 - Serve stale cache (extend TTL to 10 minutes)
 - Better than showing errors to users

Rate Limiting:

- Per-user rate limit: 10 queries/second (prevent abuse)
- Per-IP rate limit: 100 queries/second (prevent DDoS)

This multi-layer approach ensures we can handle traffic spikes without degrading user experience.”

10.6.3 Q: What's your strategy for ranking results for a brand new user?

Answer:

”New users have no history, so we use a combination of defaults and rapid learning:

Initial Strategy (Session 1):

• Age-based filtering:

- User age (from account signup) determines content appropriateness
- Age 6-12: family-friendly games only
- Age 13+: broader content, still filtered by safety policies

- **Popularity as proxy:**

- Default to globally popular games (high play time, favorites)
- Assumption: popular games are high-quality, good introduction

- **Diversity:**

- Show diverse genres (roleplay, obby, tycoon, simulator)
- Goal: learn user preferences quickly through exploration

Rapid Learning (Sessions 2-10):

- **Implicit feedback:**

- Track clicks: which games did user click?
- Track play time: which games did user play for \geq 5 minutes?
- Track favorites: which games did user save?

- **Genre preference detection:**

- If user plays 3+ roleplay games \rightarrow boost roleplay genre
- If user plays 0 scary games \rightarrow reduce scary genre

- **Creator affinity:**

- If user plays multiple games from Creator X \rightarrow boost Creator X's other games

Long-Term Personalization (Sessions 10+):

- Full personalization model (user embedding, collaborative filtering)
- Social signals: friends' games, group memberships
- Temporal patterns: play more tycoon games on weekends

Balancing Exploration vs Exploitation:

- **Exploitation:** Show games similar to past plays (80%)
- **Exploration:** Inject random high-quality games (20%)
- Prevents filter bubble, allows discovering new interests

The key is **gradual personalization**: start with safe defaults, learn quickly from early interactions, then fully personalize.”

10.6.4 Q: How do you detect and fix search quality issues?

Answer:

"Search quality degrades over time (user behavior changes, new content, model drift). We use multiple detection mechanisms:

Detection Layer 1: Automated Monitoring

- **Query-level metrics:**

- Zero result rate (alert if $\downarrow 10\%$)
- CTR drop (alert if drops $\downarrow 5\%$ week-over-week)
- Query reformulation rate (alert if $\downarrow 30\%$)

- **System-level metrics:**

- Latency (alert if P99 $\downarrow 1s$)
- Error rate (alert if $\downarrow 1\%$)

Detection Layer 2: Human Evaluation

- **Monthly audits:**

- Sample 100 top queries (cover 50%+ of search volume)
- Human raters assess: relevance (1-5 scale), quality, diversity
- Compare to baseline (previous month, competitor platforms)

- **Issue categorization:**

- Retrieval issues: relevant games not retrieved
- Ranking issues: relevant games retrieved but ranked poorly
- Quality issues: low-quality games ranked highly

Detection Layer 3: User Feedback

- "Report search result" button (crowdsourced quality)
- Support tickets mentioning search problems
- Social media monitoring (Twitter, Discord complaints)

Root Cause Analysis (Example: Clickbait Problem):

1. **Symptom:** CEO reports poor results for "Brookhaven"

2. **Hypothesis generation:**

- Retrieval issue? (relevant games not found)
- Ranking issue? (copycat games ranked too high)
- Data quality issue? (games have misleading metadata)

3. **Investigation:**

- Manual inspection: top 20 results for "Brookhaven"

- Finding: 15/20 are copycats with identical names/images
 - Root cause: **Ranking issue** (engagement signals favor clickbait)
4. **Fix:** Clickbait detection system (my project)
- Add quality signal to ranking (clickbait score)
 - Demote games with high clickbait scores
5. **Validation:**
- A/B test: treatment group sees demoted clickbait
 - Metrics: CTR, play time, user satisfaction surveys
 - Result: CTR slightly down (clickbait has high CTR by design), but play time up (users find quality games)

Continuous Improvement Loop:

- Detect → Diagnose → Fix → Validate → Repeat
- Quarterly model retraining (new data, new features)
- Annual architecture review (e.g., switch to transformer-based ranking)

The key is **multi-signal detection** (automated + human + user feedback) and **rigorous validation** (A/B tests, not just intuition)."

10.6.5 Q: How would you add a new ranking signal to the system?

Answer:

"I'll walk through this using the clickbait score as a concrete example, since I actually did this:

Step 1: Signal Definition & Feasibility (Week 1)

- **Define signal:** Clickbait score (0-1, higher = more likely clickbait)
- **Hypothesis:** Demoting clickbait will improve search quality (play time, user satisfaction)
- **Feasibility check:**
 - Can we compute it? (Yes: text rules + CLIP + YOLO)
 - At what latency? (Real-time: 10 seconds for new games)
 - Coverage? (Top 30 games initially, expandable)

Step 2: Offline Pipeline Development (Weeks 2-4)

- Build detection service (covered in Part 1)
- Backfill existing games (hundreds of millions)
- Store in Frost feature table: `clickbait_games_features_v0`
- Validate: human eval on sample (95%+ precision)

Step 3: Feature Integration (Week 5)

- **Frost schema update:**

- Add `clickbait_score` column to ranking features table
- Ensure backward compatibility (default 0.0 for missing values)

- **Ranking service integration:**

- Modify feature generation code to fetch `clickbait_score`
- Add to LTR model input (new feature column)

Step 4: Model Training (Week 6)

- **Offline training:**

- Retrain XGBoost model with new feature ($100 \rightarrow 101$ features)
- Use historical data (query logs, user interactions)
- Objective: maximize play time, minimize clickbait score

- **Offline evaluation:**

- Test set: held-out queries from last month
- Metrics: NDCG@10, MRR, clickbait in top 10 (should decrease)
- Result: NDCG improves +2%, clickbait in top 10 drops -40%

Step 5: Online A/B Testing (Weeks 7-9)

- **Experiment setup:**

- Control: Current ranking (no clickbait signal)
- Treatment: New ranking (with clickbait signal)
- Traffic split: 95% control, 5% treatment (safe rollout)

- **Metrics tracked:**

- Primary: Play time (should increase)
- Secondary: CTR (may decrease, clickbait has high CTR)
- Guardrail: Zero result rate, latency (should not regress)

- **Analysis:**

- Play time: +5% (statistically significant)
- CTR: -2% (acceptable trade-off)
- User surveys: satisfaction up +3%

- **Decision:** Ship to 100%

Step 6: Gradual Rollout (Week 10)

- Ramp: 5% \rightarrow 25% \rightarrow 50% \rightarrow 100% over 2 weeks
- Monitor for issues (latency spikes, error rate)

- Rollback plan: instant revert to control if critical issue

Step 7: Post-Launch Monitoring (Ongoing)

- Weekly metrics review (play time, clickbait in top 10)
- Monthly model retraining (incorporate new data)
- Iterate on clickbait detection (expand PAL, improve precision)

Key Lessons:

- **Validate early:** Offline eval before A/B (avoids wasted user exposure)
- **Gradual rollout:** Catch issues before full launch
- **Multiple metrics:** Don't optimize for single metric (CTR alone would favor clickbait)
- **Cross-team coordination:** Needed DSA for backfill, Game Discovery for ranking integration

”

10.6.6 Q: What's your biggest challenge with search at Roblox's scale?

Answer:

”The biggest challenge is the **tension between freshness and quality signals**.

The Problem:

- **New games** (created every second):
 - Need to surface quickly (creators expect instant visibility)
 - But have no engagement data (no CTR, play time, favorites)
 - Can't rely on quality signals (clickbait detection requires metadata)
- **Quality signals** (engagement, ratings):
 - Take time to accumulate (days/weeks)
 - New games start at a disadvantage (low ranking)
 - Creates barrier for new creators (can't break through)

Specific Example (Illustrative):

- Creator X publishes high-quality game at 10:00 AM
- Game is indexed in Elasticsearch within seconds (real-time)
- But ranking model sees:
 - Engagement: 0 clicks, 0 play time (no data yet)
 - Quality: 0 favorites, no clickbait score yet (processing delay)
 - Creator reputation: unknown (new creator)
- Result: Game ranks poorly despite being high-quality

- Creator frustrated: "My game doesn't show up in search!"

Our Mitigations (Partial Solutions):

1. Freshness Boost:

- New games (published \leq 24 hours) get temporary ranking boost
- Decays exponentially over time (day 1: +10%, day 7: +1%)
- Gives new games a chance to accumulate engagement

2. Content-Based Signals (No Engagement Required):

- Thumbnail quality (CLIP embedding \rightarrow aesthetic score)
- Description completeness (has detailed description?)
- Genre tagging (properly categorized?)
- These signals available immediately

3. Creator Reputation Transfer:

- If Creator X has 3 high-quality games previously, new game inherits reputation
- Assumption: good creators make good games
- Helps established creators, but not new ones

4. Rapid Engagement Acceleration:

- Show new games in "New & Trending" section (separate from main search)
- Users browsing this section provide early engagement signals
- Signals feed back into main search ranking within hours

Why This Remains Hard:

- **Cold start is unsolvable:** Can't predict quality without data
- **Gaming the system:** Freshness boost incentivizes spam (publish many low-quality games)
- **Balance:** Too much freshness boost \rightarrow low-quality games surface; too little \rightarrow new creators can't break through

Future Directions I'd Explore:

• Pre-launch quality prediction:

- Analyze game before publish (in Studio editor)
- Predict quality from: code complexity, asset quality, playtesting data
- Give creators feedback: "Add more detail to your description to improve searchability"

• Multi-armed bandit:

- Treat new game ranking as exploration-exploitation problem
- Dynamically adjust freshness boost based on early engagement
- If game gets high CTR in first hour \rightarrow increase boost; if low CTR \rightarrow decrease boost

This challenge touches multiple areas: ranking algorithms, creator incentives, platform economics - exactly the kind of **cross-functional systems problem** I enjoy solving."

10.7 Connection to DataHub (Search Context)

💡 Key Insight

How Your Search Experience Applies to DataHub:

"DataHub's core product is **metadata search and discovery** - finding the right dataset, understanding lineage, discovering related data. This maps directly to my search experience:

- **Roblox Game Search ↔ DataHub Dataset Search**

- Games have metadata (title, description, creator) ↔ Datasets have metadata (schema, owner, tags)
- Users search for games (keywords, concepts) ↔ Data scientists search for datasets (column names, business context)
- Quality signals (engagement, clickbait) ↔ Quality signals (data quality scores, usage frequency)

- **RSS (Vector Database) ↔ Semantic Dataset Search**

- CLIP embeddings for game thumbnails ↔ Embeddings for schema/column descriptions
- Semantic search: "anime style games" ↔ "customer transaction data"
- Both solve: keywords alone don't capture conceptual similarity

- **RAG Platform ↔ DataHub Documentation/Lineage**

- RAG over Roblox docs (policies, guides) ↔ RAG over dataset documentation (data dictionaries, ETL logic)
- Retrieval: find relevant context ↔ Lineage: find upstream/downstream dependencies
- Generation: answer questions ↔ Explain: "How is this field computed?"

- **Quality Signals (Clickbait Detection) ↔ Data Quality Signals**

- Game quality: clickbait score, engagement, ratings ↔ Dataset quality: completeness, freshness, usage
- Both require: real-time signal generation, integration with ranking
- Both face: trade-offs between quality and discoverability

I'm excited about DataHub because it's the **same core problem** - search, discovery, quality - but in the data infrastructure domain. The techniques I've built (vector search, RAG, quality signals) directly transfer."

A Quick Reference: Key Numbers

Metric	Value
Previous Approach Precision	~50% (near-random)
Title Heuristics Precision	99%+
CLIP Image Similarity Precision	95%+
YOLO + CLIP Logo Precision	99.5%+
Real-Time Latency	~10 seconds end-to-end
SQS Throughput	~10 QPS game updates
PAL Size (Initial)	~30-100 entries
PAL Size (Current)	Hundreds of entries
Query Coverage	20-30% of top queries
Protected Games (Top)	30 games (expandable)
Backfill Scale	Hundreds of millions of games
Project Timeline	Q1 2024 (~3 months)
Teams Coordinated	5 (DSA, Human Eval, Brand Safety, ML Platform, Game Discovery)
Creator Backlash	Zero (policy success)

B Quick Reference: Architecture Components

- **PAL (Protected Assets Library):** Centralized registry (text, image, logo)
- **SNS Topics:** PlaceEntityChangeEvents, UniverseDisplayInformationChangeEvents
- **SQS Queue:** Buffer for game update events (~10 QPS)
- **Queue Processor:** Consumes SQS, calls Detection Service
- **Detection Service (bedev2):** Multi-signal analysis (text, CLIP, YOLO)
- **RSS (Vector DB):** CLIP embeddings storage and similarity search
- **Frost Feature Store:** Table clickbait_games_features_v0
- **Elasticsearch:** Indexed for search ranking integration
- **MLP Pipelines:** Offline backfill (per-feature: title, thumbnail, logo)
- **Powerhouse:** Aggregates scores, publishes to Frost in batch

C Quick Reference: Contact Info

- **Interviewer:** Abe (DataHub)
- **Interview Type:** Day 2 Onsite - Project Deep Dive + Search Architecture
- **Format:** 45-60 minutes
- **Your Role:** Principal Software Engineer, Roblox
- **Project:** Game Clickbait Detection System (Q1 2024)