

SNAP (Snapchat) Backend Engineer

Interview Preparation

Comprehensive Problem Set with Solutions

Alex Yang

Prepared: November 2024

Contents

1 Interview Overview	3
1.1 About SNAP Inc.	3
1.2 Interview Process	3
1.3 Technical Focus Areas	4
1.4 Frequently Asked Questions	4
2 Problem 1: Shortest Path in Maze	5
2.1 Problem Description	5
2.2 Solution Approach	5
2.3 Implementation - Approach 1: BFS (No Wall Breaking)	6
2.4 Implementation - Approach 2: Dijkstra (With Wall Breaking)	6
2.5 Complexity Analysis	7
3 Problem 2: Ephemeral Message Queue	8
3.1 Problem Description	8
3.2 Solution	8
3.3 Complexity Analysis	9
4 Problem 3: LRU Cache	11
4.1 Problem Description	11
4.2 Solution	11
4.3 Complexity Analysis	12
5 System Design: Snapchat Stories	13
5.1 Requirements	13
5.2 High-Level Architecture	13
5.3 Data Model	13
5.4 System Flow	14
5.5 Optimizations	15
5.6 Scale Calculations	15

6 Interview Tips and Resources	16
6.1 Key Success Factors	16
6.2 Study Resources	16
6.3 Common Pitfalls	17
6.4 Behavioral Preparation	18
7 Final Checklist	19
7.1 Technical Preparation	19
7.2 Day Before Interview	19
7.3 During Interview	20
8 Conclusion	20

1 Interview Overview

1.1 About SNAP Inc.

⌚ Product Focus

SNAP Inc. (founded 2011, Nasdaq: **SNAP**) is a camera company that created Snapchat, a multimedia messaging app with 400M+ daily active users worldwide.

Core Products:

- **Ephemeral Content:** Photos and videos that disappear after viewing
- **AR Technology:** Advanced face filters and lenses using computer vision
- **Stories:** 24-hour temporary content timelines
- **Spotlight:** Short-form video discovery platform
- **Real-time Communication:** Low-latency messaging and video streaming

1.2 Interview Process

Based on actual candidate experiences:

1. **Recruiter Screen** (30 min): Background, interest in camera/AR/social products
2. **Technical Phone Screen** (45-60 min): 1-2 LeetCode medium problems, expects *Runnable code* (no pseudocode)
3. **Onsite/Virtual** (4-6 hours):
 - 2-4 Coding Rounds: LeetCode medium/hard, emphasis on speed and correctness
 - 1-2 System Design: Snapchat features or scalable infrastructure
 - Behavioral: Integrated throughout, values "Kind, Smart, Creative"

⚠ Critical Point

SNAP's Unique Expectations:

- **Speed Matters:** They explicitly value fast problem-solving
- **Runnable Code:** No pseudocode—must compile and run
- **Product Knowledge:** Must understand and use Snapchat
- **Values Alignment:** "Kind, Smart, Creative" evaluated throughout

1.3 Technical Focus Areas

💡 Concept Review

Most Commonly Tested Topics:

- **Graphs:** BFS/DFS, shortest path (Dijkstra, A*), grid traversal
- **Linked Lists:** Deep copy with random pointers, reversal
- **Trees:** Serialization/deserialization, traversals
- **Arrays/Stacks:** Rainwater trapping, monotonic stack
- **Hash Maps:** LRU cache, frequency counting
- **Heaps:** Priority queues, top K, scheduling
- **System Design:** CDN, real-time messaging, ephemeral content, microservices

1.4 Frequently Asked Questions

💡 Key Insight

Actual questions from Glassdoor, LeetCode, and candidate reports:

- Shortest Path in Maze (with wall breaking)
- Copy List with Random Pointers
- Serialize and Deserialize Binary Tree
- Trapping Rainwater
- LRU Cache
- Number of Islands (variations)
- Design Stories feature / Real-time messaging / AR filter system

2 Problem 1: Shortest Path in Maze

2.1 Problem Description

⌚ Product Focus

Why SNAP Asks This:

SNAP commonly asks graph/maze problems related to video games, AR navigation features, and Snap Map functionality. This appears frequently in phone screens and onsite interviews.

Given: A 2D grid where 0 represents walkable path and 1 represents wall/obstacle.

Tasks:

1. Find the shortest path from start to end position
2. Support wall-breaking (cost = 1 per wall, walk cost = 1)
3. Return minimum distance or -1 if no path exists

Example:

```

1 grid = [
2     [0, 0, 1, 0],
3     [0, 1, 0, 0],
4     [0, 0, 0, 1],
5     [1, 0, 0, 0]
6 ]
7 start = (0, 0), end = (3, 3)
8
9 shortest_path(grid, start, end) # Returns: 6
10 shortest_path_with_breaking(grid, start, end) # Returns: 5

```

Constraints:

- $1 \leq \text{rows}, \text{cols} \leq 100$
- $\text{grid}[i][j] \in \{0, 1\}$
- start and end are valid positions

2.2 Solution Approach

💡 Key Insight

Algorithm Selection:

- **BFS:** Use when all edges have equal weight (standard maze)
- **Dijkstra:** Use when edges have different costs (wall-breaking variant)
- BFS is simpler and faster for the basic problem
- Dijkstra handles the weighted variant elegantly

2.3 Implementation - Approach 1: BFS (No Wall Breaking)

```

1 from collections import deque
2 from typing import List, Tuple
3
4 def shortest_path(grid: List[List[int]],
5                  start: Tuple[int, int],
6                  end: Tuple[int, int]) -> int:
7     """BFS to find shortest path without breaking walls."""
8     if not grid or not grid[0]:
9         return -1
10
11    rows, cols = len(grid), len(grid[0])
12    if grid[start[0]][start[1]] == 1 or grid[end[0]][end[1]] == 1:
13        return -1
14
15    queue = deque([(start[0], start[1], 0)])
16    visited = {start}
17    directions = [(0, 1), (1, 0), (0, -1), (-1, 0)]
18
19    while queue:
20        r, c, dist = queue.popleft()
21
22        if (r, c) == end:
23            return dist
24
25        for dr, dc in directions:
26            nr, nc = r + dr, c + dc
27            if (0 <= nr < rows and 0 <= nc < cols and
28                grid[nr][nc] == 0 and (nr, nc) not in visited):
29                visited.add((nr, nc))
30                queue.append((nr, nc, dist + 1))
31
32    return -1

```

2.4 Implementation - Approach 2: Dijkstra (With Wall Breaking)

```

1 import heapq
2
3 def shortest_path_with_breaking(grid: List[List[int]],
4                                 start: Tuple[int, int],
5                                 end: Tuple[int, int]) -> int:
6     """Dijkstra's algorithm allowing wall breaking."""
7     if not grid or not grid[0]:
8         return -1
9
10    rows, cols = len(grid), len(grid[0])
11    heap = [(0, start[0], start[1])] # (cost, row, col)
12    visited = {}
13    directions = [(0, 1), (1, 0), (0, -1), (-1, 0)]
14
15    while heap:
16        cost, r, c = heapq.heappop(heap)
17
18        if (r, c) == end:
19            return cost
20

```

```

21     if (r, c) in visited and visited[(r, c)] <= cost:
22         continue
23     visited[(r, c)] = cost
24
25     for dr, dc in directions:
26         nr, nc = r + dr, c + dc
27         if 0 <= nr < rows and 0 <= nc < cols:
28             # Cost: 1 for walk, 1 for breaking wall
29             new_cost = cost + 1
30             if (nr, nc) not in visited or visited[(nr, nc)] > new_cost:
31                 heapq.heappush(heap, (new_cost, nr, nc))
32
33     return -1

```

2.5 Complexity Analysis

BFS Approach:

- **Time:** $O(R \times C)$ - visit each cell at most once
- **Space:** $O(R \times C)$ - queue and visited set

Dijkstra Approach:

- **Time:** $O(R \times C \times \log(R \times C))$ - heap operations
- **Space:** $O(R \times C)$ - heap and visited map

Action Item

Interview Execution:

1. Clarify: Can we break walls? How many? Cost model?
2. Start with BFS for simple case
3. Explain upgrade to Dijkstra if weighted
4. Test with examples: empty grid, no path, single cell
5. Discuss optimization: bidirectional search if asked

3 Problem 2: Ephemeral Message Queue

3.1 Problem Description

⌚ Product Focus

Snapchat's Core Feature:

Snapchat's defining characteristic is ephemeral messaging—messages that automatically disappear after viewing or after a time limit. This tests understanding of time-based data structures and efficient queue management.

Requirements:

- Messages have TTL (time-to-live)
- Auto-delete after viewing (Snap behavior)
- Auto-delete after expiration
- Efficient cleanup of expired messages

3.2 Solution

```

1 import time
2 import heapq
3 from typing import List, Dict, Optional
4 from collections import defaultdict
5
6 class EphemeralMessageQueue:
7     def __init__(self):
8         self.messages = defaultdict(list) # recipient -> [messages]
9         self.message_map = {} # message_id -> message data
10        self.expiration_heap = [] # (expires_at, message_id)
11        self.next_id = 1
12
13    def send_message(self, sender: str, recipient: str,
14                     content: str, ttl_seconds: int) -> int:
15        """Send message with TTL."""
16        message_id = self.next_id
17        self.next_id += 1
18
19        expires_at = time.time() + ttl_seconds
20        message = {
21            'id': message_id,
22            'sender': sender,
23            'recipient': recipient,
24            'content': content,
25            'expires_at': expires_at,
26            'viewed': False
27        }
28
29        self.messages[recipient].append(message)
30        self.message_map[message_id] = message
31        heapq.heappush(self.expiration_heap, (expires_at, message_id))
32
33    return message_id

```

```

34
35     def get_messages(self, recipient: str) -> List[Dict]:
36         """Get all non-expired, unviewed messages."""
37         self._cleanup_expired()
38
39         current_time = time.time()
40         valid_messages = []
41
42         for msg in self.messages[recipient]:
43             if (not msg['viewed']) and
44                 msg['expires_at'] > current_time):
45                 valid_messages.append(msg)
46
47         return valid_messages
48
49     def mark_viewed(self, recipient: str, message_id: int) -> None:
50         """Mark as viewed and delete (Snapchat behavior)."""
51         if message_id in self.message_map:
52             msg = self.message_map[message_id]
53             msg['viewed'] = True
54             # Remove from recipient's list
55             self.messages[recipient] = [
56                 m for m in self.messages[recipient]
57                 if m['id'] != message_id
58             ]
59             del self.message_map[message_id]
60
61     def _cleanup_expired(self) -> None:
62         """Remove expired messages."""
63         current_time = time.time()
64
65         while self.expiration_heap:
66             expires_at, msg_id = self.expiration_heap[0]
67             if expires_at > current_time:
68                 break
69
70             heapq.heappop(self.expiration_heap)
71             if msg_id in self.message_map:
72                 msg = self.message_map[msg_id]
73                 recipient = msg['recipient']
74                 self.messages[recipient] = [
75                     m for m in self.messages[recipient]
76                     if m['id'] != msg_id
77                 ]
78                 del self.message_map[msg_id]

```

3.3 Complexity Analysis

- `send_message`: $O(\log n)$ for heap push
- `get_messages`: $O(n + k \log k)$ where k is expired messages
- `mark_viewed`: $O(m)$ where m is messages for recipient
- **Space**: $O(n)$ for storing n messages

💡 Key Insight

Design Trade-offs:

- **Min-heap:** Efficient $O(\log n)$ insertion, $O(1)$ peek for next expiration
- **Hash map:** $O(1)$ message lookup by ID
- **Lazy deletion:** Only clean up when checking messages (avoid background threads)
- **Alternative:** Could use Redis with TTL for production scale

4 Problem 3: LRU Cache

4.1 Problem Description

Product Focus

Why SNAP Asks This:

SNAP frequently asks this problem to assess understanding of caching mechanisms critical to their infrastructure. Media content (photos/videos) must be cached efficiently for fast retrieval.

Must implement get() and put() operations in $O(1)$ time.

4.2 Solution

```

1  class ListNode:
2      def __init__(self, key=0, value=0):
3          self.key = key
4          self.value = value
5          self.prev = None
6          self.next = None
7
8  class LRUCache:
9      def __init__(self, capacity: int):
10         self.capacity = capacity
11         self.cache = {} # key -> ListNode
12         # Dummy head and tail
13         self.head = ListNode()
14         self.tail = ListNode()
15         self.head.next = self.tail
16         self.tail.prev = self.head
17
18     def _remove(self, node: ListNode) -> None:
19         """Remove node from linked list."""
20         node.prev.next = node.next
21         node.next.prev = node.prev
22
23     def _add_to_head(self, node: ListNode) -> None:
24         """Add node right after head (most recently used)."""
25         node.next = self.head.next
26         node.prev = self.head
27         self.head.next.prev = node
28         self.head.next = node
29
30     def get(self, key: int) -> int:
31         if key not in self.cache:
32             return -1
33
34         node = self.cache[key]
35         # Move to head (most recently used)
36         self._remove(node)
37         self._add_to_head(node)
38         return node.value
39
40     def put(self, key: int, value: int) -> None:
41         if key in self.cache:
42             # Update existing

```

```
43         node = self.cache[key]
44         node.value = value
45         self._remove(node)
46         self._add_to_head(node)
47     else:
48         # Add new
49         if len(self.cache) >= self.capacity:
50             # Remove LRU (before tail)
51             lru = self.tail.prev
52             self._remove(lru)
53             del self.cache[lru.key]
54
55         new_node = ListNode(key, value)
56         self.cache[key] = new_node
57         self._add_to_head(new_node)
```

4.3 Complexity Analysis

- **get:** $O(1)$ - hash lookup + linked list operations
- **put:** $O(1)$ - hash operations + linked list operations
- **Space:** $O(\text{capacity})$

Concept Review

Key Data Structure Insight:
Why Hash Map + Doubly Linked List?

- **Hash Map:** $O(1)$ lookup by key
- **Doubly Linked List:** $O(1)$ removal and insertion
- **Head:** Most recently used
- **Tail:** Least recently used (eviction candidate)

Alternative approaches (worse):

- Array + timestamp: $O(n)$ to find LRU
- Singly linked list: $O(n)$ to remove node
- OrderedDict (Python): Correct but less impressive in interviews

5 System Design: Snapchat Stories

5.1 Requirements

Design the Stories feature that powers one of Snapchat's core experiences:

- Users can post photo/video Stories (24-hour lifespan)
- Friends can view Stories in chronological order
- Stories auto-delete after 24 hours
- Support 400M+ DAU with low latency
- Handle high upload/view traffic during peak hours

5.2 High-Level Architecture

Concept Review

System Components:

1. **Upload Service:** Handle media uploads, compression
2. **Storage Layer:** S3 for media, Cassandra for metadata
3. **CDN:** Cloudflare/Akamai for global distribution
4. **Story Service:** Manage Story creation, retrieval, deletion
5. **TTL Service:** Background job to delete expired Stories
6. **Feed Service:** Generate personalized Story feeds

5.3 Data Model

Cassandra Schema:

```

1 Story {
2     story_id: UUID
3     user_id: UUID
4     media_url: String
5     created_at: Timestamp
6     expires_at: Timestamp # created_at + 24h
7     view_count: Int
8     thumbnail_url: String
9 }
10
11 StoryView {
12     story_id: UUID
13     viewer_id: UUID
14     viewed_at: Timestamp
15 }
16
17 UserFeed {
18     user_id: UUID

```

```
19     friend_stories: List<StoryMetadata> # Denormalized  
20     last_updated: Timestamp  
21 }
```

5.4 System Flow

Story Creation Flow:

1. User uploads photo/video → Upload Service
2. Compress/process media → FFmpeg workers
3. Store in S3, create Story record in Cassandra
4. Push to CDN for distribution
5. Notify friends via WebSocket/FCM
6. Update friend feeds asynchronously

Story Viewing Flow:

1. User requests friend Stories
2. Feed Service queries Cassandra for recent Stories
3. Filter expired Stories (check expires_at)
4. Serve media URLs from CDN (nearest edge location)
5. Record view in StoryView table
6. Update view_count asynchronously

TTL/Expiration Flow:

1. Background cron job runs every hour
2. Scans Stories where expires_at < current_time
3. Batch delete from Cassandra
4. Queue S3 deletions (async)
5. Update CDN cache invalidation

5.5 Optimizations

💡 Key Insight

Performance Optimizations:

- **CDN Caching:** Cache popular Stories near users (99% hit rate)
- **Adaptive Bitrate:** Serve different qualities based on bandwidth
- **Preloading:** Prefetch friend Stories in background
- **Thumbnail Generation:** Show thumbnails before video loads
- **Sharding:** Partition Cassandra by user_id for horizontal scaling
- **Feed Denormalization:** Pre-compute friend feeds for fast reads

5.6 Scale Calculations

Capacity Estimation (400M DAU):

- Average 2 Stories/user/day: 800M Stories/day
- Average Story size: 5MB (video) or 500KB (photo)
- Weighted average: 2MB per Story
- Storage: $800M \times 2MB = 1.6PB/\text{day}$
- With 24h TTL: 1.6PB active storage
- Bandwidth: $1.6PB / 86400s = 18.5 \text{ GB/s upload}$
- View rate (10x upload): 185 GB/s download

⚠ Critical Point

Bottlenecks to Address:

1. **Upload Spikes:** Peak hours (evenings) see 5-10x normal traffic
 - Solution: Auto-scaling upload workers, queue-based processing
2. **Hot Stories:** Viral content causes uneven CDN load
 - Solution: Consistent hashing, cache warming, rate limiting
3. **Delete Operations:** 800M deletions/day at 24h mark
 - Solution: Batch deletes, lazy deletion, S3 lifecycle policies

6 Interview Tips and Resources

6.1 Key Success Factors

✓ Action Item

1. Speed Matters

- SNAP explicitly values fast problem-solving
- Practice solving medium problems in 20-25 minutes
- Have templates ready for common patterns (BFS, Dijkstra, DFS)
- Type fast, think faster

✓ Action Item

2. Runnable Code Required

- No pseudocode—must compile and run
- Test with examples during interview
- Handle edge cases: null, empty, single element
- Syntax matters: missing colons/brackets = bad impression

✓ Action Item

3. Product Knowledge

- Use Snapchat daily and understand features deeply
- Know: Stories, Spotlight, Lenses, Chat, Discover, Snap Map
- Discuss trade-offs in context of their products
- Show genuine passion for camera/AR technology

6.2 Study Resources

Coding Practice:

- **LeetCode**: Company tag "Snapchat" or "SNAP" (Premium required)
- **Focus Topics**: Graphs (BFS/DFS), Trees, Linked Lists, Arrays
- **Difficulty**: Medium (70%), Hard (30%)
- **Daily Practice**: 2-3 problems per day for 4-6 weeks

System Design:

- **SNAP Engineering Blog:** eng.snap.com
- **Books:** "Designing Data-Intensive Applications" (Kleppmann)
- **Focus Areas:** CDN, real-time systems, ephemeral content, AR pipelines

Interview Experiences:

- Glassdoor SNAP reviews (filter by Backend Engineer)
- 1Point3Acres - Chinese forum with detailed reports
- Prepfully SNAP interview guide
- LeetCode Discuss: Snapchat interview threads

6.3 Common Pitfalls

⚠ Critical Point

What Gets Candidates Rejected:

1. **Slow Coding:** Taking 40+ min for a medium problem
 - SNAP values speed—practice timed coding religiously
2. **Not Testing:** Writing code without running through examples
 - Always test with 2-3 examples, including edge cases
3. **Ignoring Edge Cases:** Null inputs, empty arrays, single elements
 - Show defensive programming mindset
4. **Poor Communication:** Silent coding without explanation
 - Explain approach before coding, talk through logic
5. **Lack of Product Knowledge:** Never used Snapchat
 - Genuine interest in their products is non-negotiable

6.4 Behavioral Preparation

❑ Concept Review

SNAP's Three Core Values:

SNAP explicitly evaluates candidates on three dimensions:

1. **Kind:** Collaborative, respectful, inclusive
 - Prepare examples of helping teammates
 - Show empathy and active listening
 - Discuss inclusive engineering practices
2. **Smart:** Problem-solving, technical depth, learning mindset
 - Complex technical problems you solved
 - How you learn new technologies quickly
 - Times you debugged tricky issues
3. **Creative:** Innovation, bold ideas, thinking differently
 - Novel solutions to engineering challenges
 - Times you questioned status quo
 - Side projects showing creativity

Prepare STAR stories for:

- Times you demonstrated kindness in team settings
- Complex technical problems you solved creatively
- Learning from failures
- Why SNAP/camera/AR technology excites you
- Conflict resolution and collaboration

7 Final Checklist

7.1 Technical Preparation

✓ Action Item

Must-Practice Topics:

- ✓ Graph algorithms: BFS, DFS, Dijkstra, A*
- ✓ Tree traversals: inorder, preorder, postorder, level-order
- ✓ Linked list manipulation: reversal, cycle detection, deep copy
- ✓ Dynamic programming: 1D/2D DP, memoization
- ✓ String algorithms: sliding window, two pointers
- ✓ Design patterns: LRU cache, pub-sub, observer

7.2 Day Before Interview

✓ Action Item

Final Prep Checklist:

- ✓ Review 5-10 medium problems similar to SNAP questions
- ✓ Practice talking out loud while coding
- ✓ Test your setup: IDE, internet, video/audio
- ✓ Review SNAP products: use Snapchat, check latest features
- ✓ Prepare questions to ask interviewer about SNAP's tech stack
- ✓ Review your resume—be ready to discuss every project
- ✓ Get good sleep (critical for coding speed)

7.3 During Interview

✓ Action Item

Interview Execution:

- Ask clarifying questions (constraints, edge cases, examples)
- Explain approach before coding
- Write clean, commented code
- Test with examples as you go
- Discuss complexity analysis
- Propose optimizations if time permits
- Show enthusiasm for SNAP's products

8 Conclusion

SNAP interviews test strong algorithmic skills, system design thinking, and alignment with their "Kind, Smart, Creative" values.

💡 Key Insight

Success Formula:

1. **Speed and Accuracy:** Practice timed coding daily
2. **Graph Algorithms:** Very common at SNAP (BFS/DFS mastery)
3. **Clean Code:** Production-ready, runnable, well-tested
4. **Real-time Systems:** Understand WebSockets, CDN, caching
5. **Product Passion:** Show genuine interest in camera/AR/ephemeral content
6. **Values Alignment:** Demonstrate Kind, Smart, Creative through examples

Good luck with your SNAP interview!